

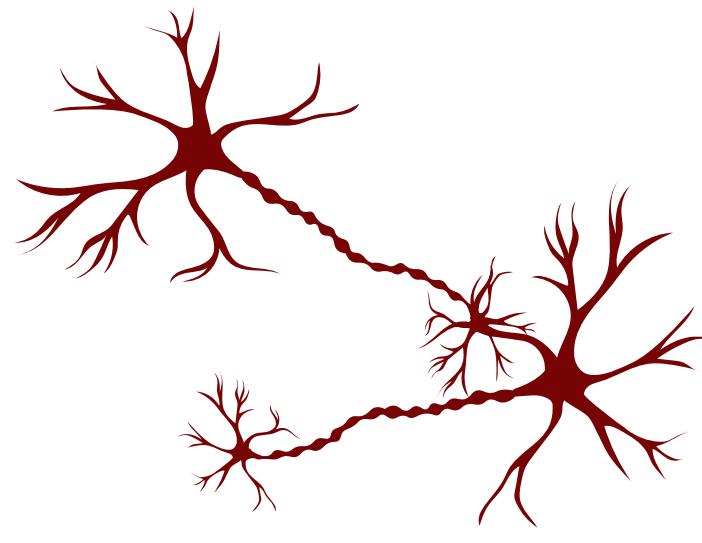
# Understanding Artificial Neural Networks

Presented by

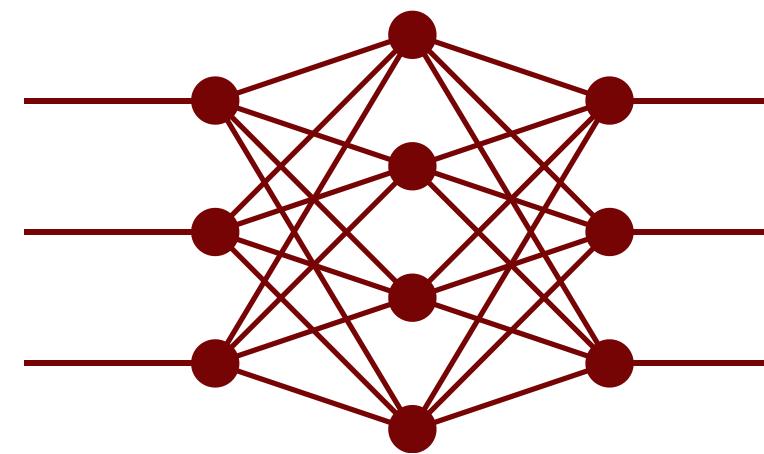


# What Are Neural Networks?

- A neural network (ANN) is a machine learning model inspired by the human brain's neural structure.
- It consists of interconnected nodes, or artificial neurons, that mimic brain neurons.

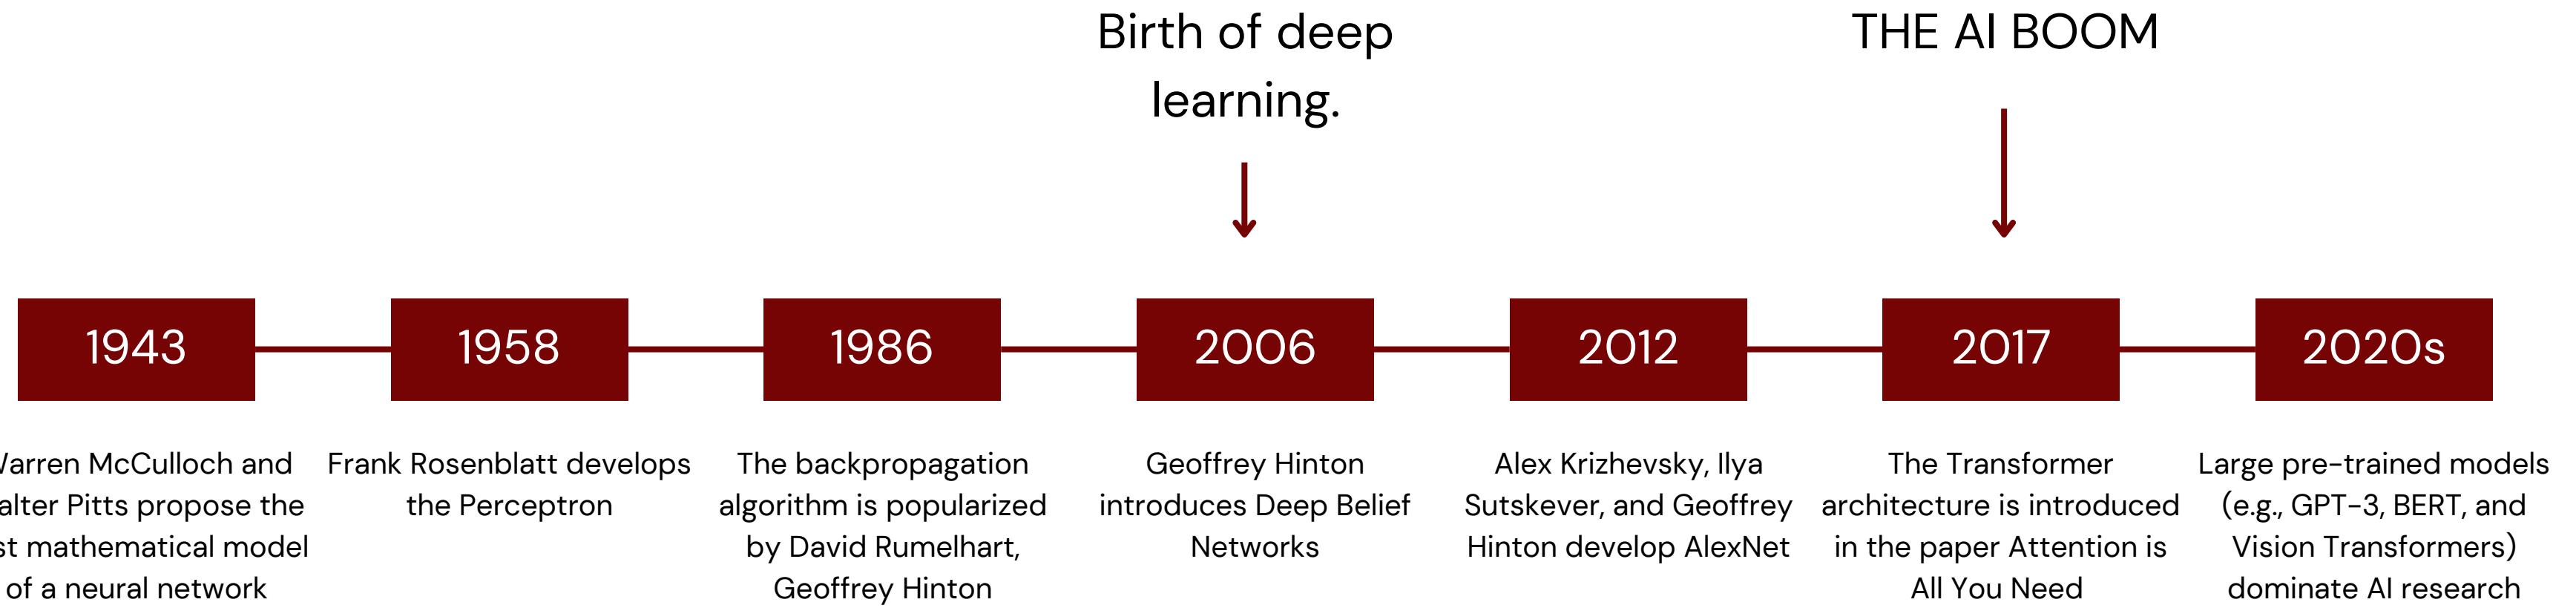


Biological Neurons



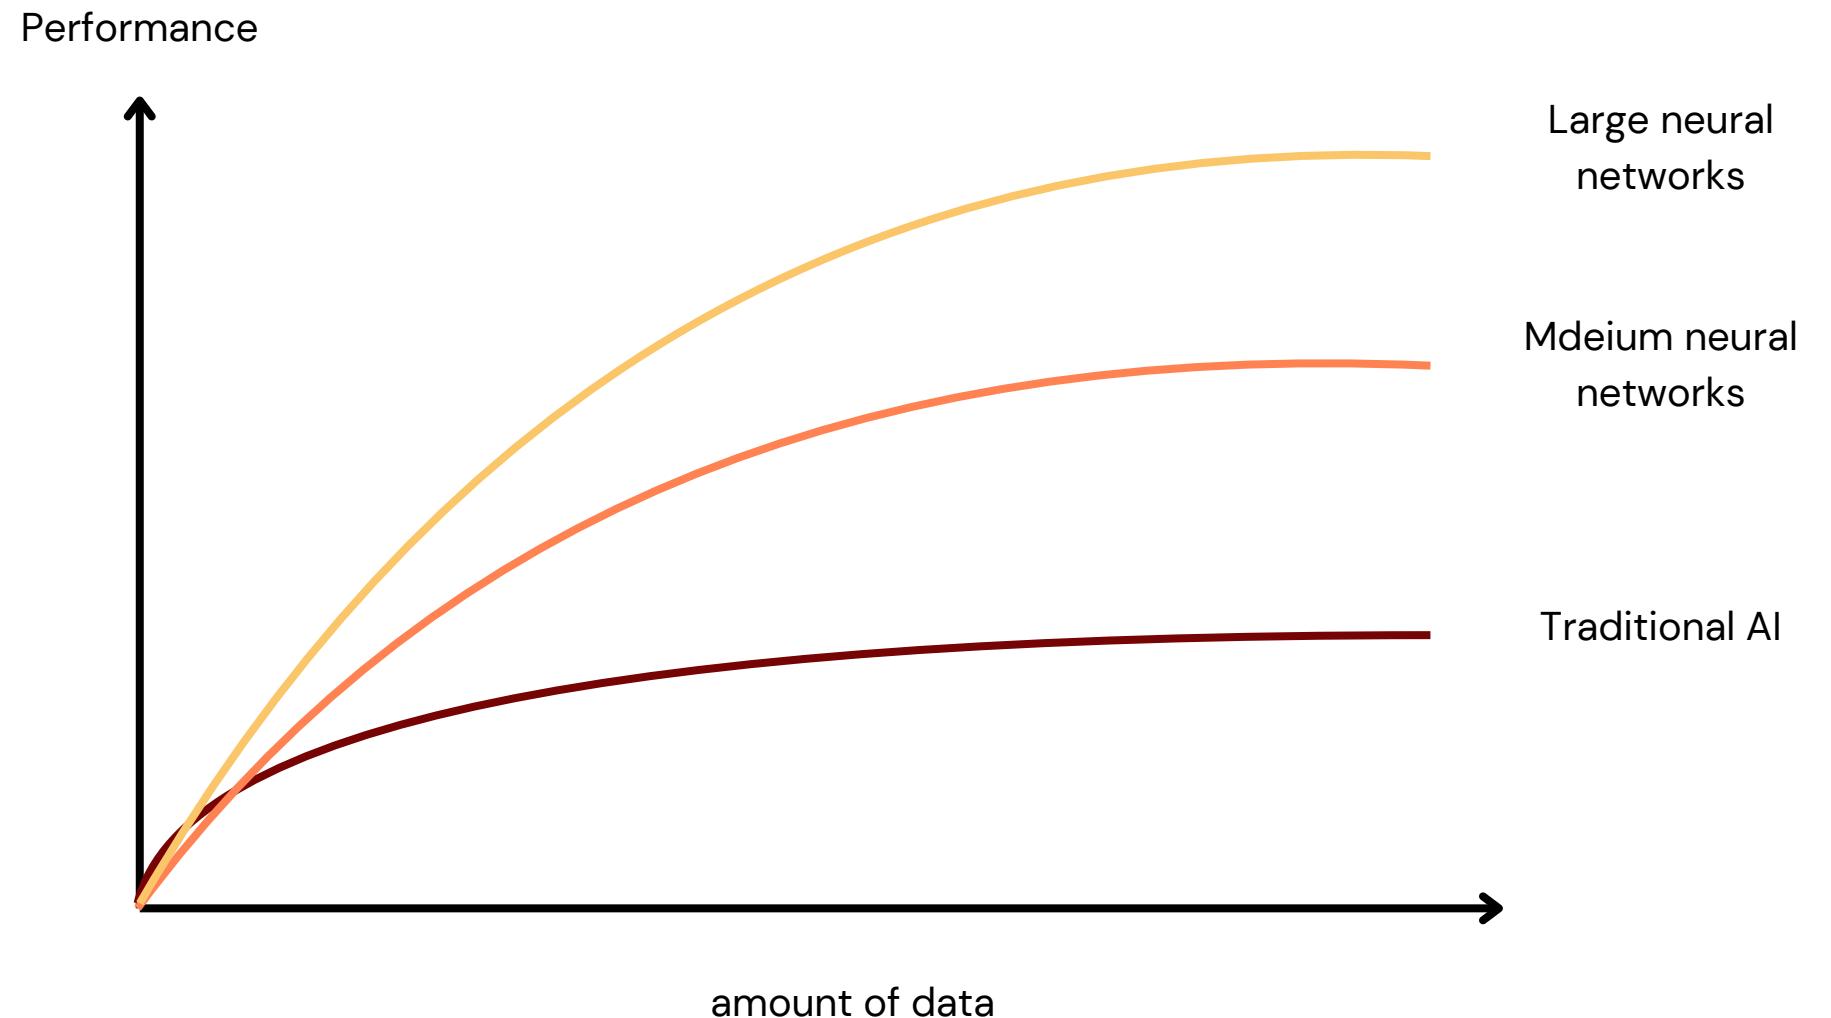
Artificial Neurons

# History of Neural Networks



# Why Are Neural Networks Relevant Today?

- Explosion of data: Neural networks thrive on large datasets.
- Advancements in computing power: GPUs and TPUs.

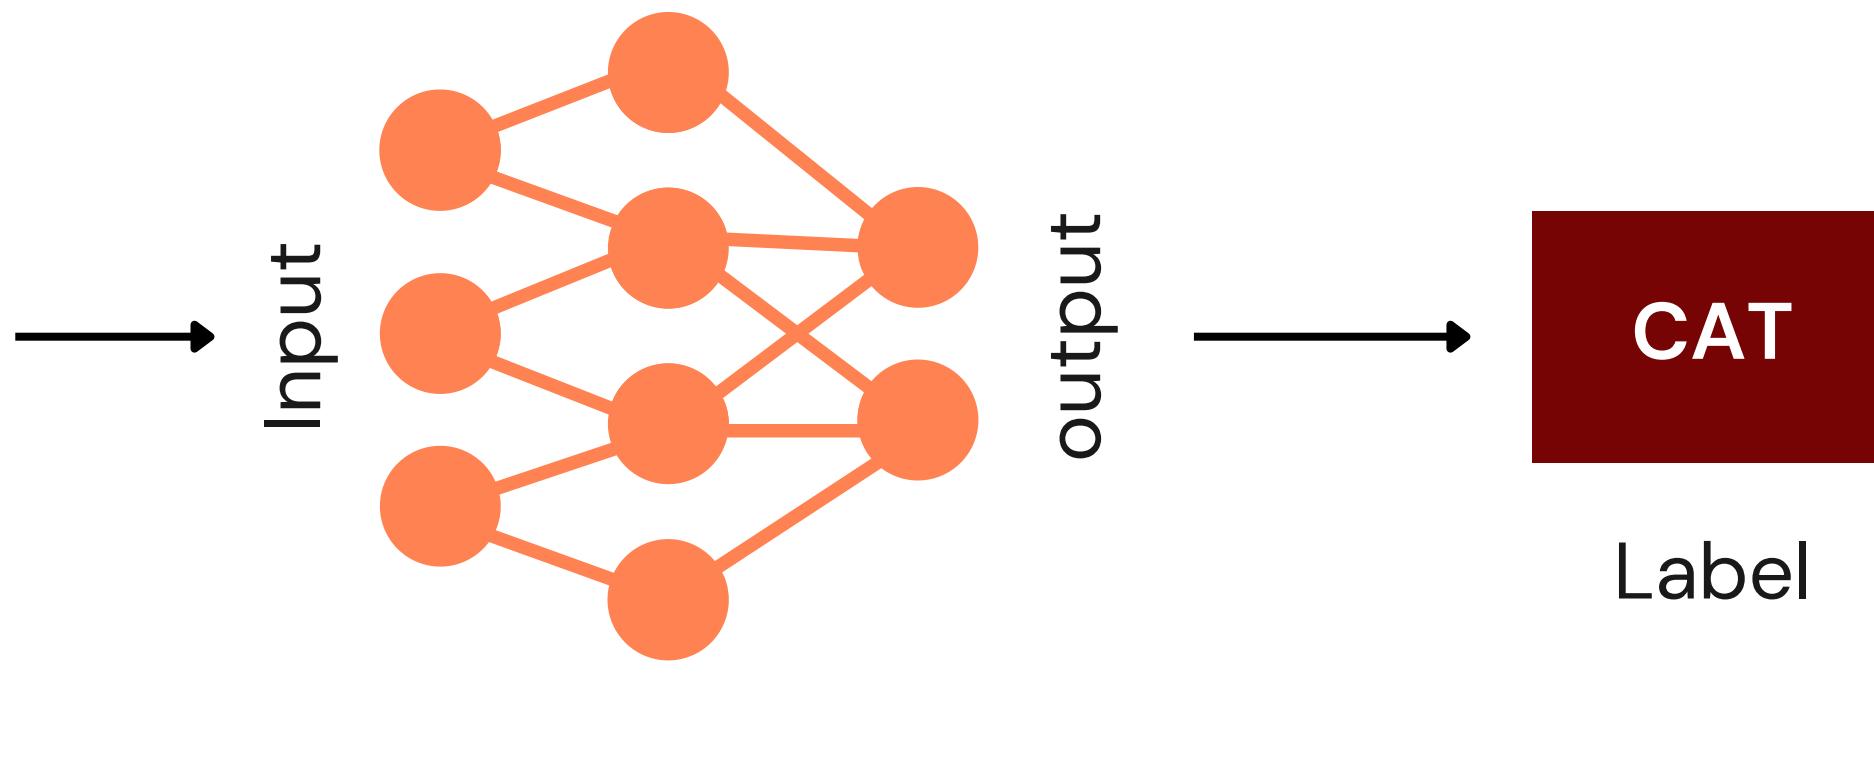


# What Can We Do With Neural Networks?

## Classification

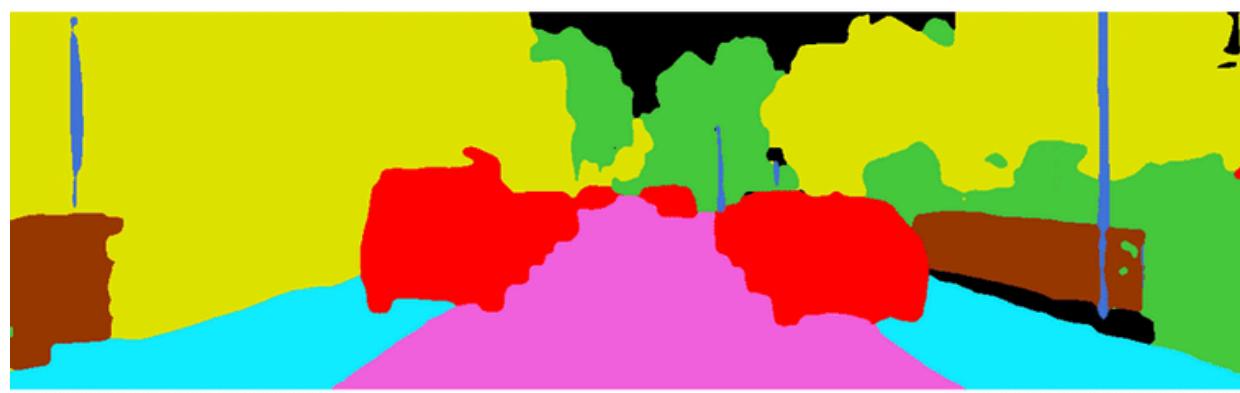


image



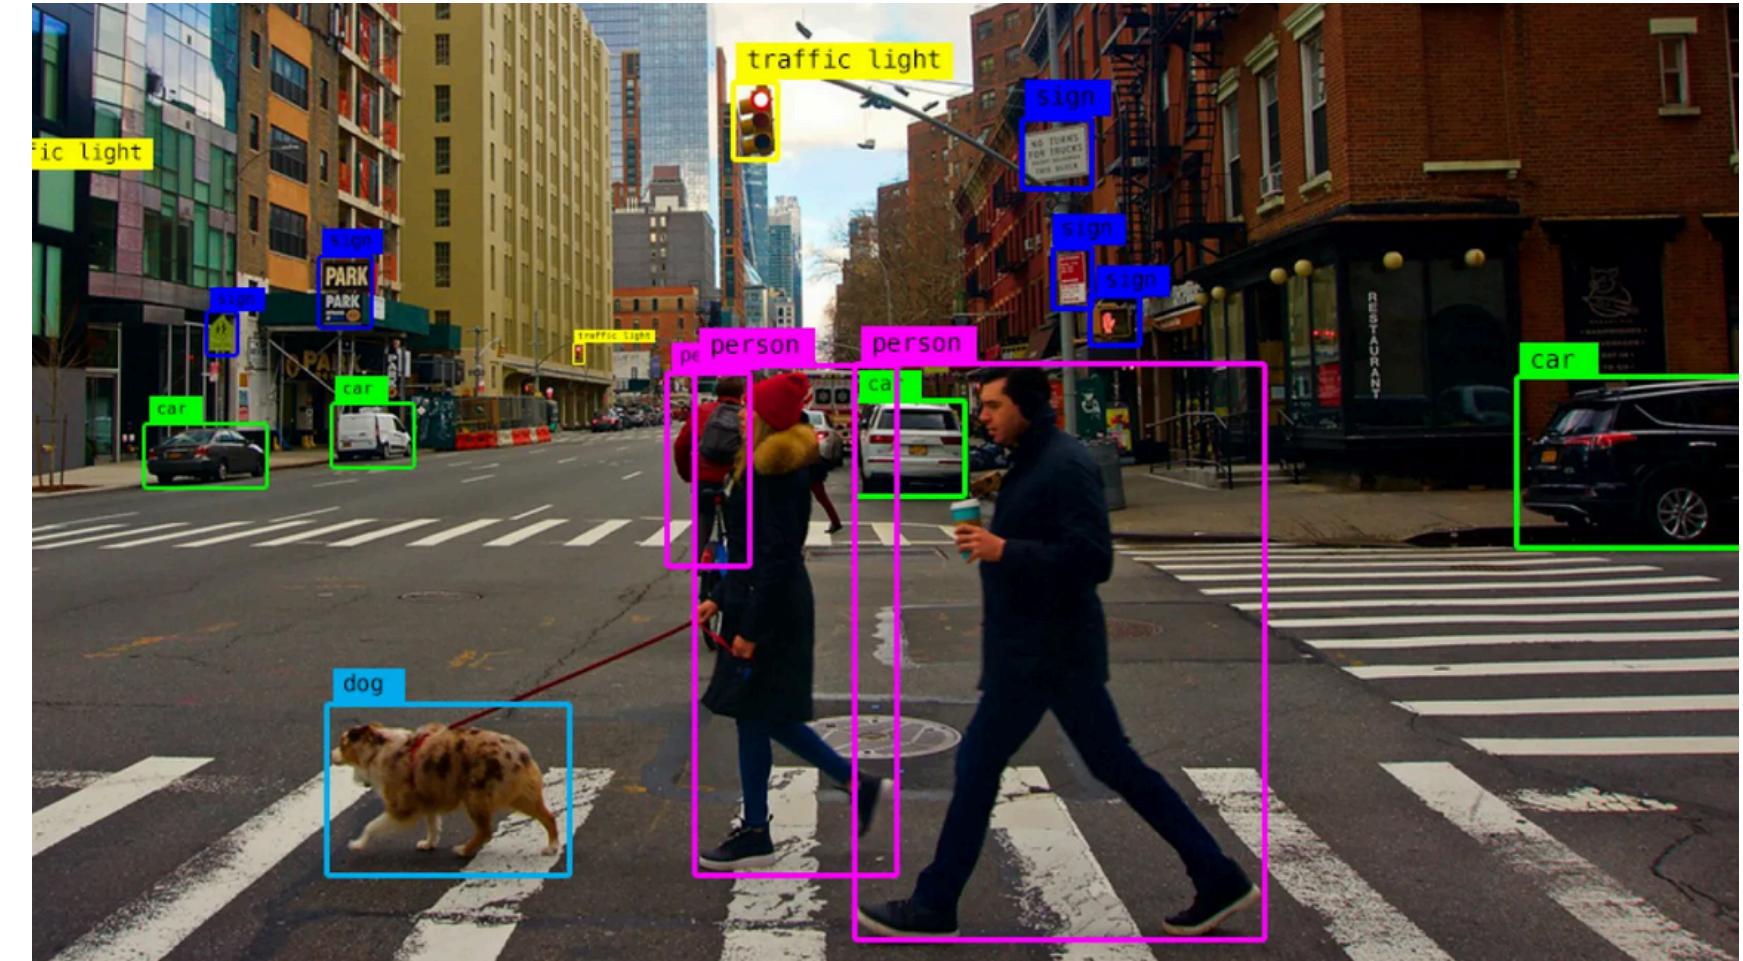
# What Can We Do With Neural Networks?

## Tasks beyond classification



Road	Sidewalk	Building	Fence
Pole	Vegetation	Vehicle	Unlabel

Semantic segmentation



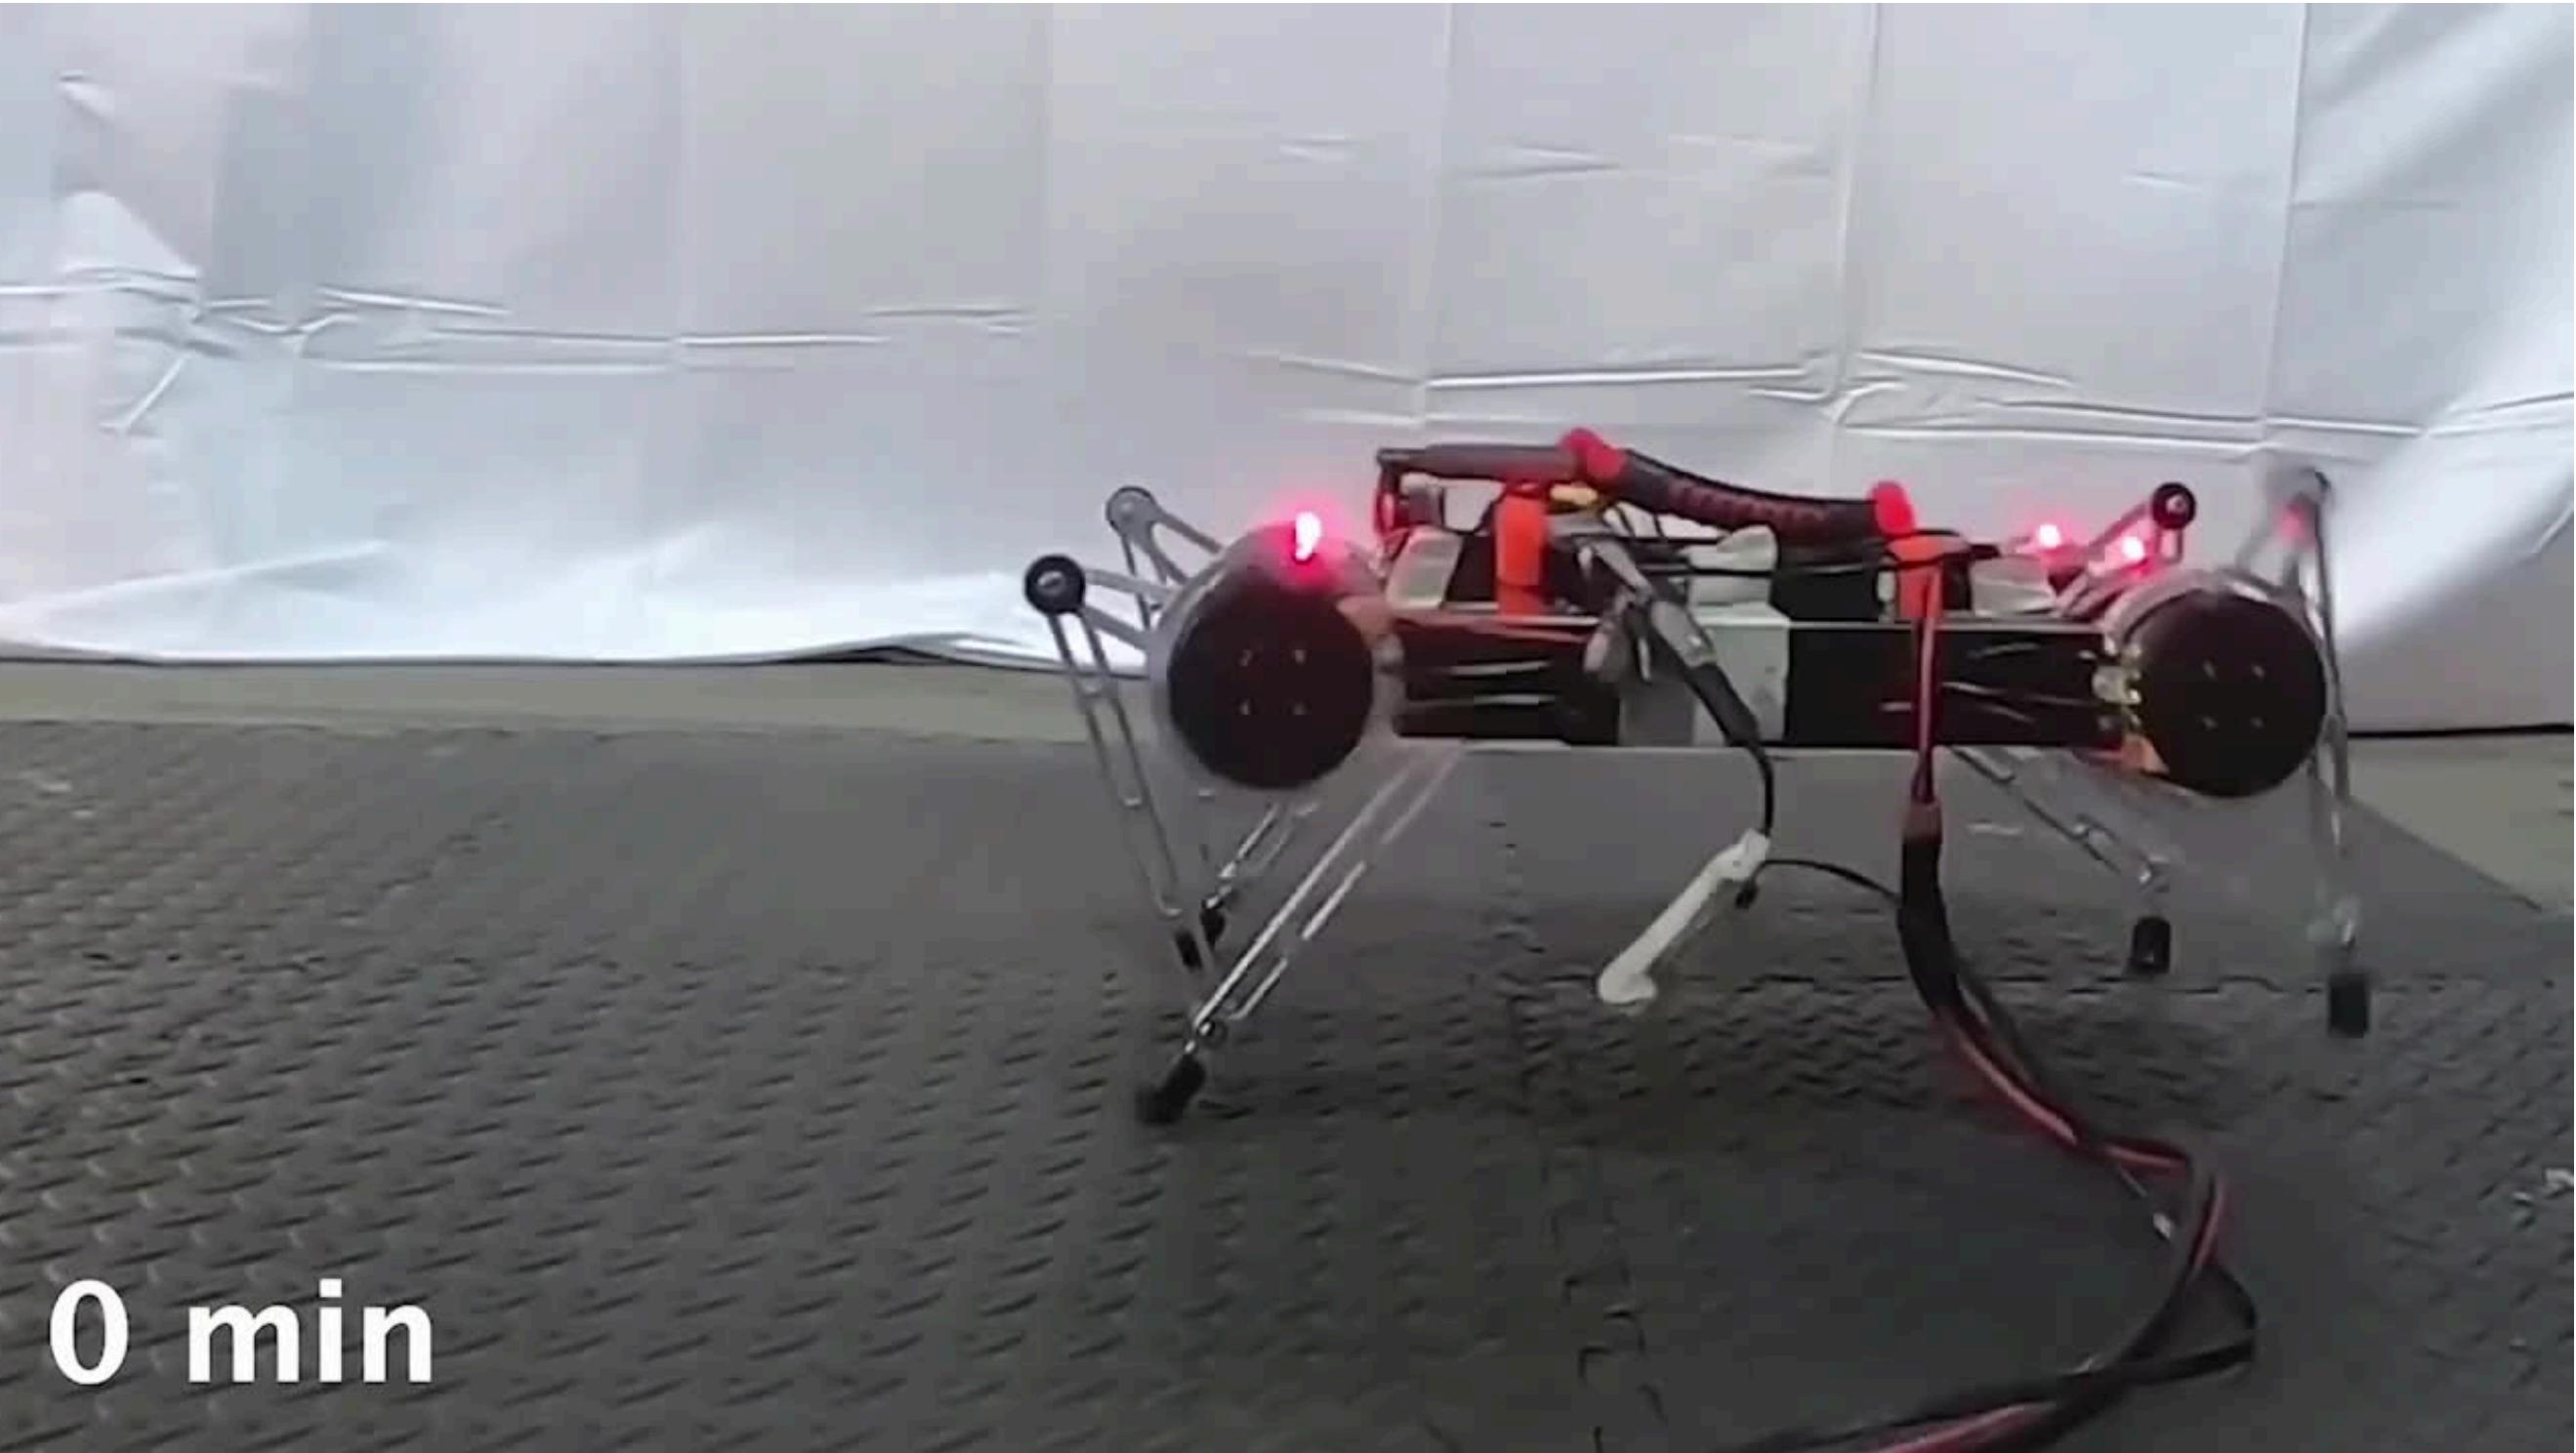
Object detection

# What Can We Do With Neural Networks?



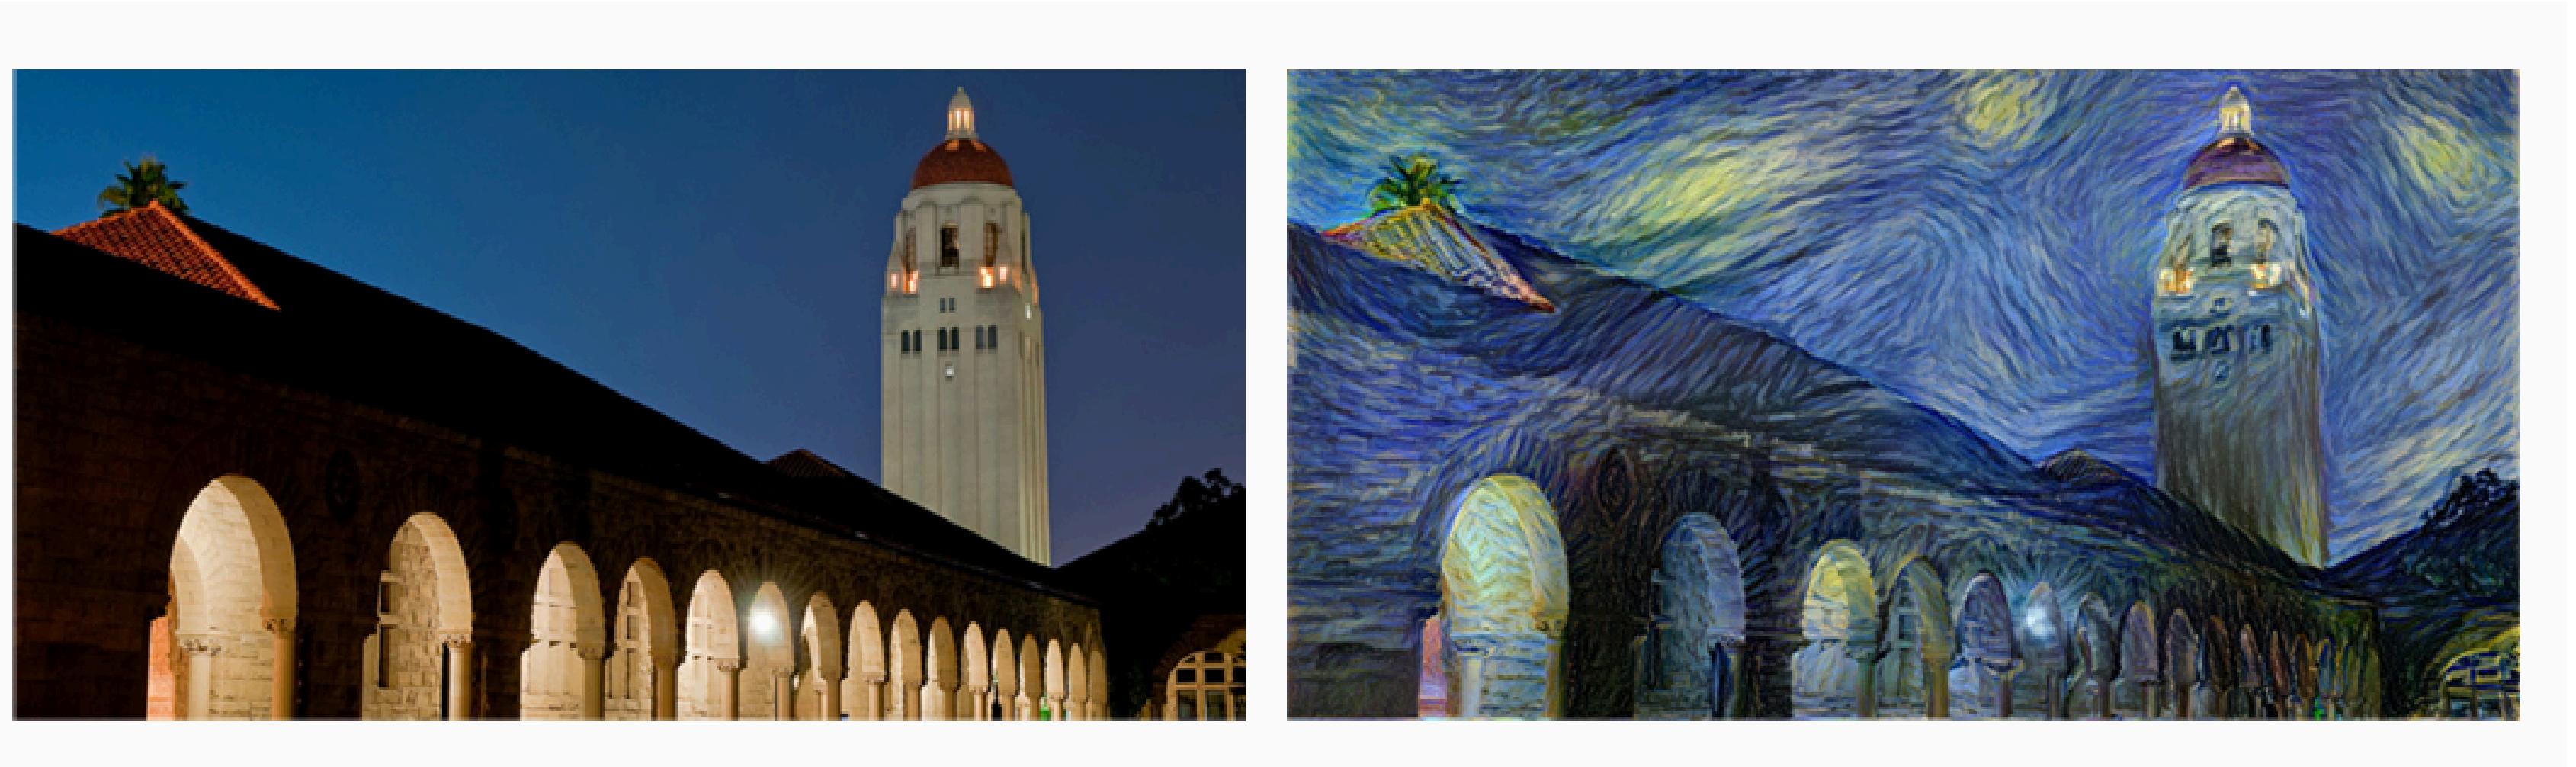
Video Understanding

# What Can We Do With Neural Networks?



Controlling physical systems

# What Can We Do With Neural Networks?



Style transfe

# What Can We Do With Neural Networks?

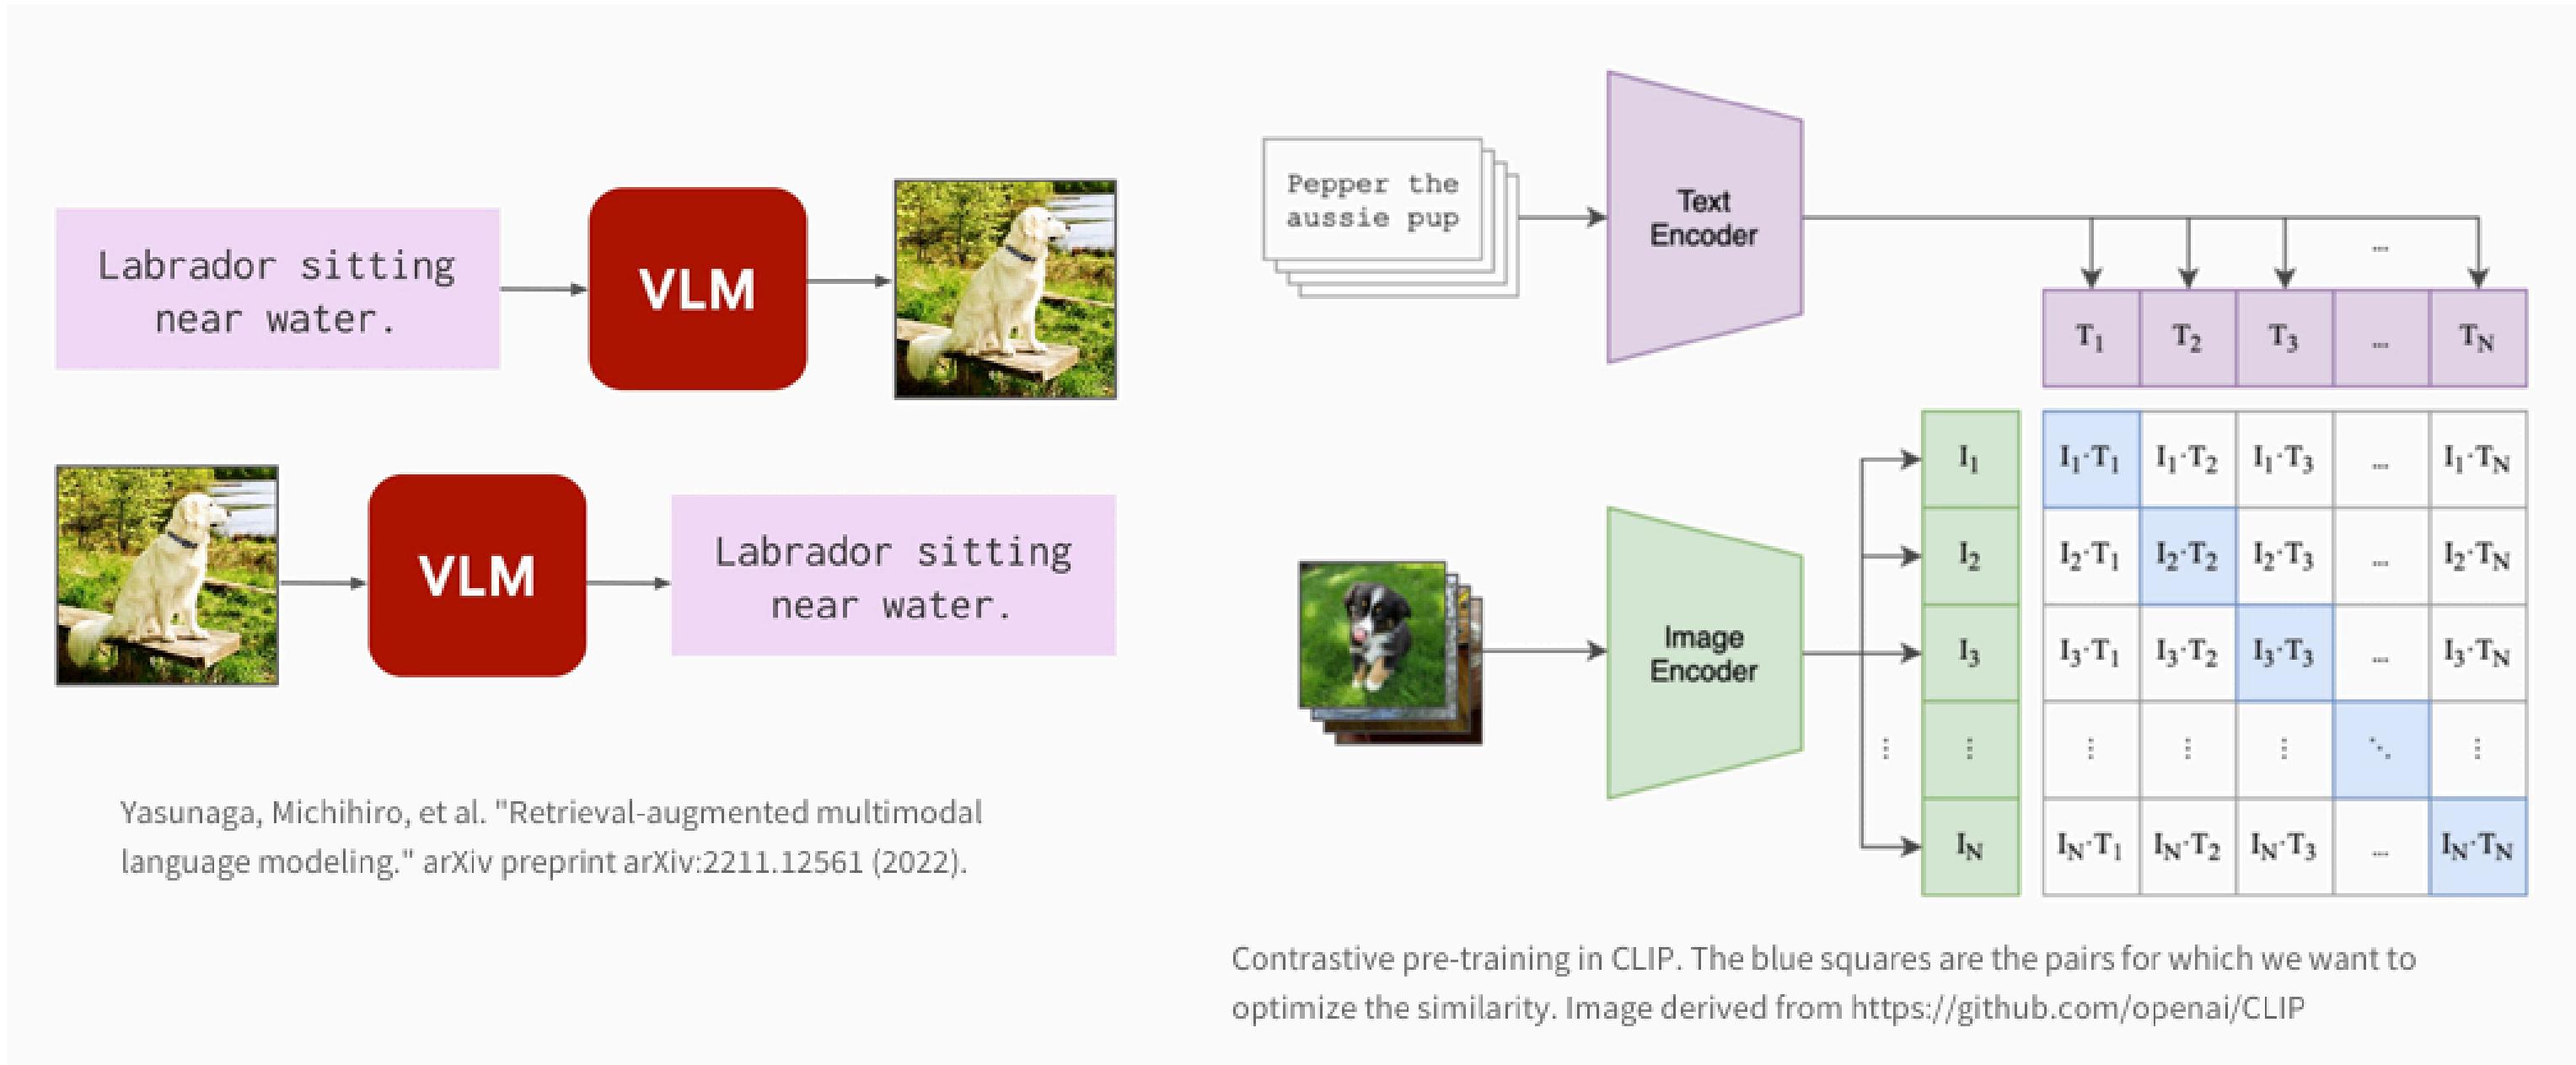


“Teddy bears working on new AI research underwater with 1990s technology”

DALL-E 2

Generative modeling

# What Can We Do With Neural Networks?

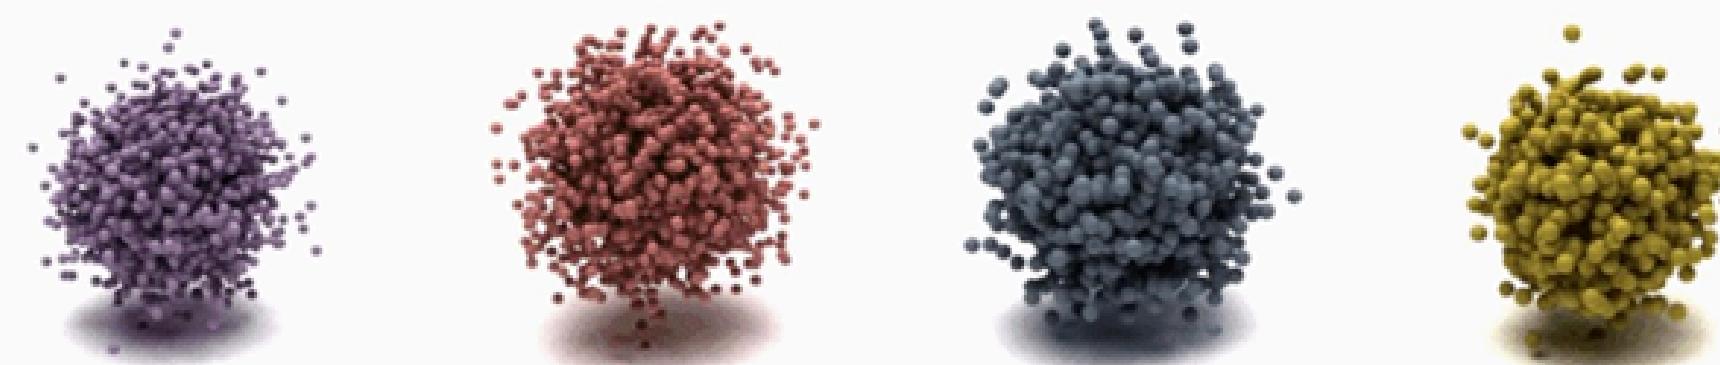


Vision Language Modeling

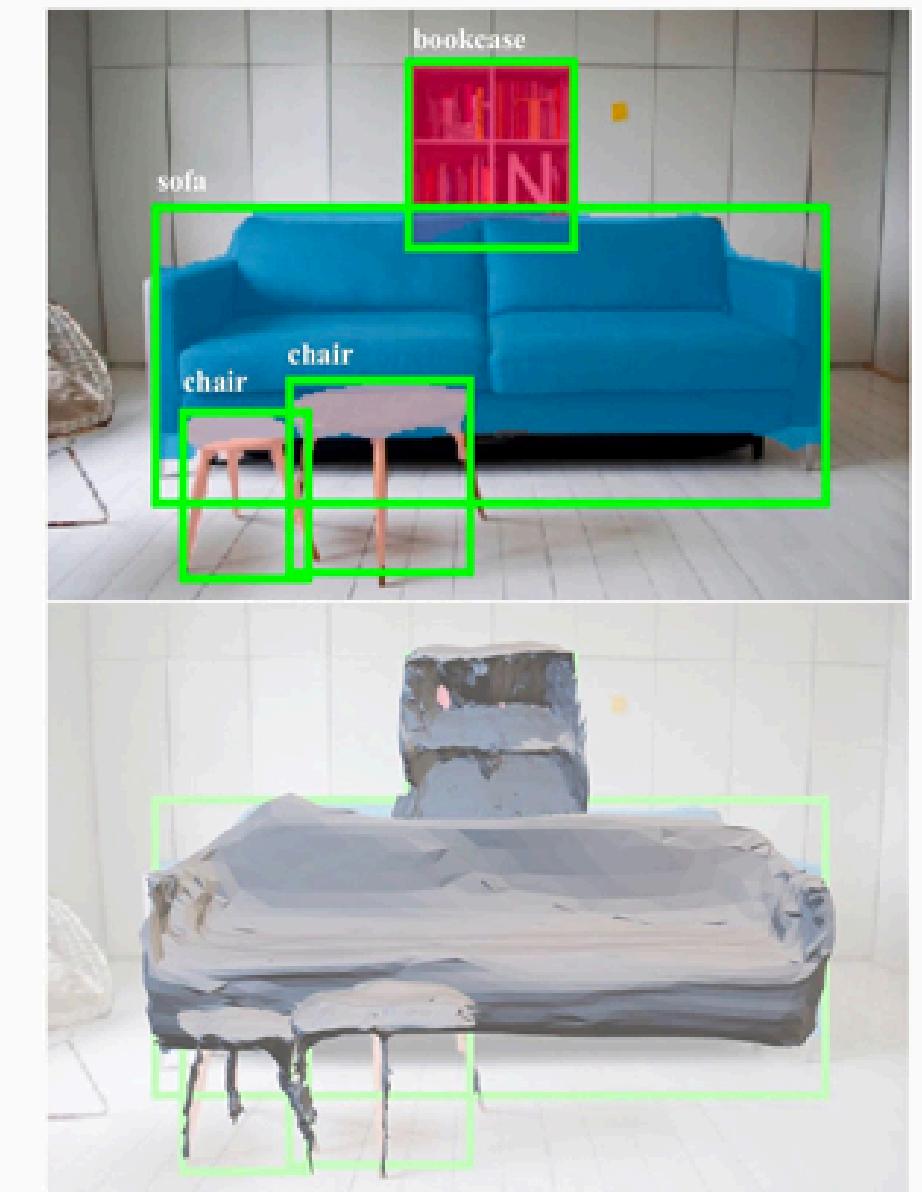
# What Can We Do With Neural Networks?



Choy et al., 3D-R2N2: Recurrent Reconstruction Neural Network (2016)



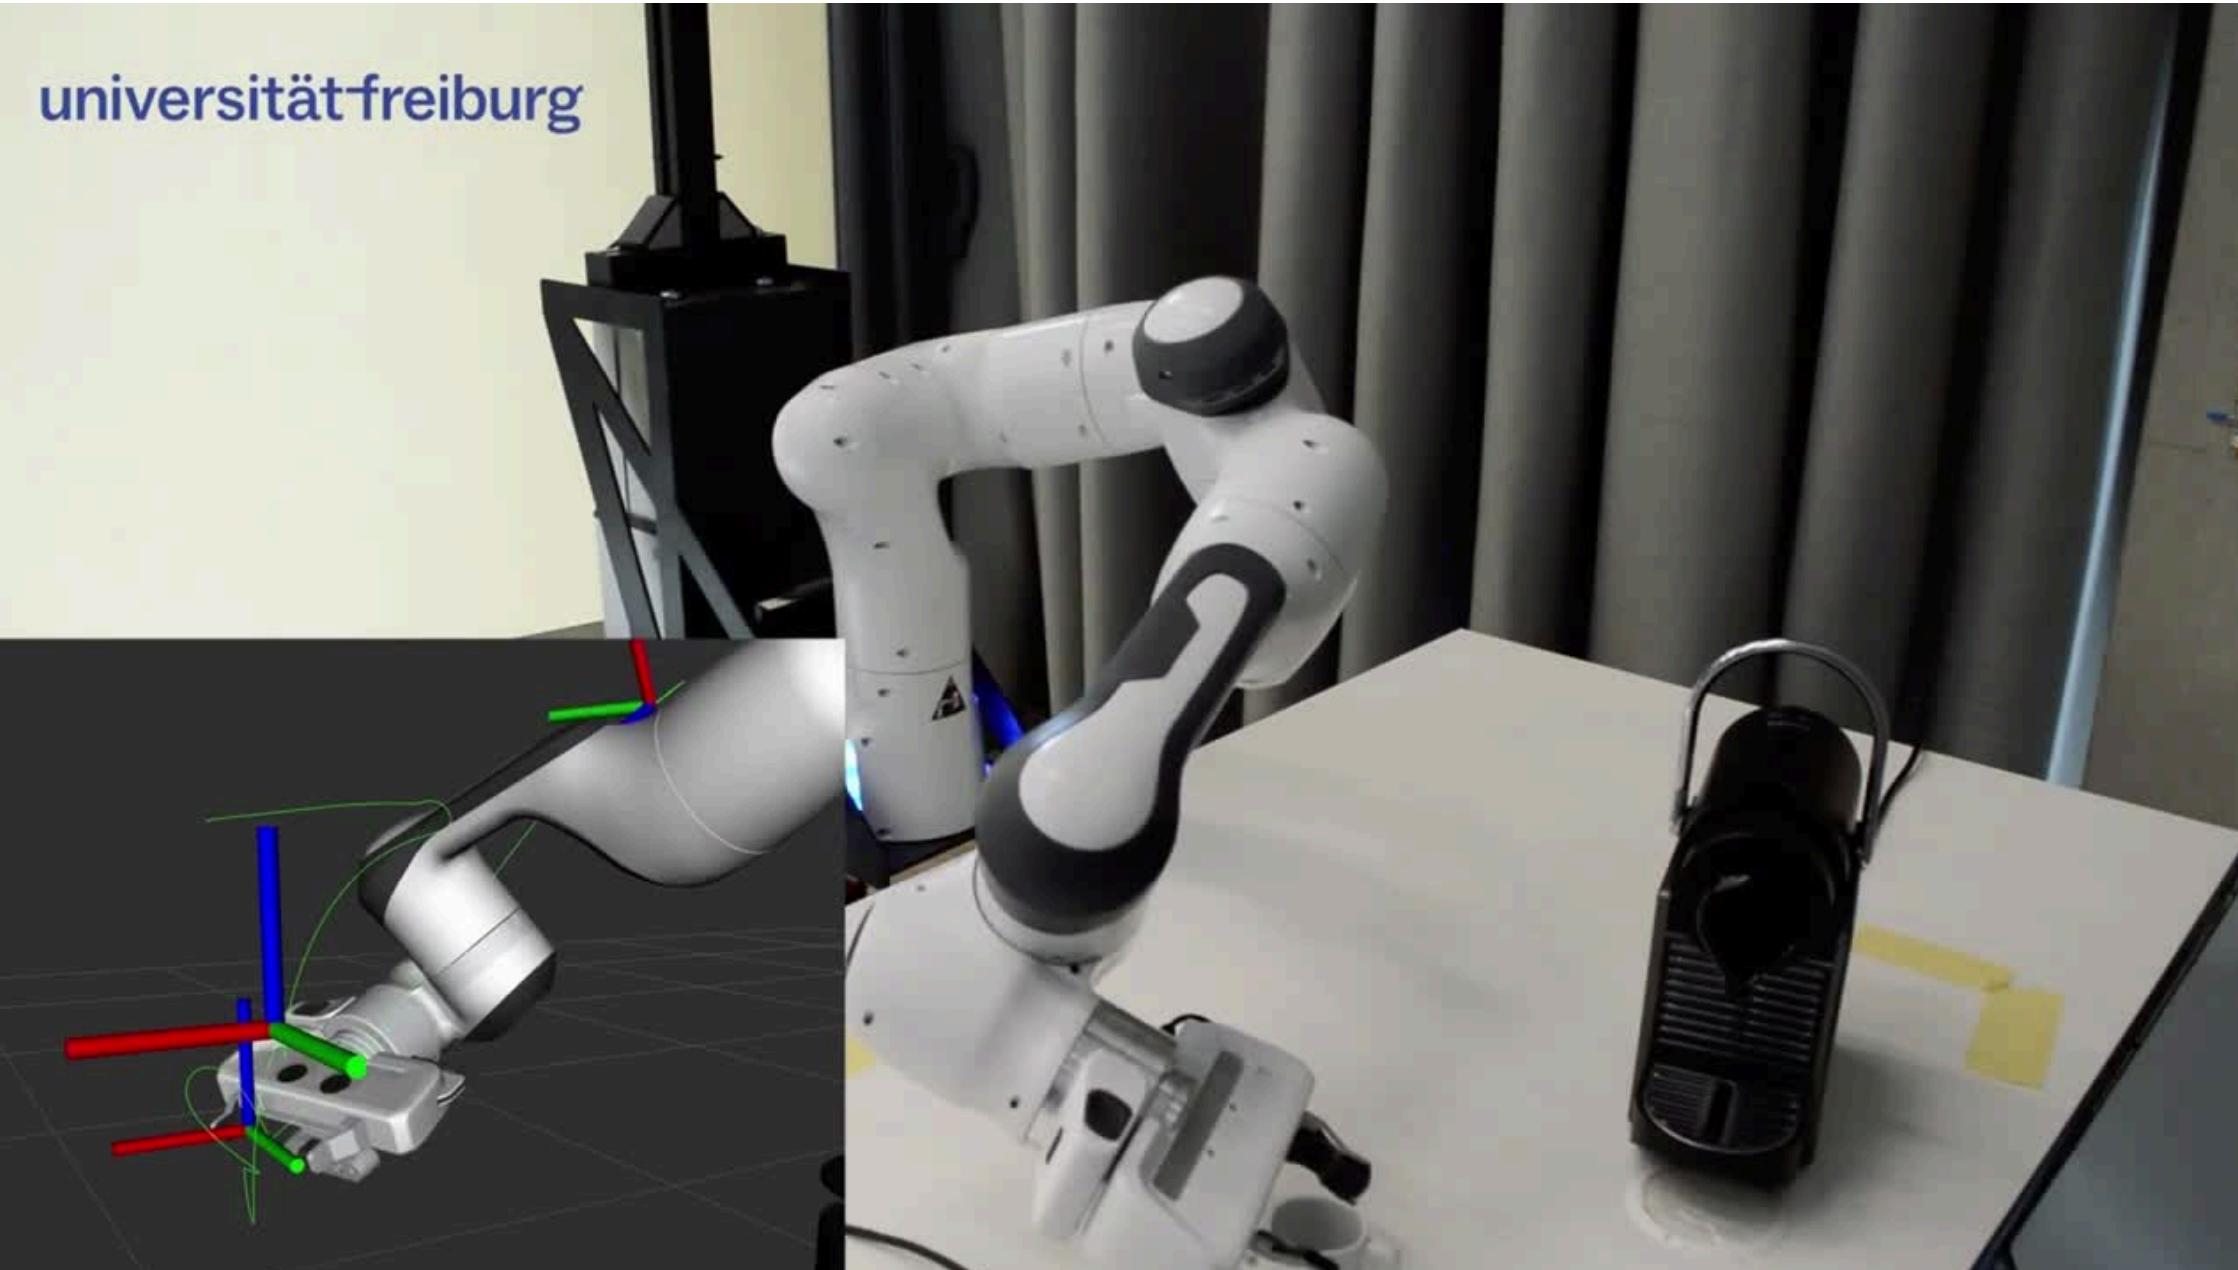
Zhou et al., 3D Shape Generation and Completion through Point-Voxel Diffusion (2021)



Gkioxari et al., "Mesh R-CNN", ICCV 2019

3D Vision

# What Can We Do With Neural Networks?

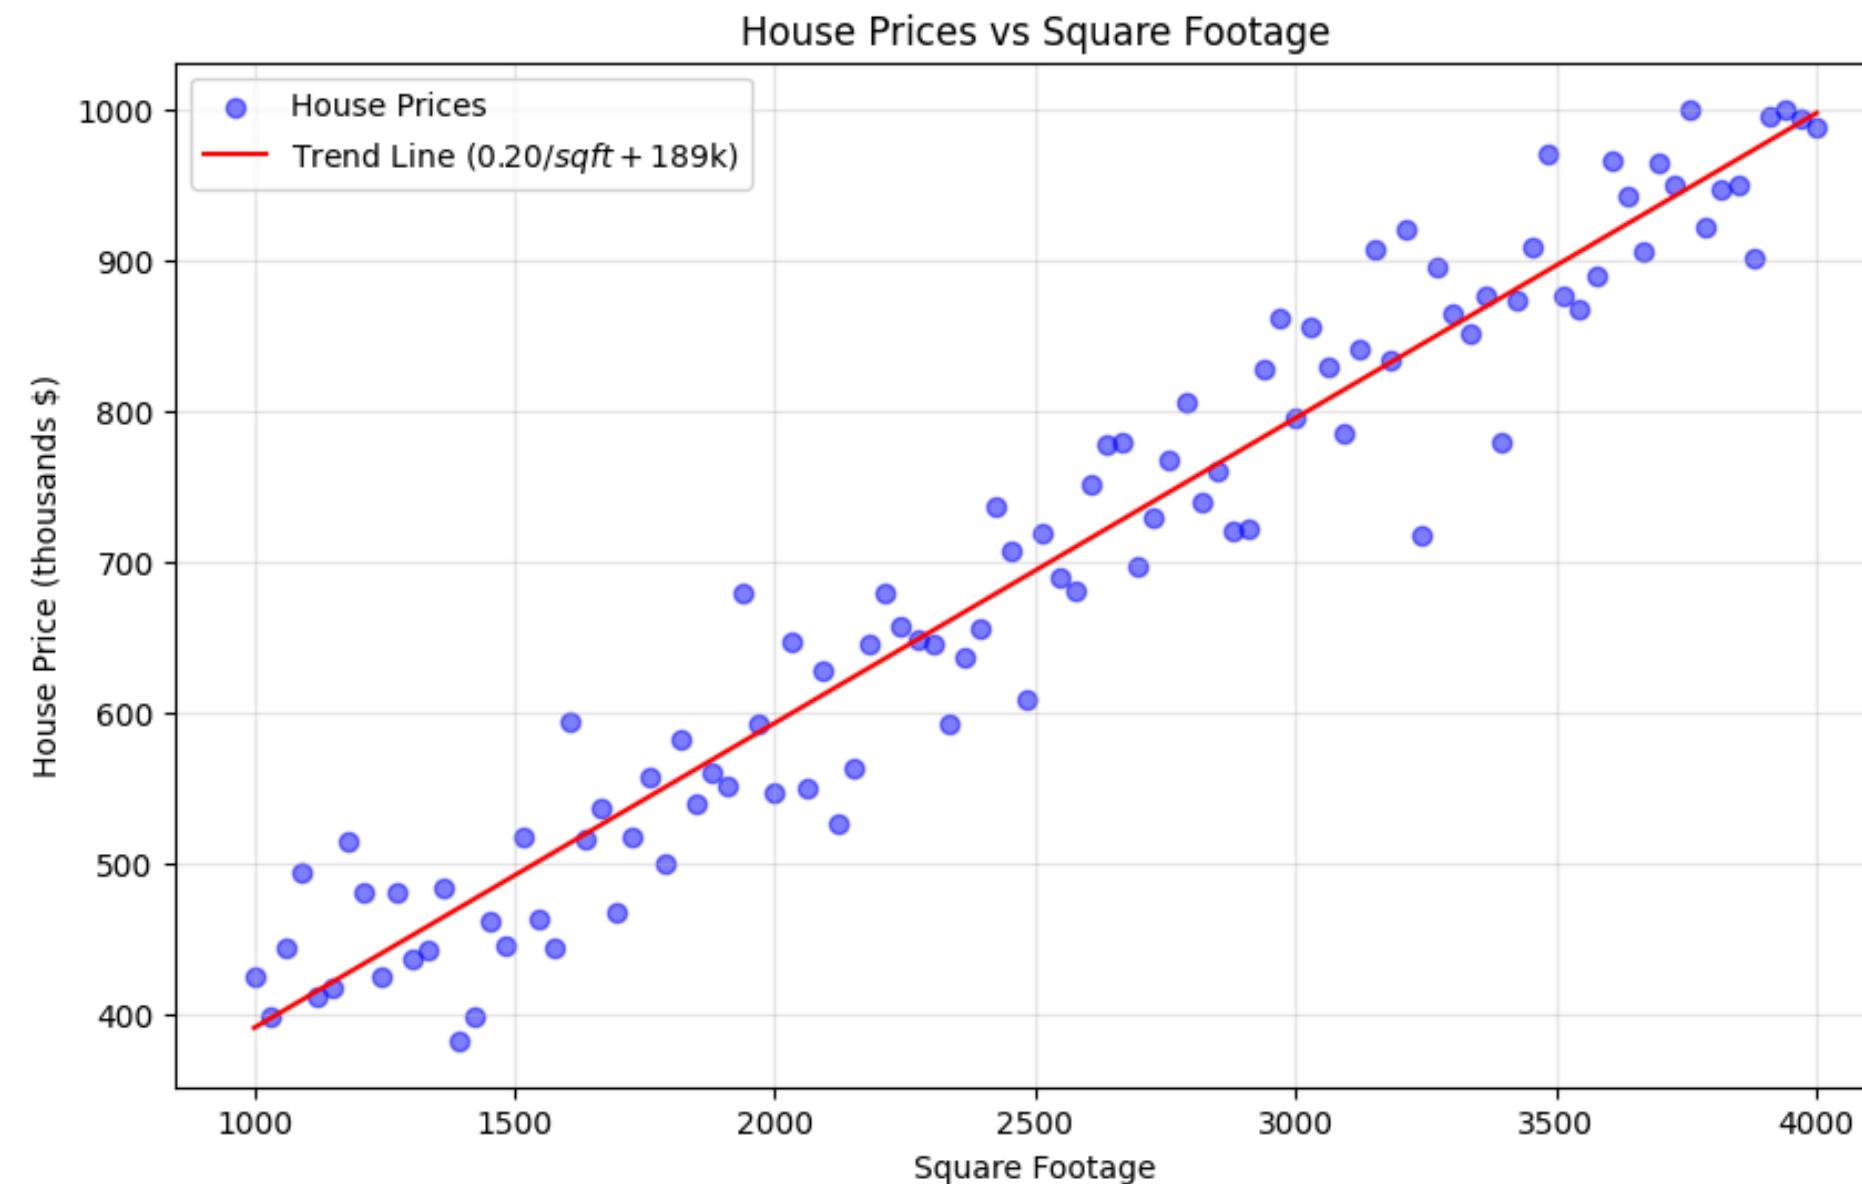


Multi-modal intelligence

# Why Neural Networks ?

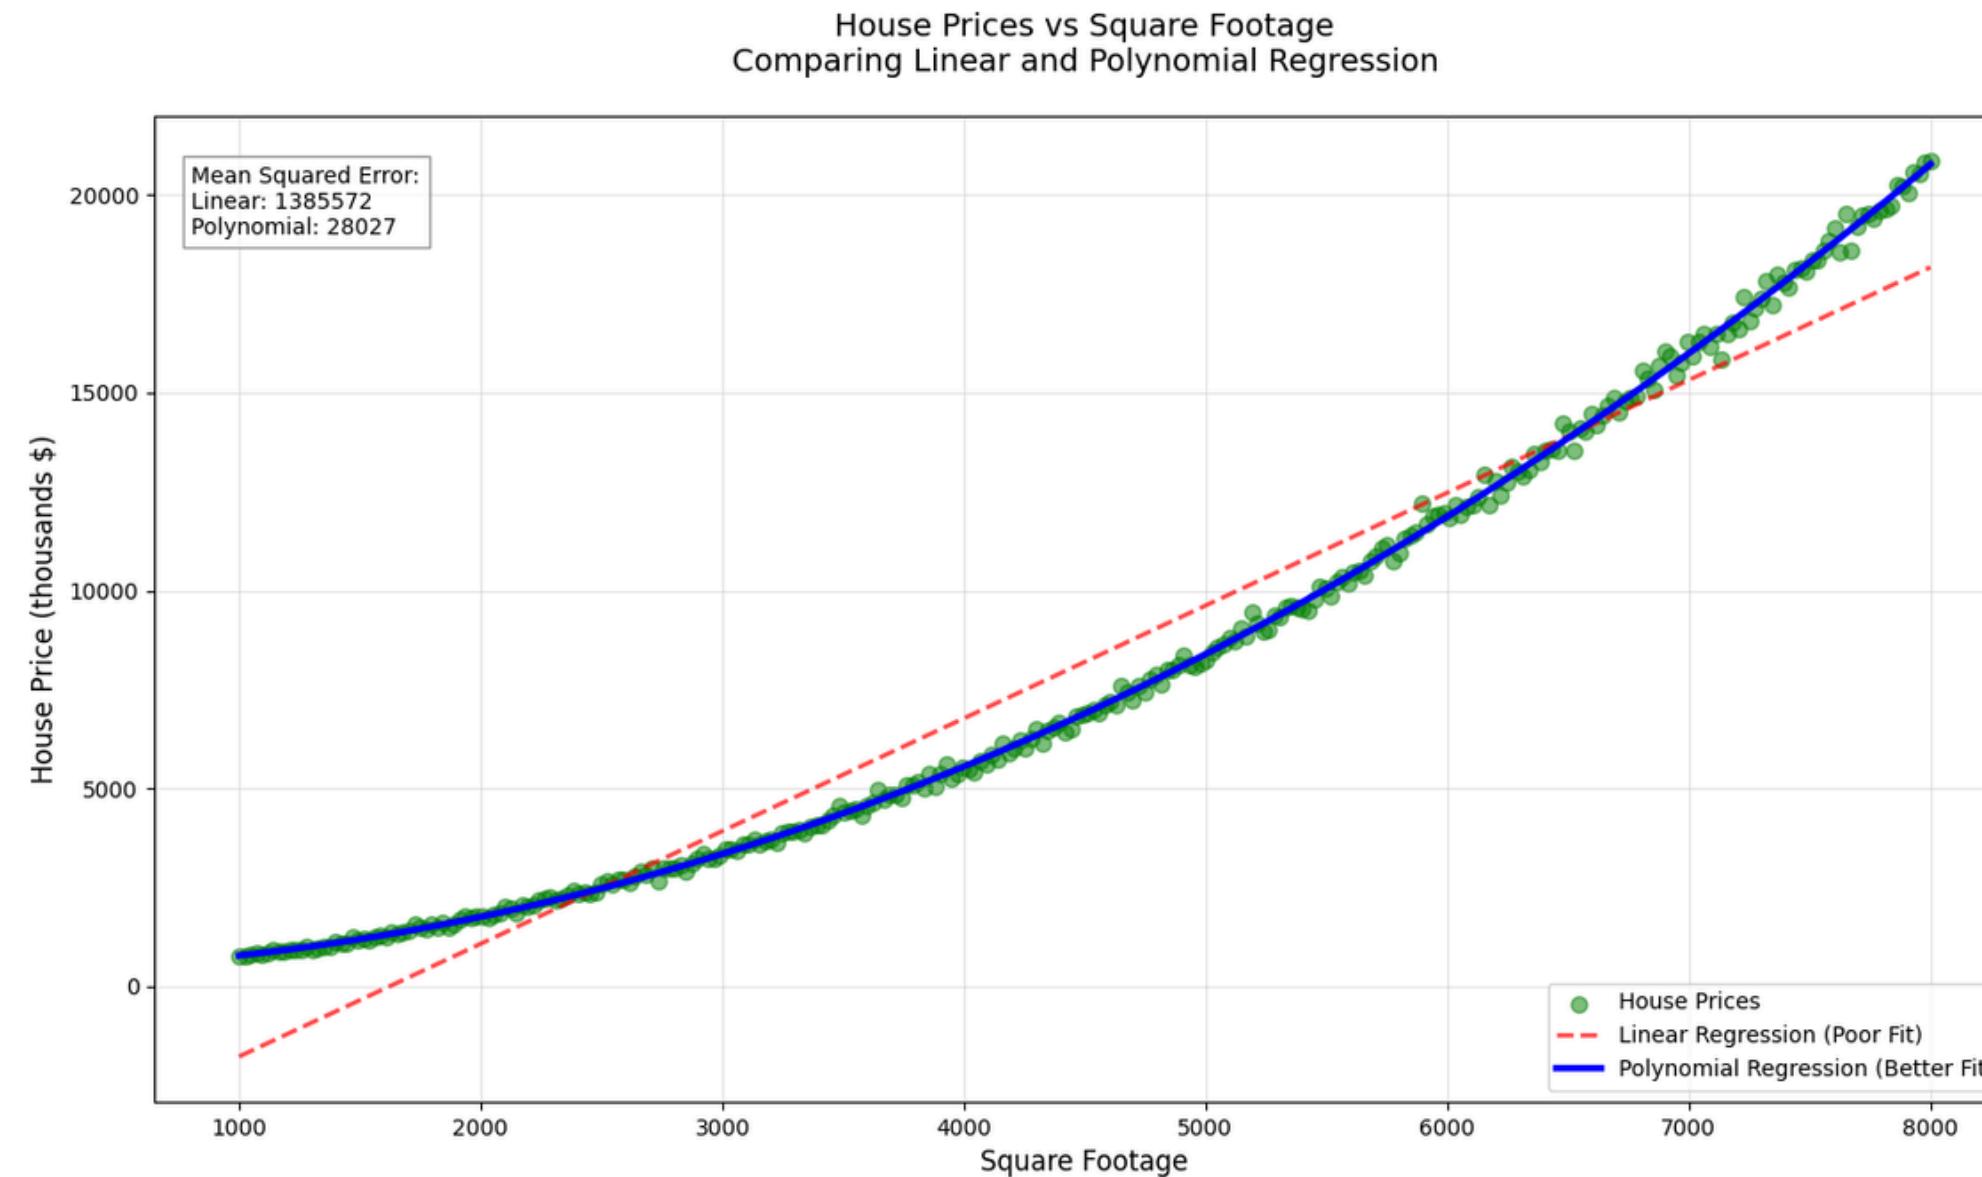
# Linear regression

- Fits a straight line to data.
- It works well when the relationship between input and output is linear.



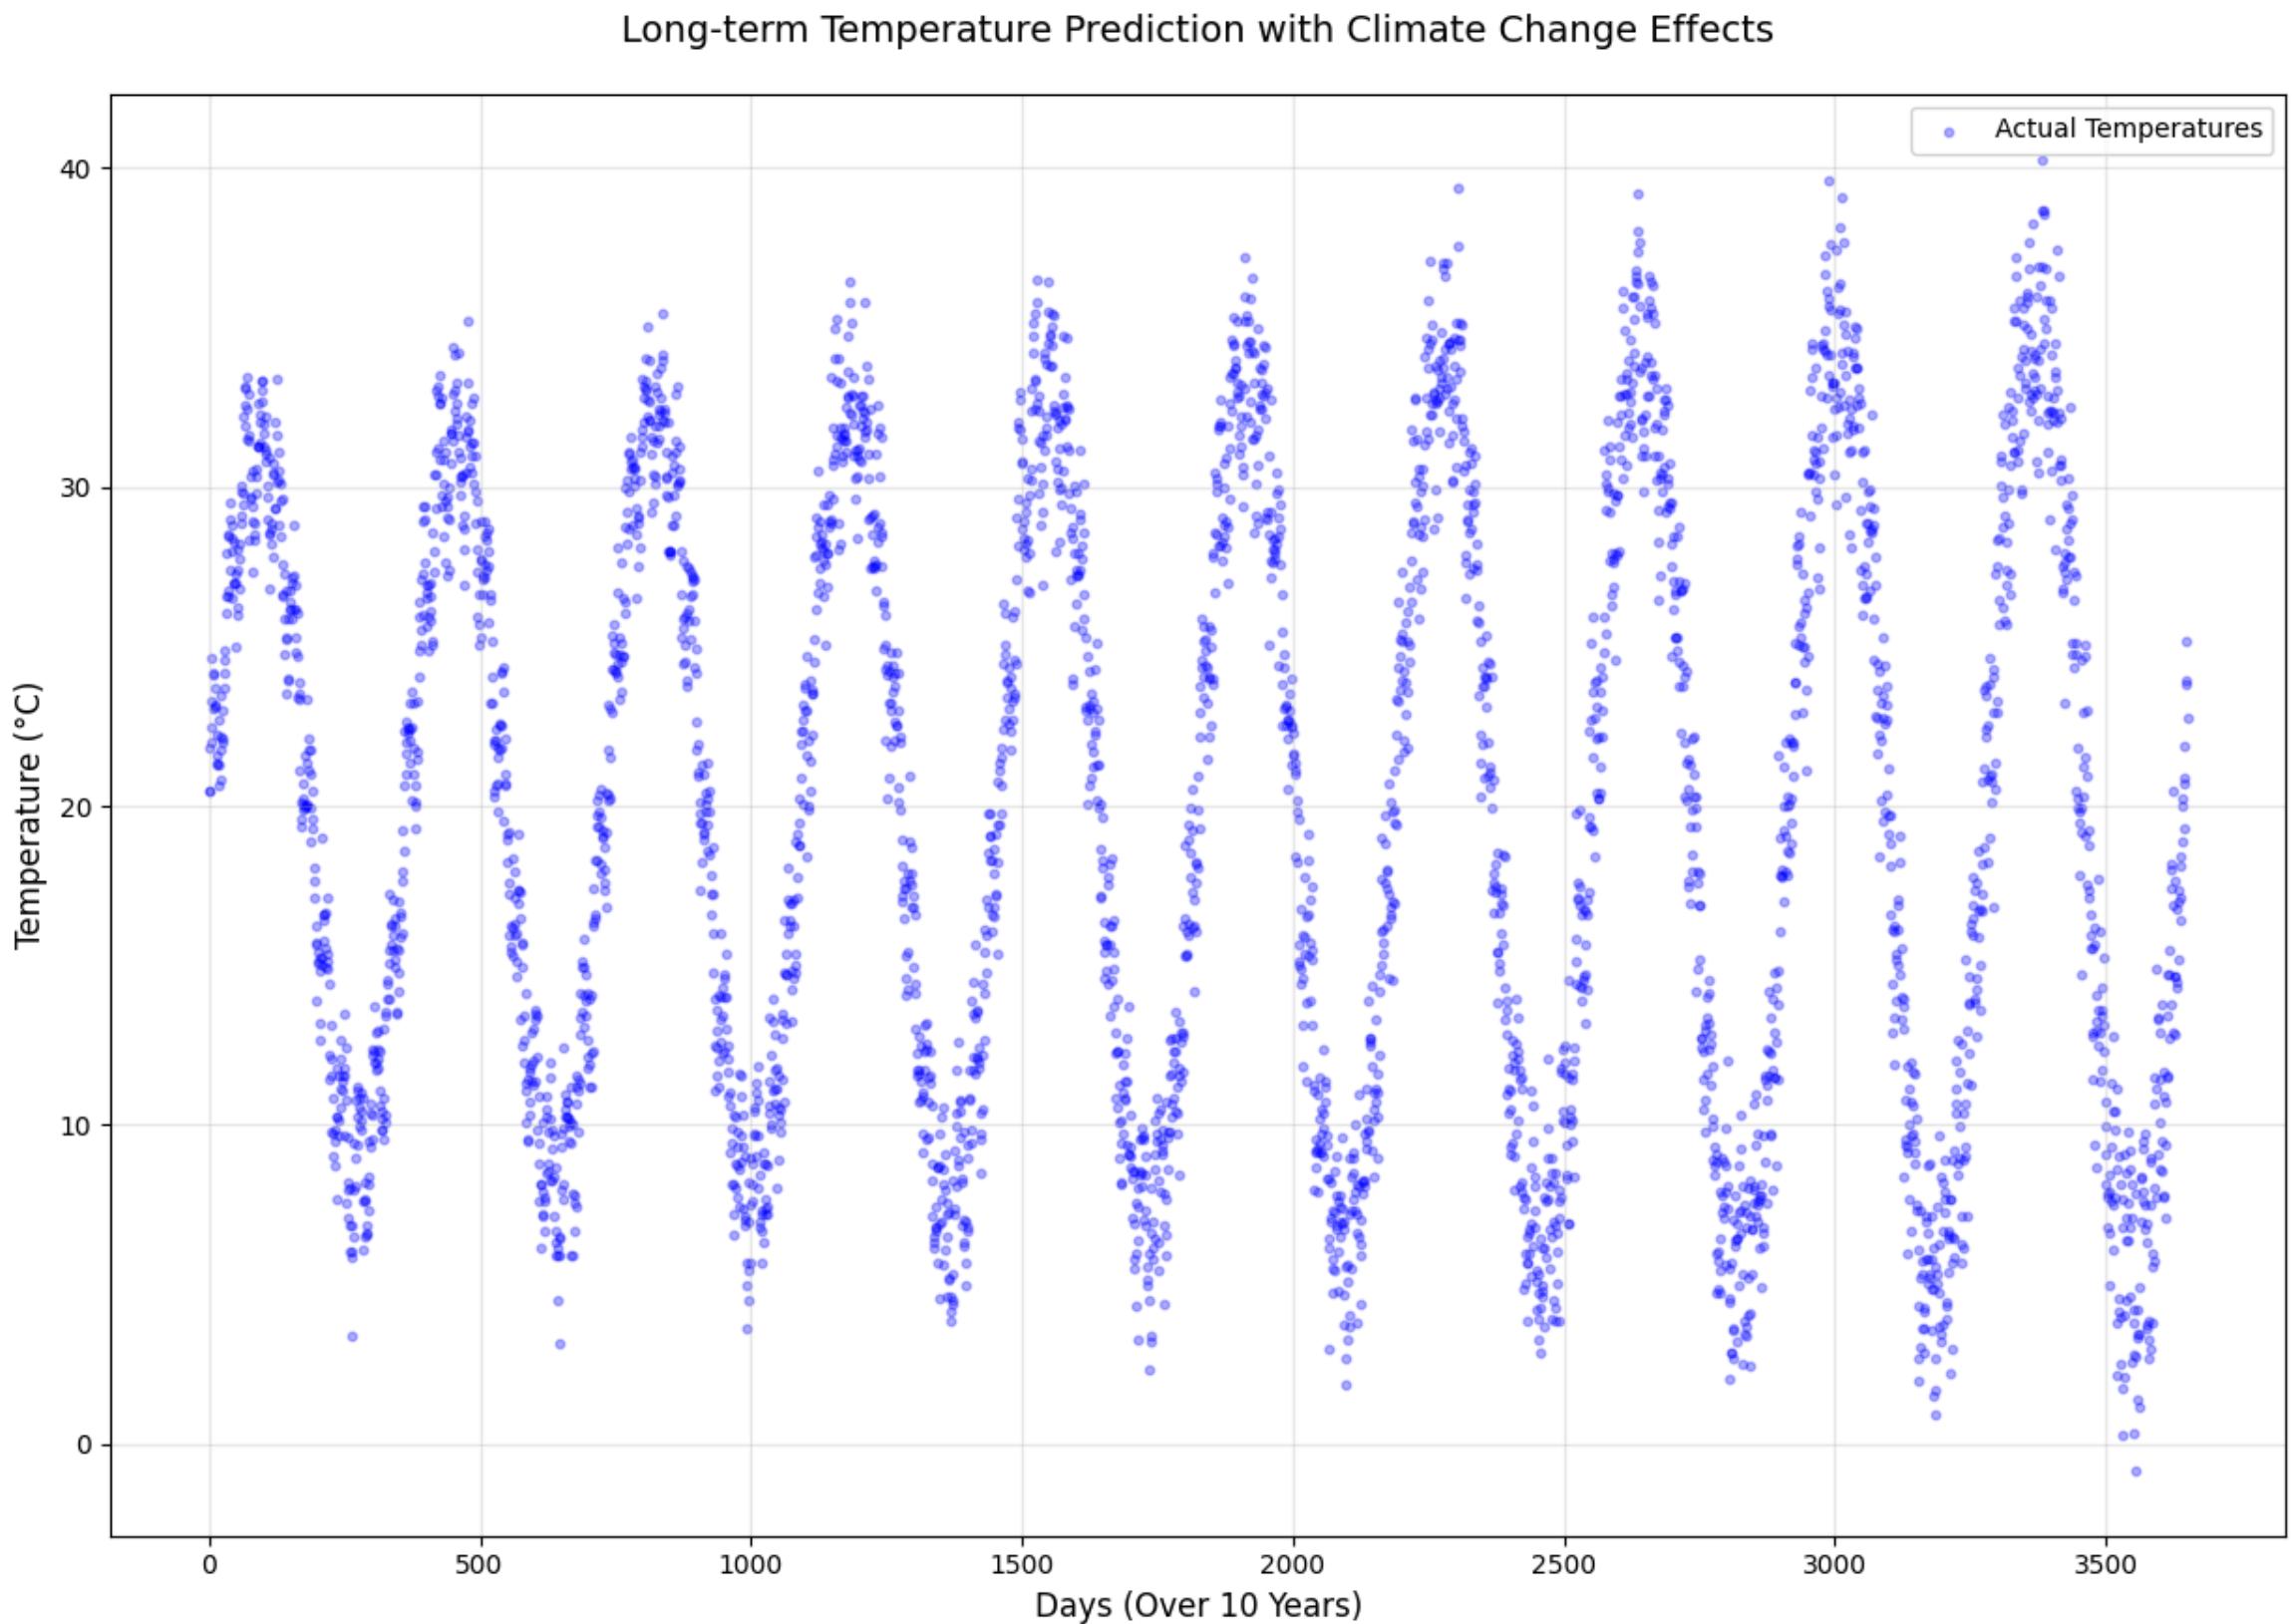
# Polynomial Regression

- Fits data with a polynomial equation with degree greater than 1
- Works on more complex data.



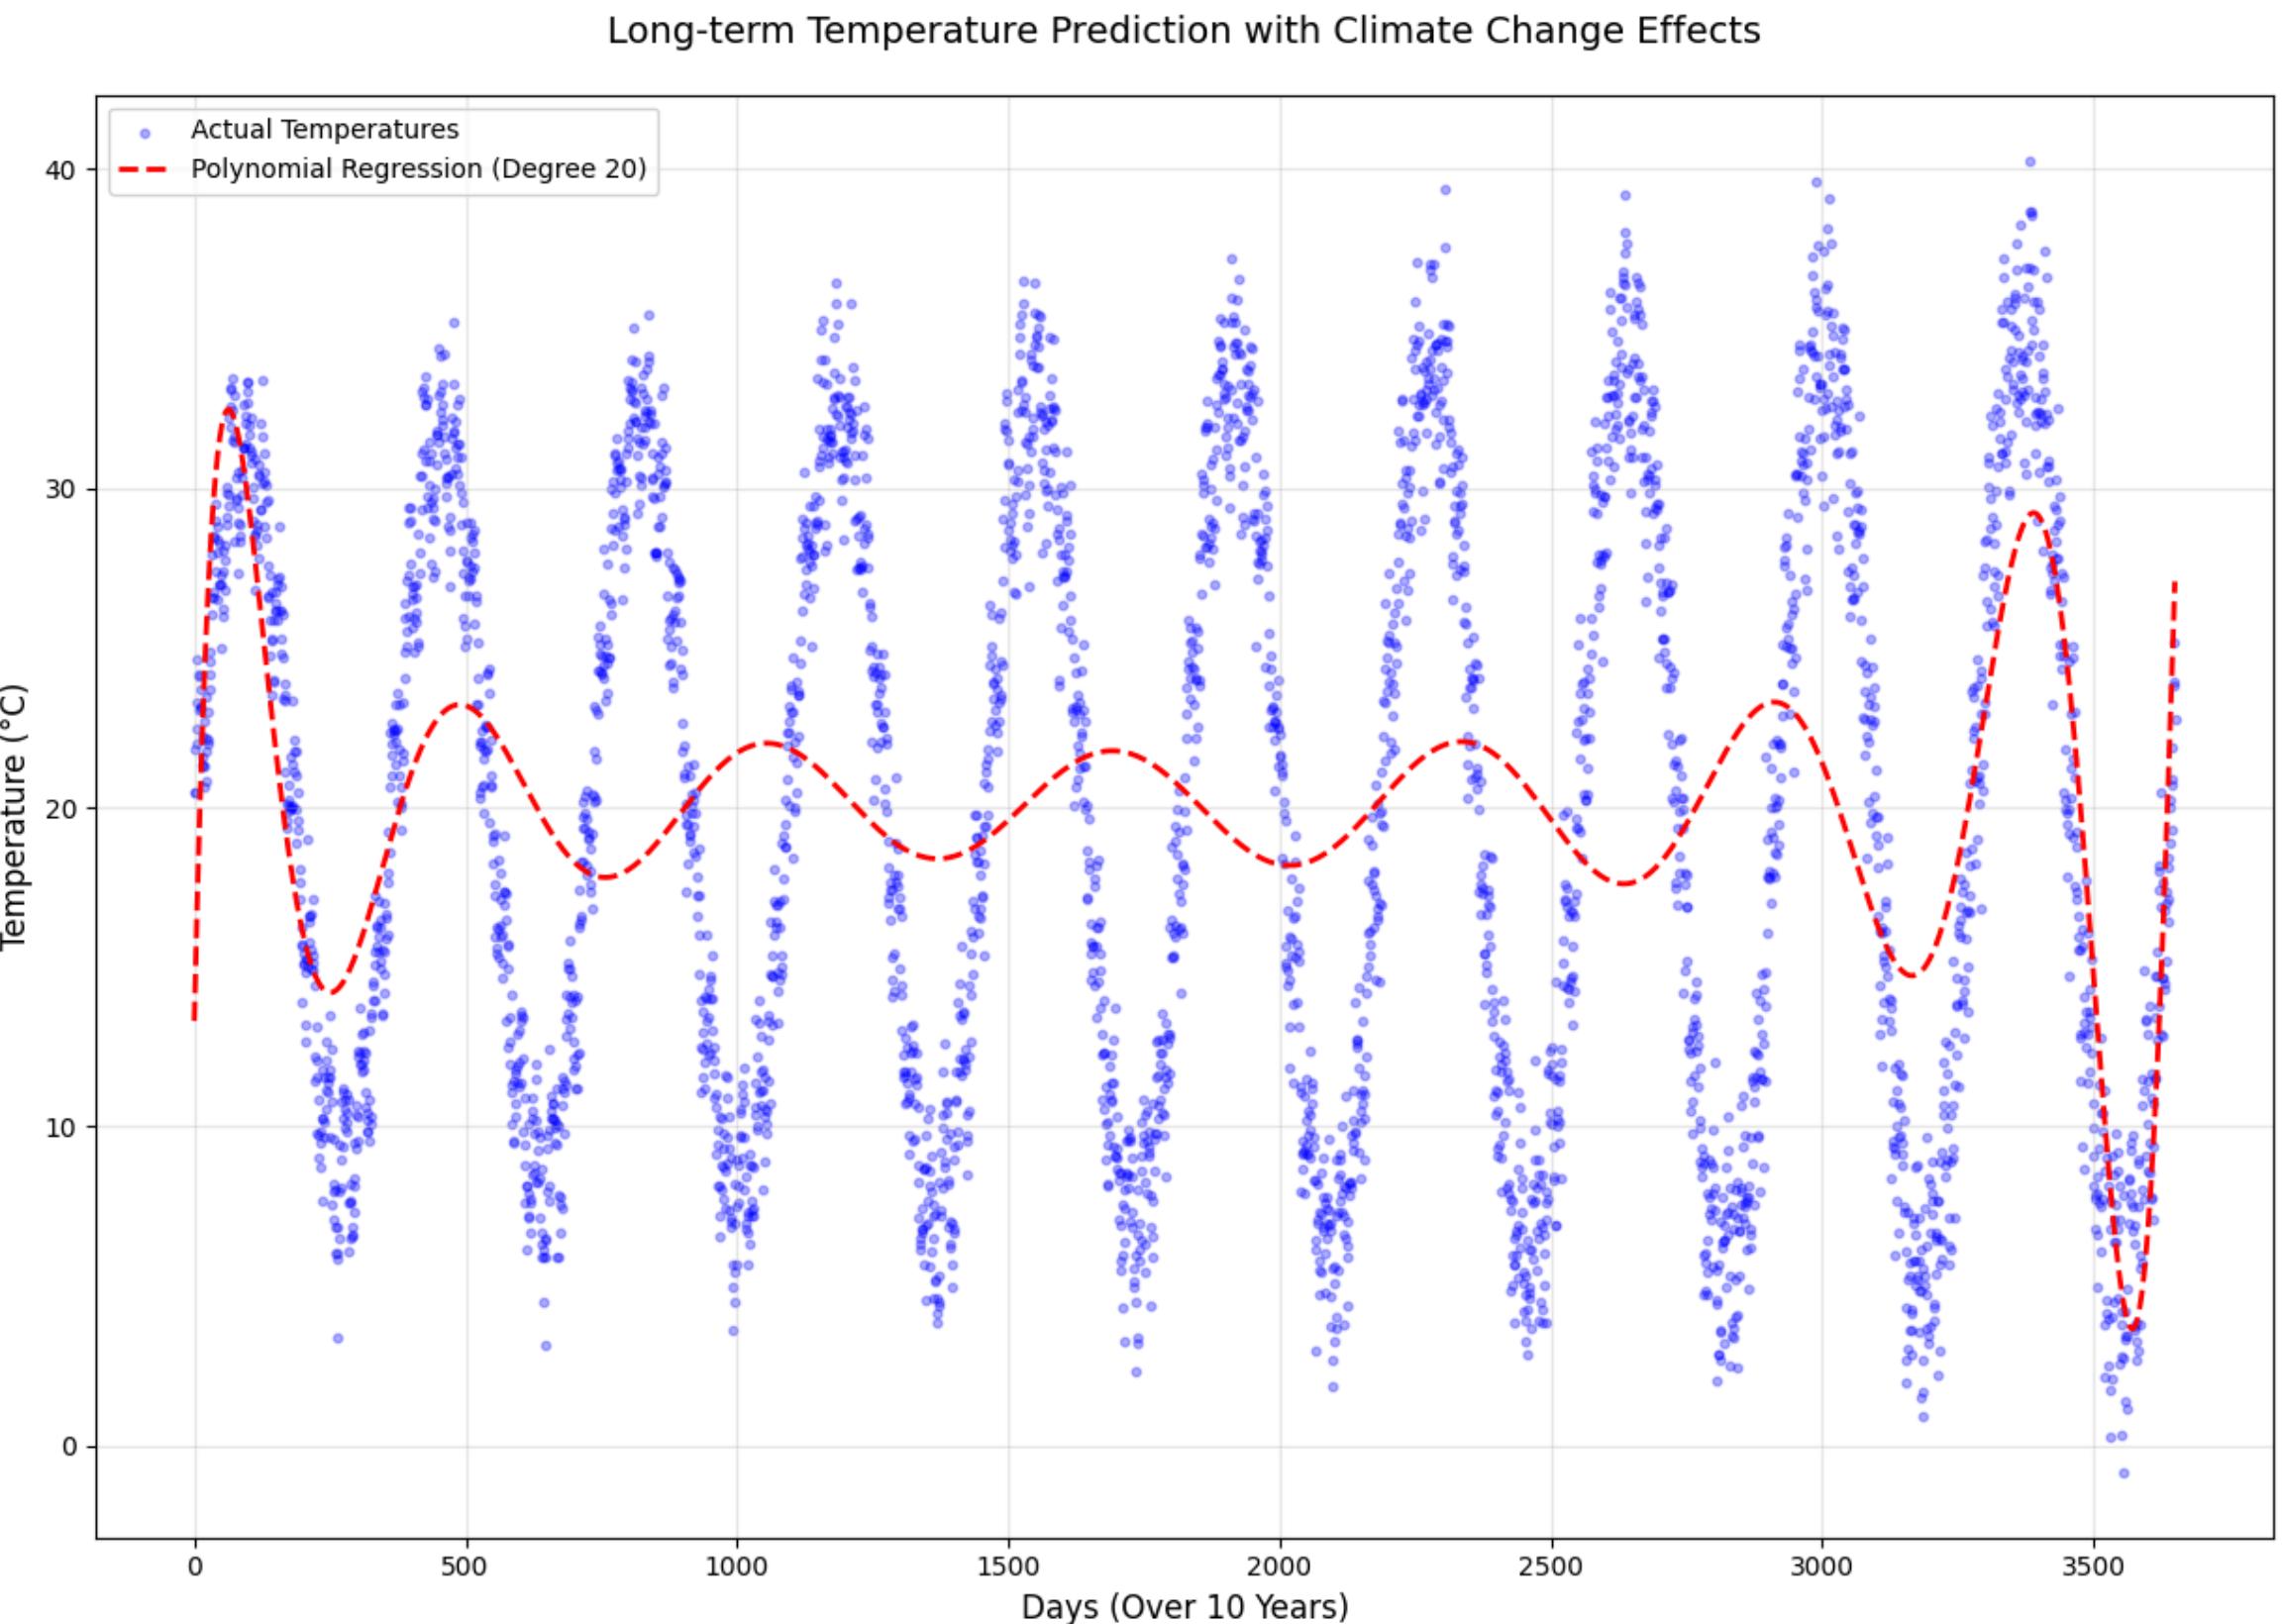
# Can we fit this data?

- How about taking much more complex data?
- Can polynomial models fit the data?



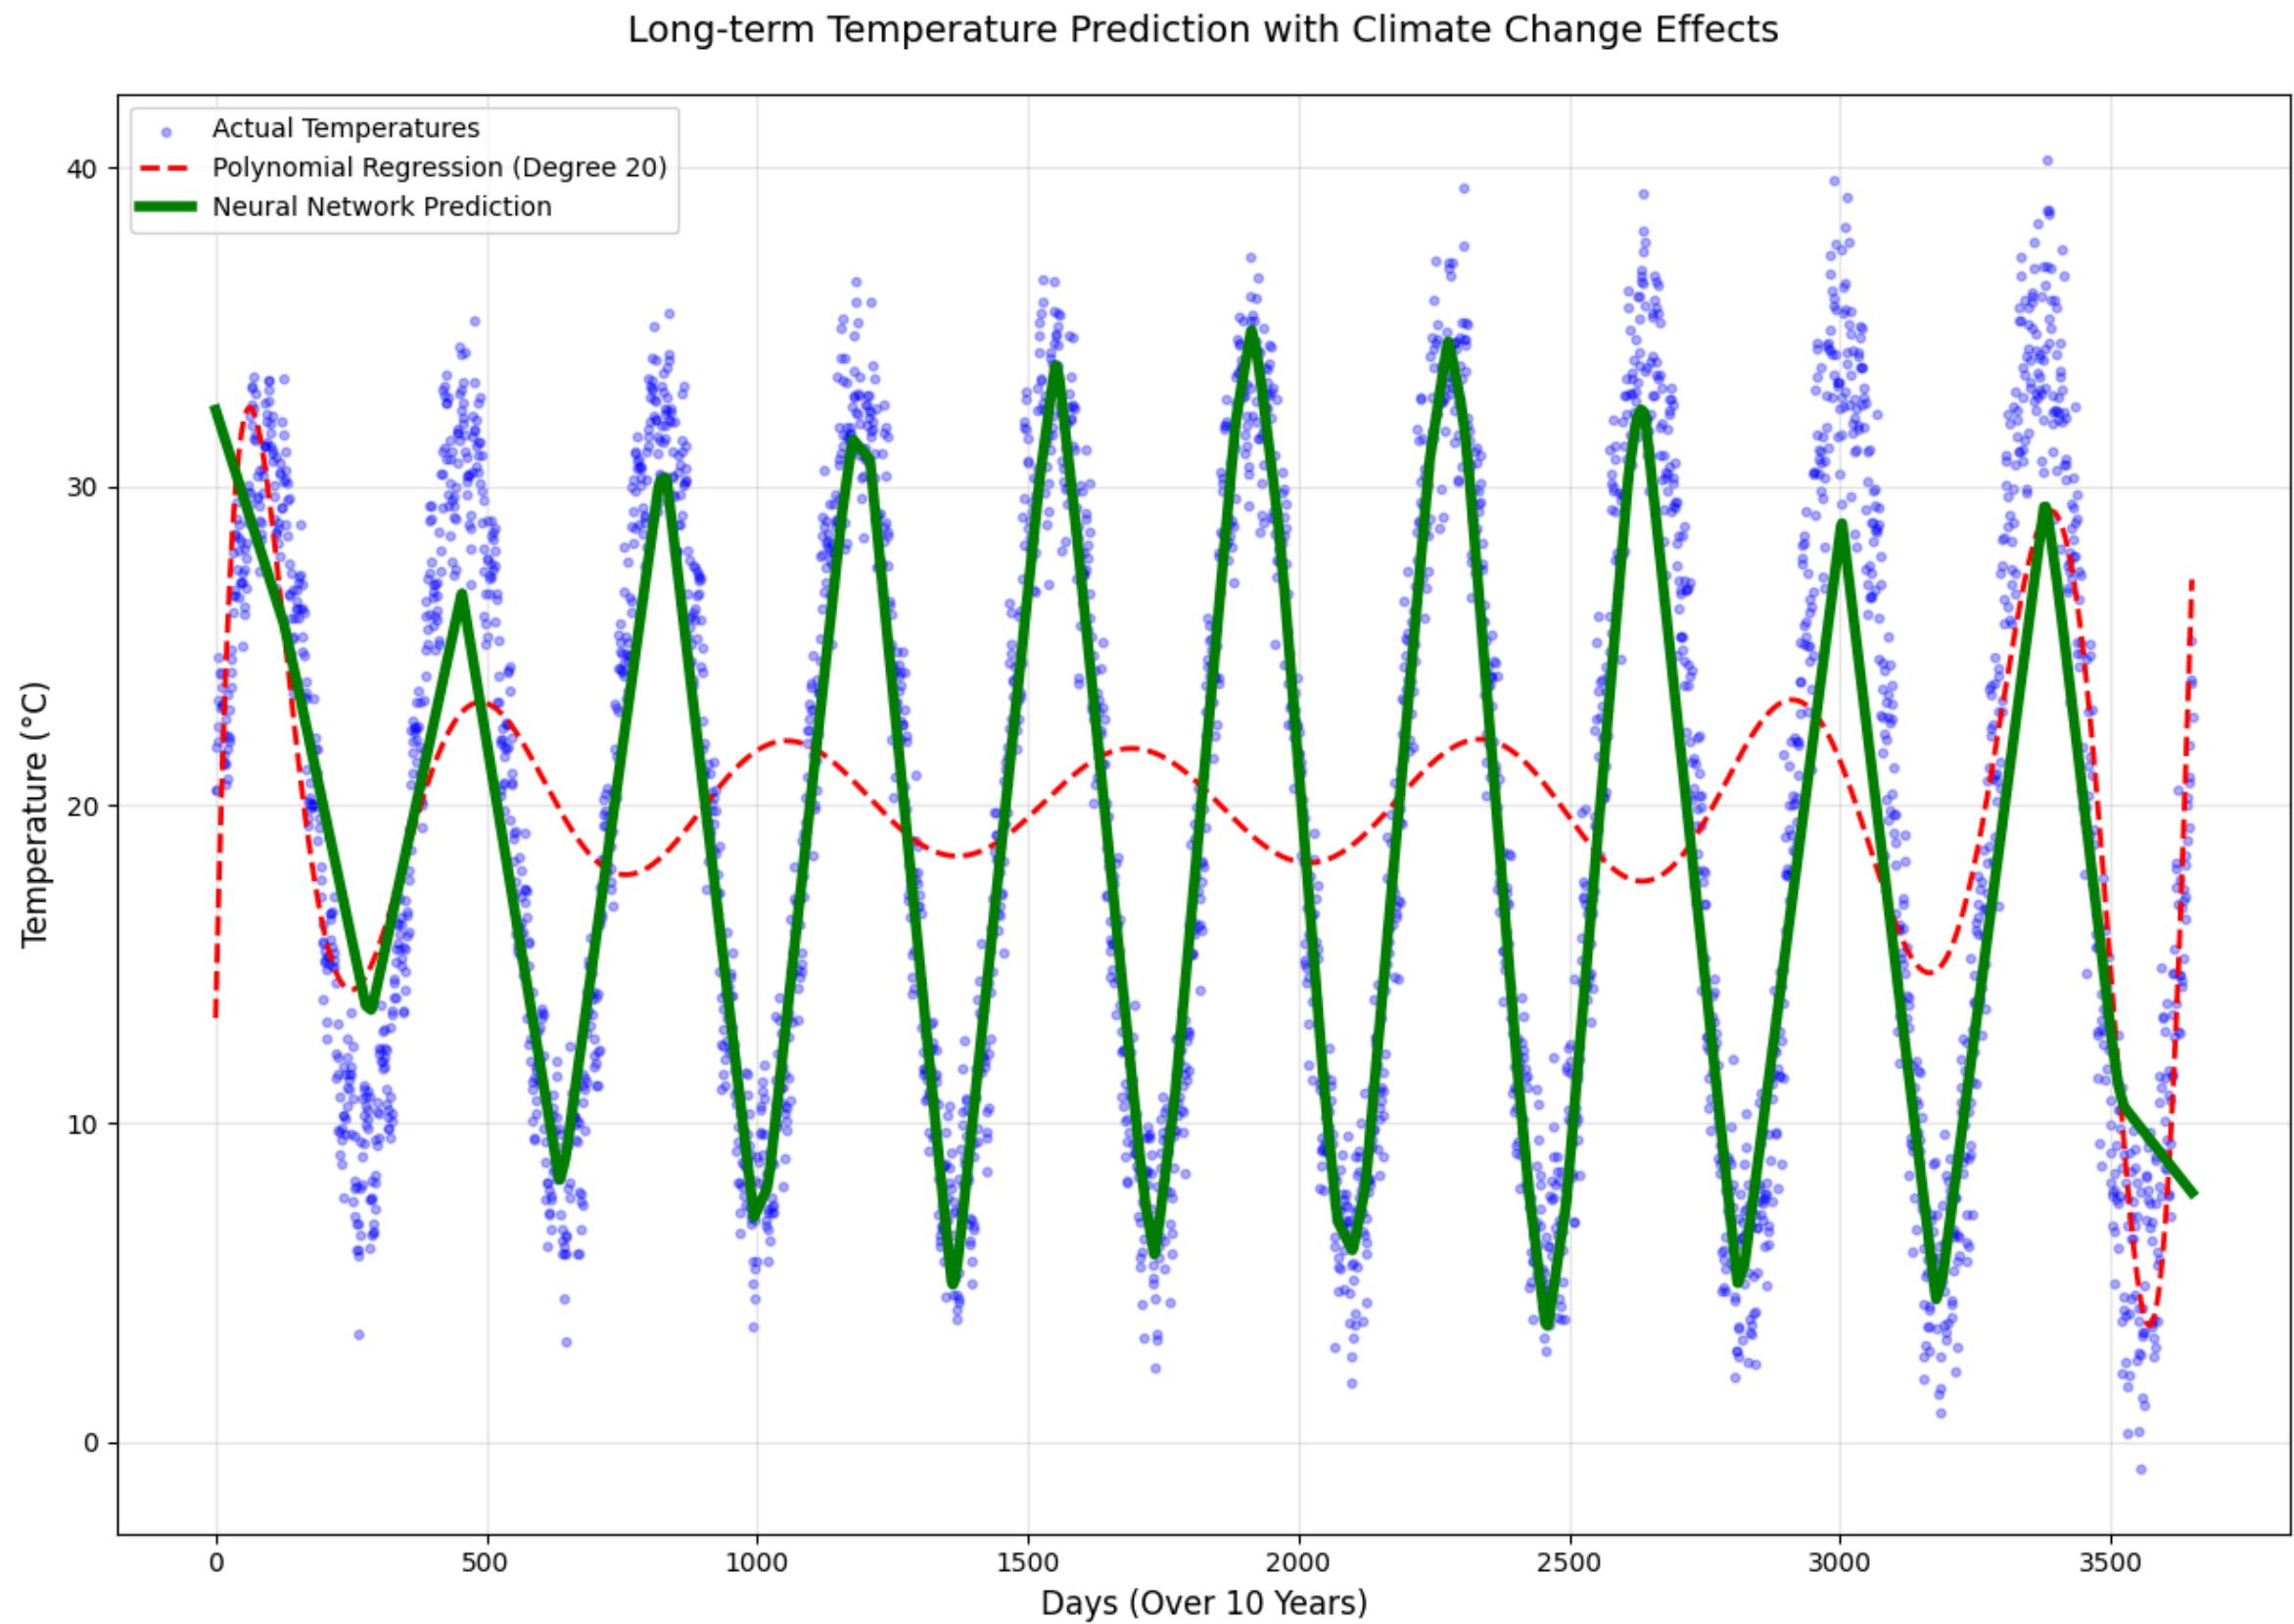
# Can we fit this data?

- Let's try polynomial regression
- It seems like the model failed to fit the data



# Can we fit this data?

- Let's try Neural Networks
- It appears that the neural network excels at adapting to this complex data!



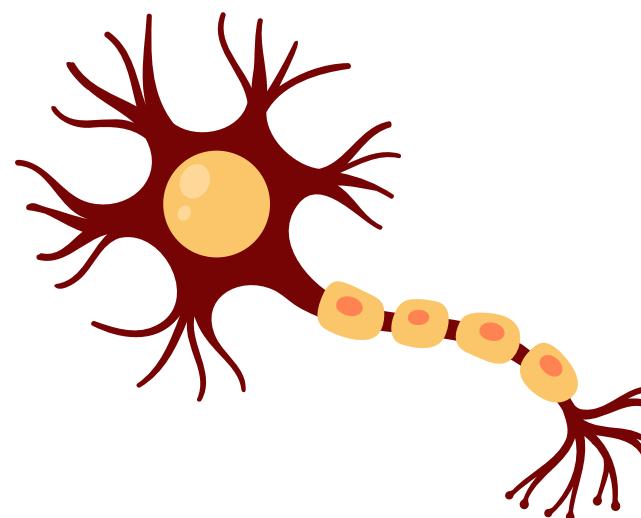
# Neural networks emerge as the champion.

- Neural networks can adapt to evolving patterns over time.
- They can automatically learn and model complex interactions between multiple factors

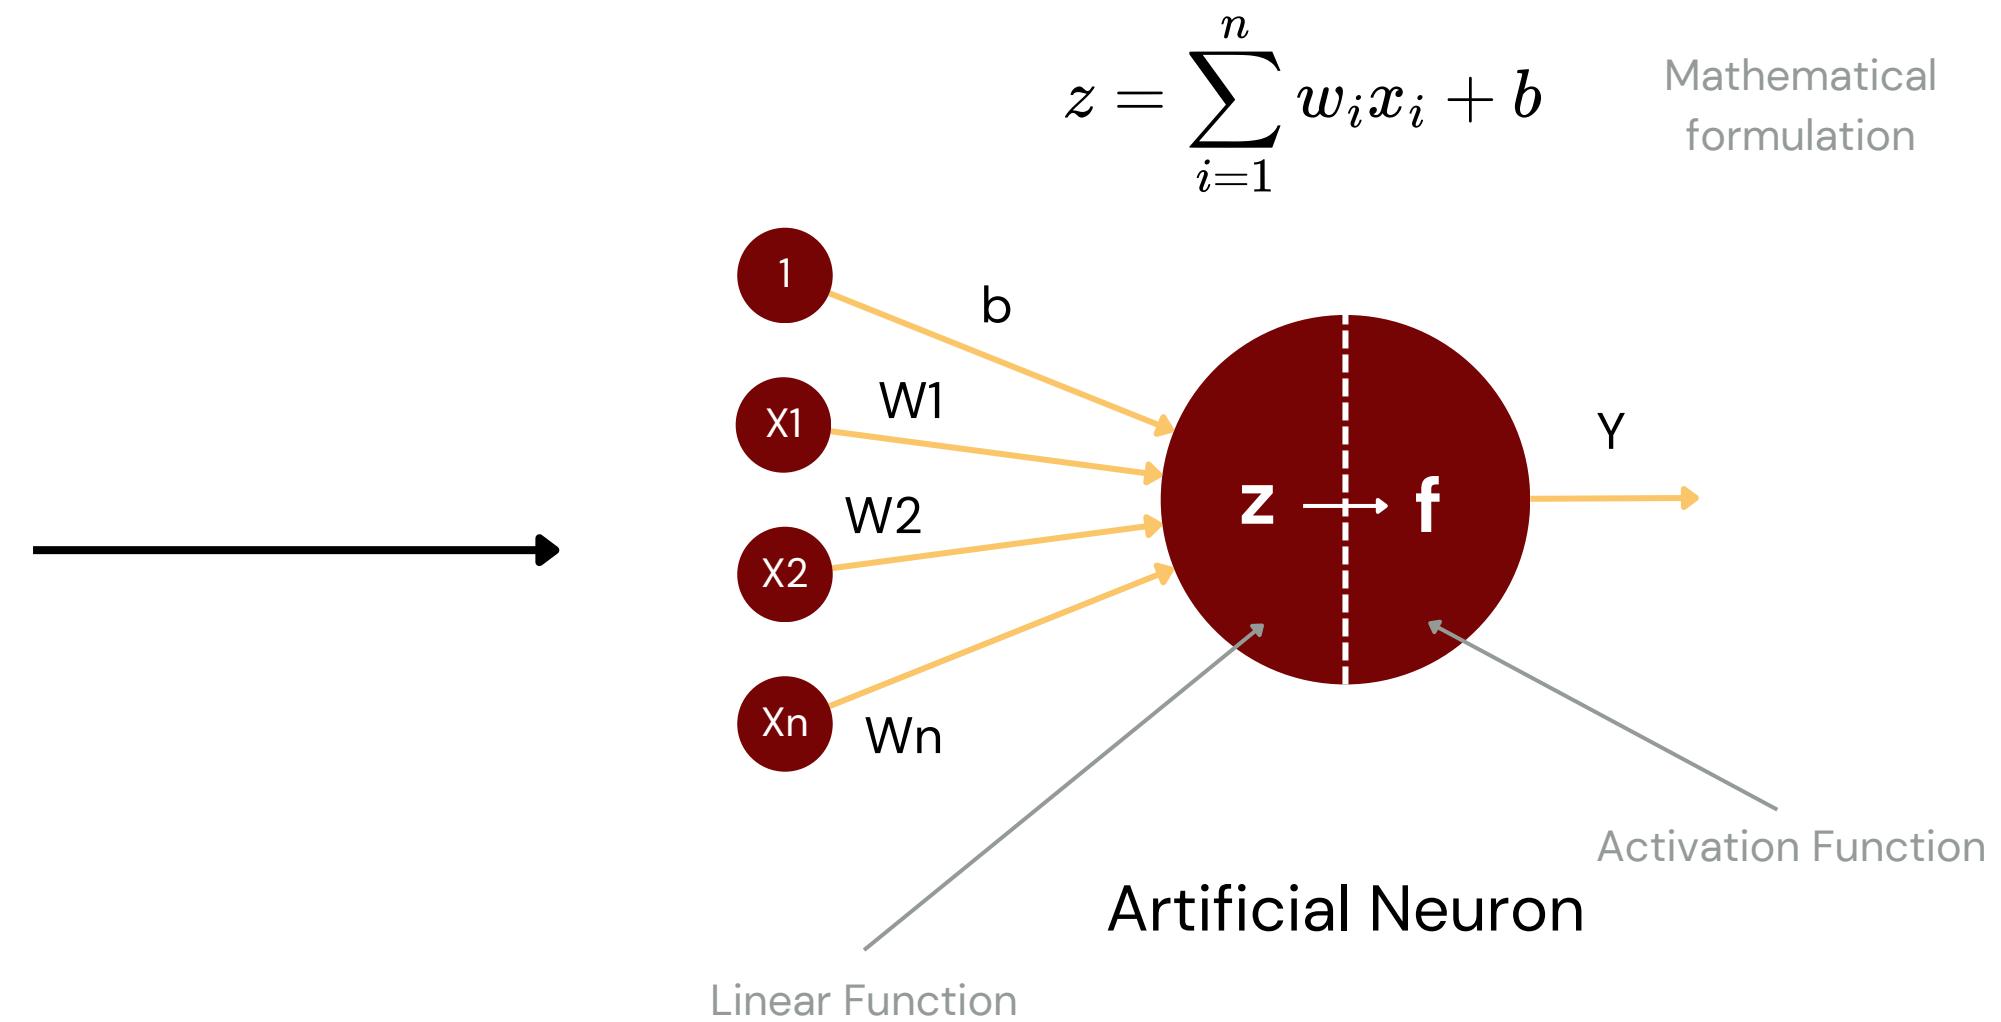
# How Do Neural Networks Function?

# A Single Neuron

- A single artificial neuron is a mathematical model or function that takes an input, does a specific calculation based on weights and adds a bias

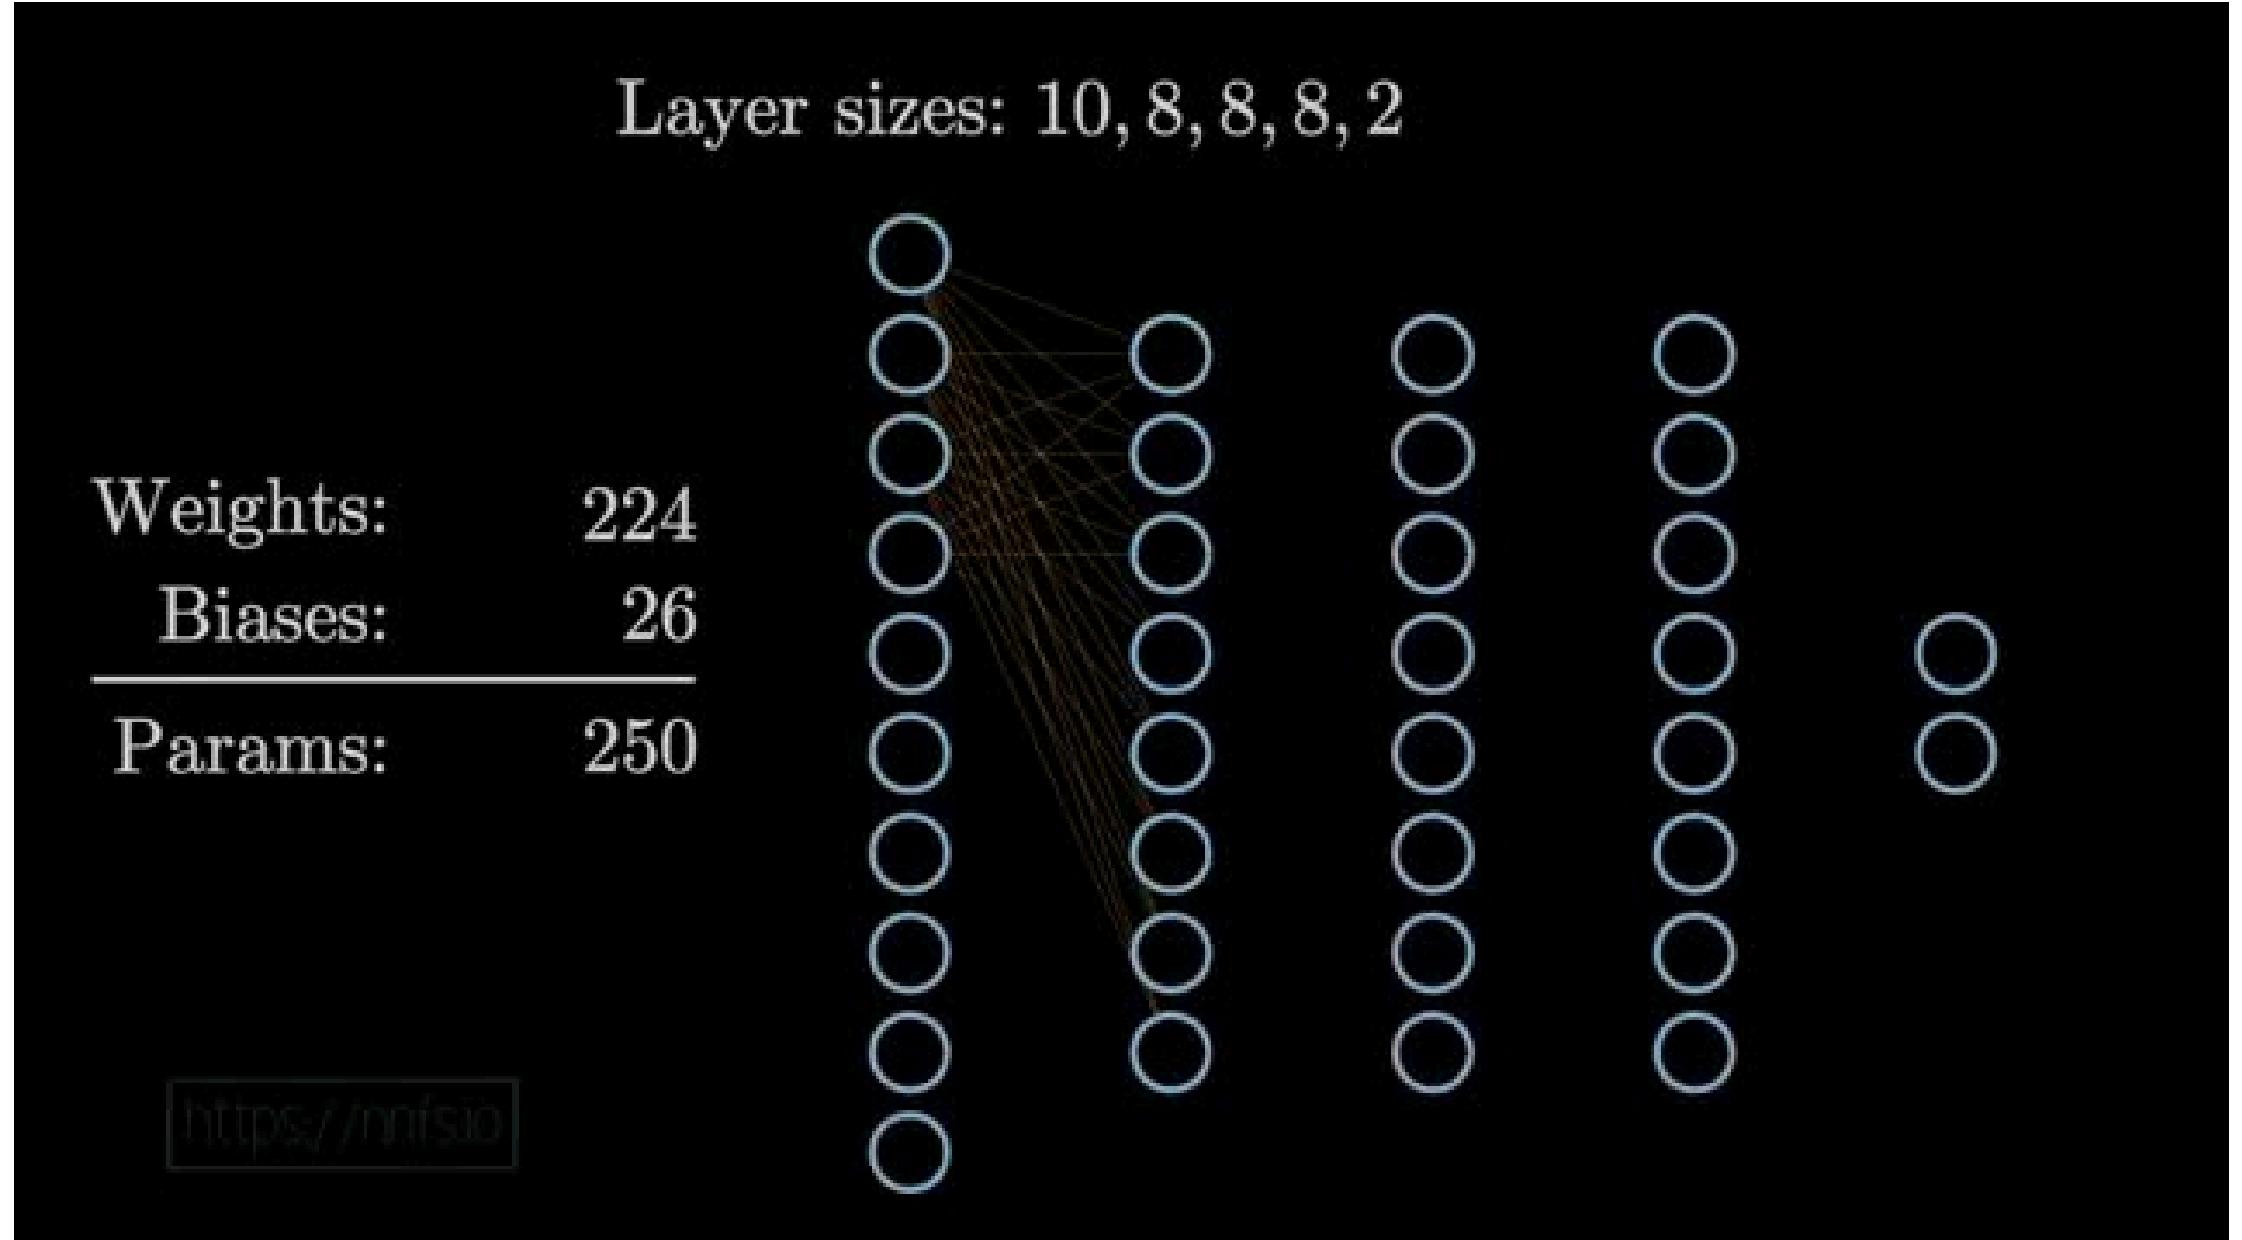


Biological Neuron



# Multiple Neurons

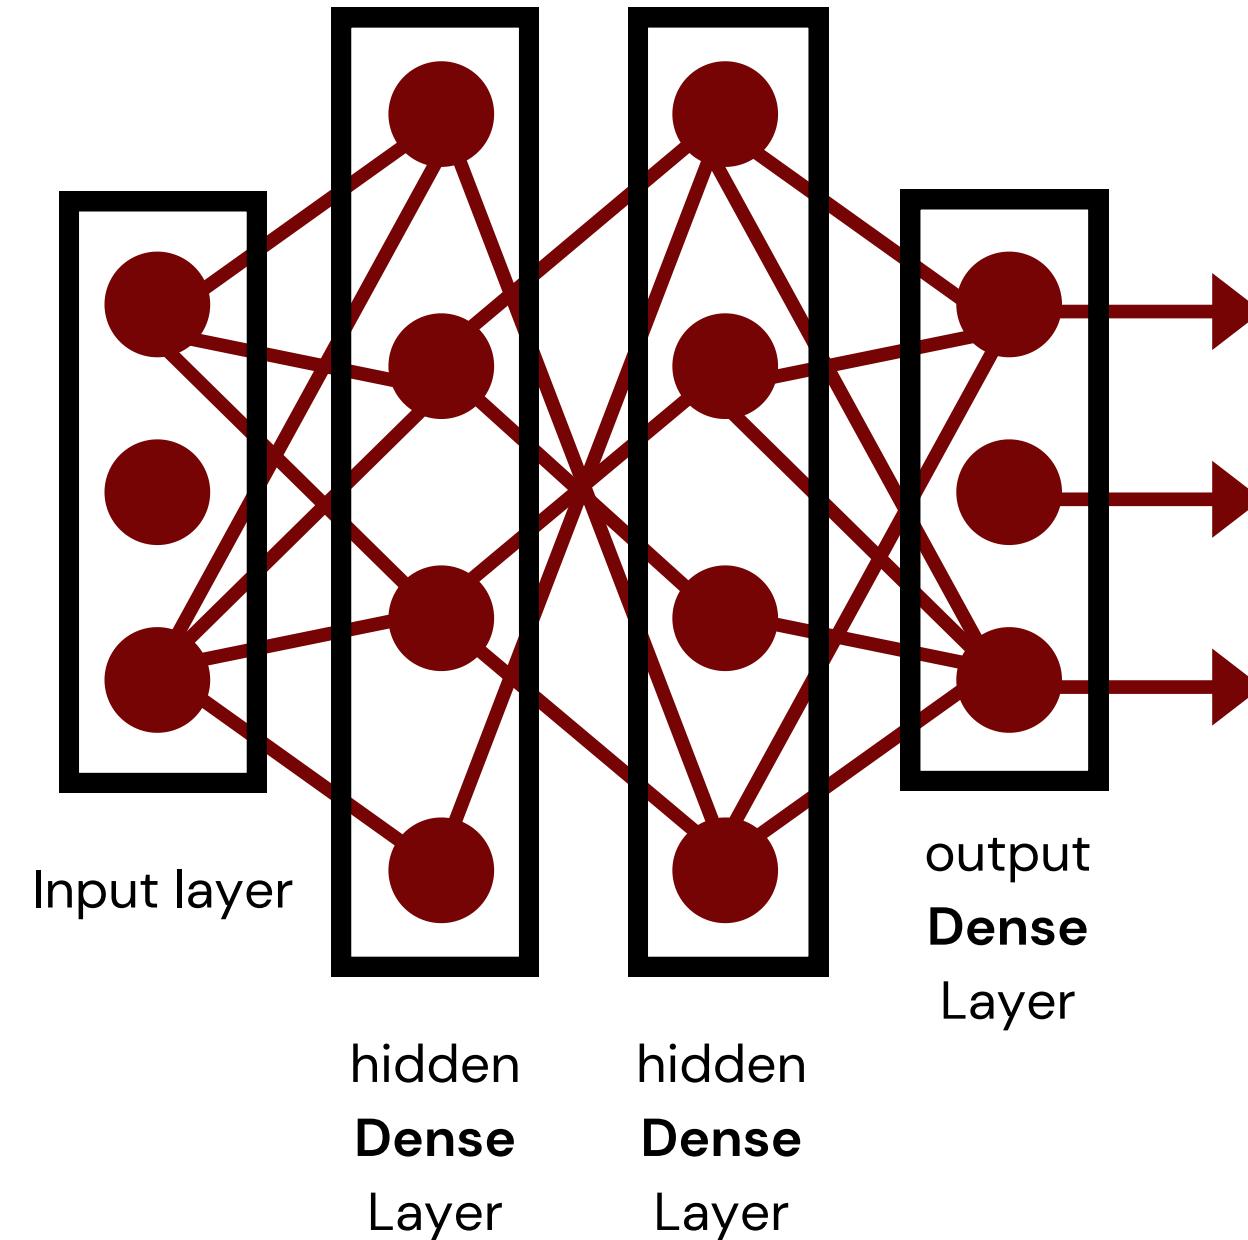
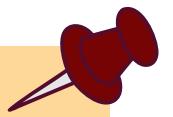
- A single neuron by itself is relatively useless.
- When interconnected, numerous neurons can produce results that often exceed other machine learning methods.
- Due to their complexity, neural networks are often seen as black boxes, as their decision-making processes are unclear.



# First Technical terminology, "Dense Layers"

- In a dense layer, every neuron connects to all neurons in the following layer.
- Each connection between neurons has a weight associated with it, which is a **trainable parameter**.
- A bias is added to help fit the real-world types of dynamic data.

What allows a Neural Network to fit non-linear data is not just the network itself, but the **activation functions** used within it!

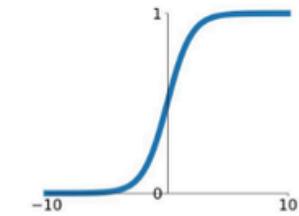


# What is an Activation Function ?

- Activation functions help a neural network learn and model complex patterns by introducing **non-linearity**.
- They enable the network to capture non-linear relationships in data, like **curves** and **complex shapes**.
- Without them, the network would behave like simple linear regression.

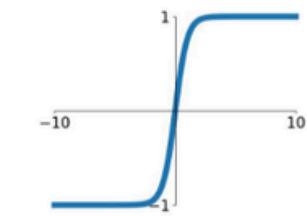
**Sigmoid**

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



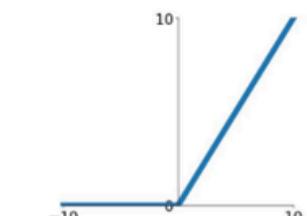
**tanh**

$$\tanh(x)$$



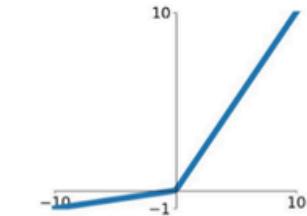
**ReLU**

$$\max(0, x)$$



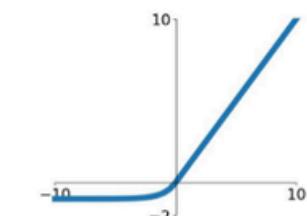
**Leaky ReLU**

$$\max(0.1x, x)$$



**Maxout**

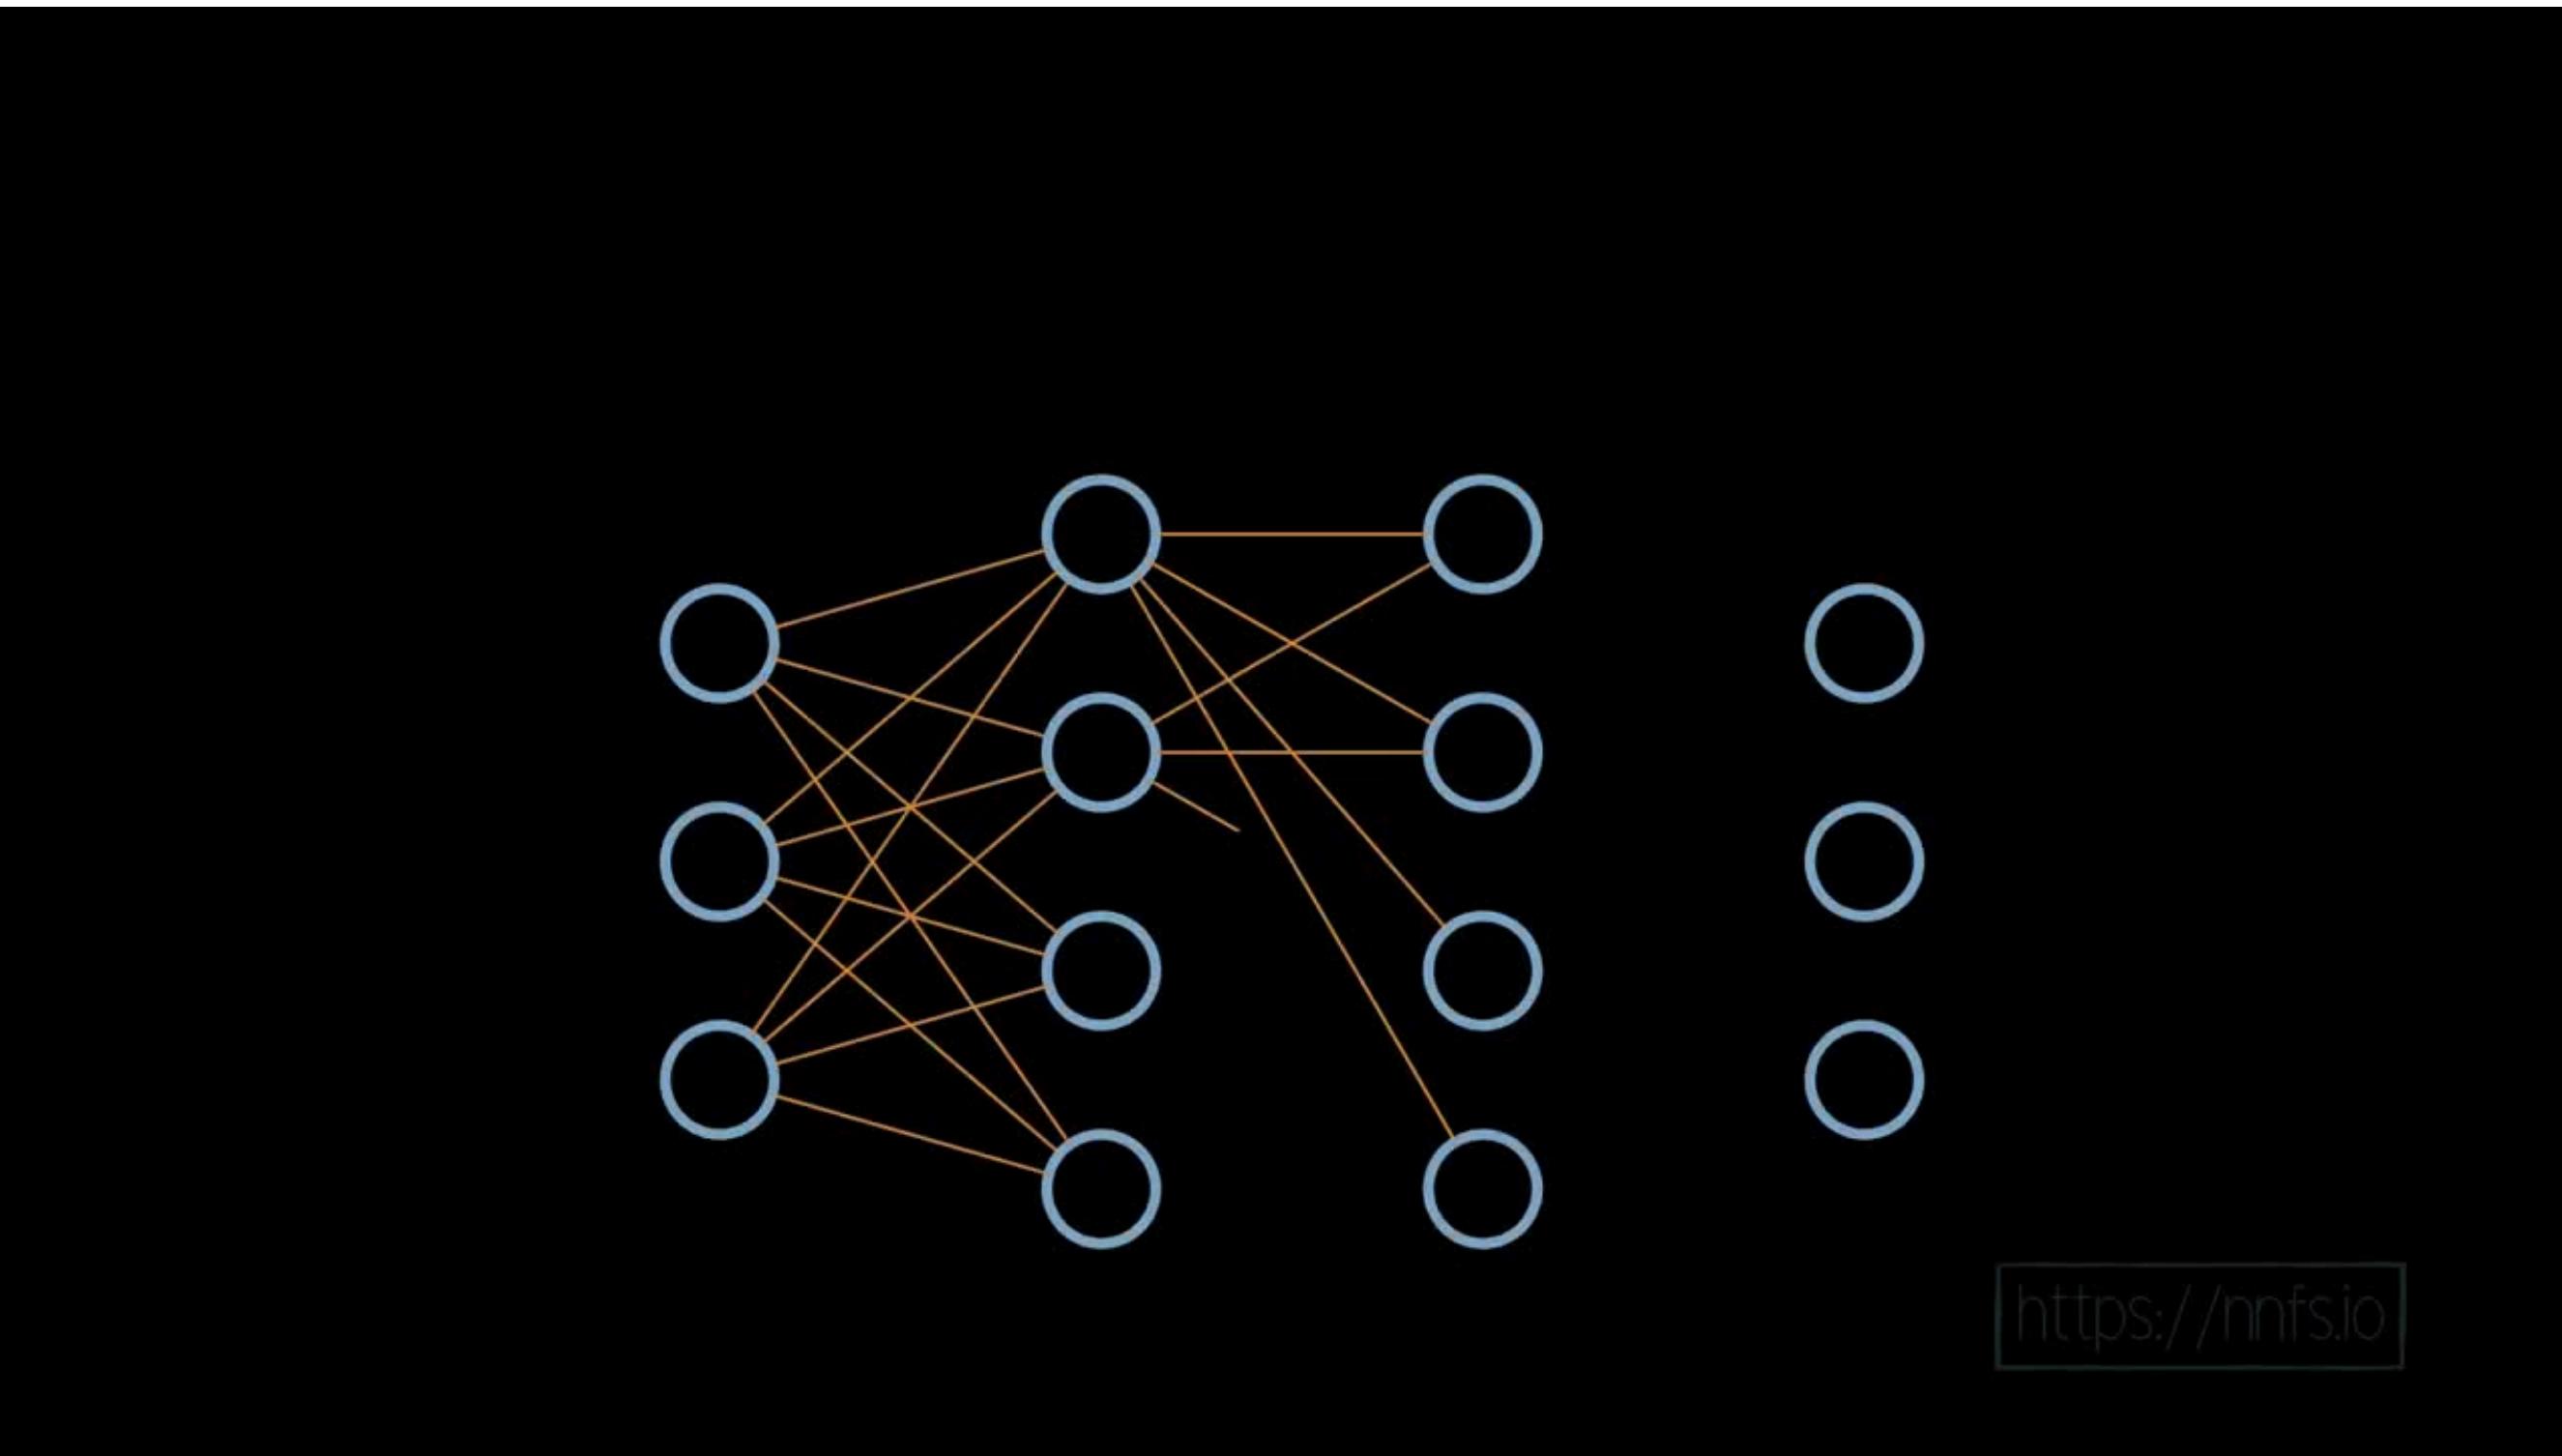
$$\max(w_1^T x + b_1, w_2^T x + b_2)$$



**ELU**

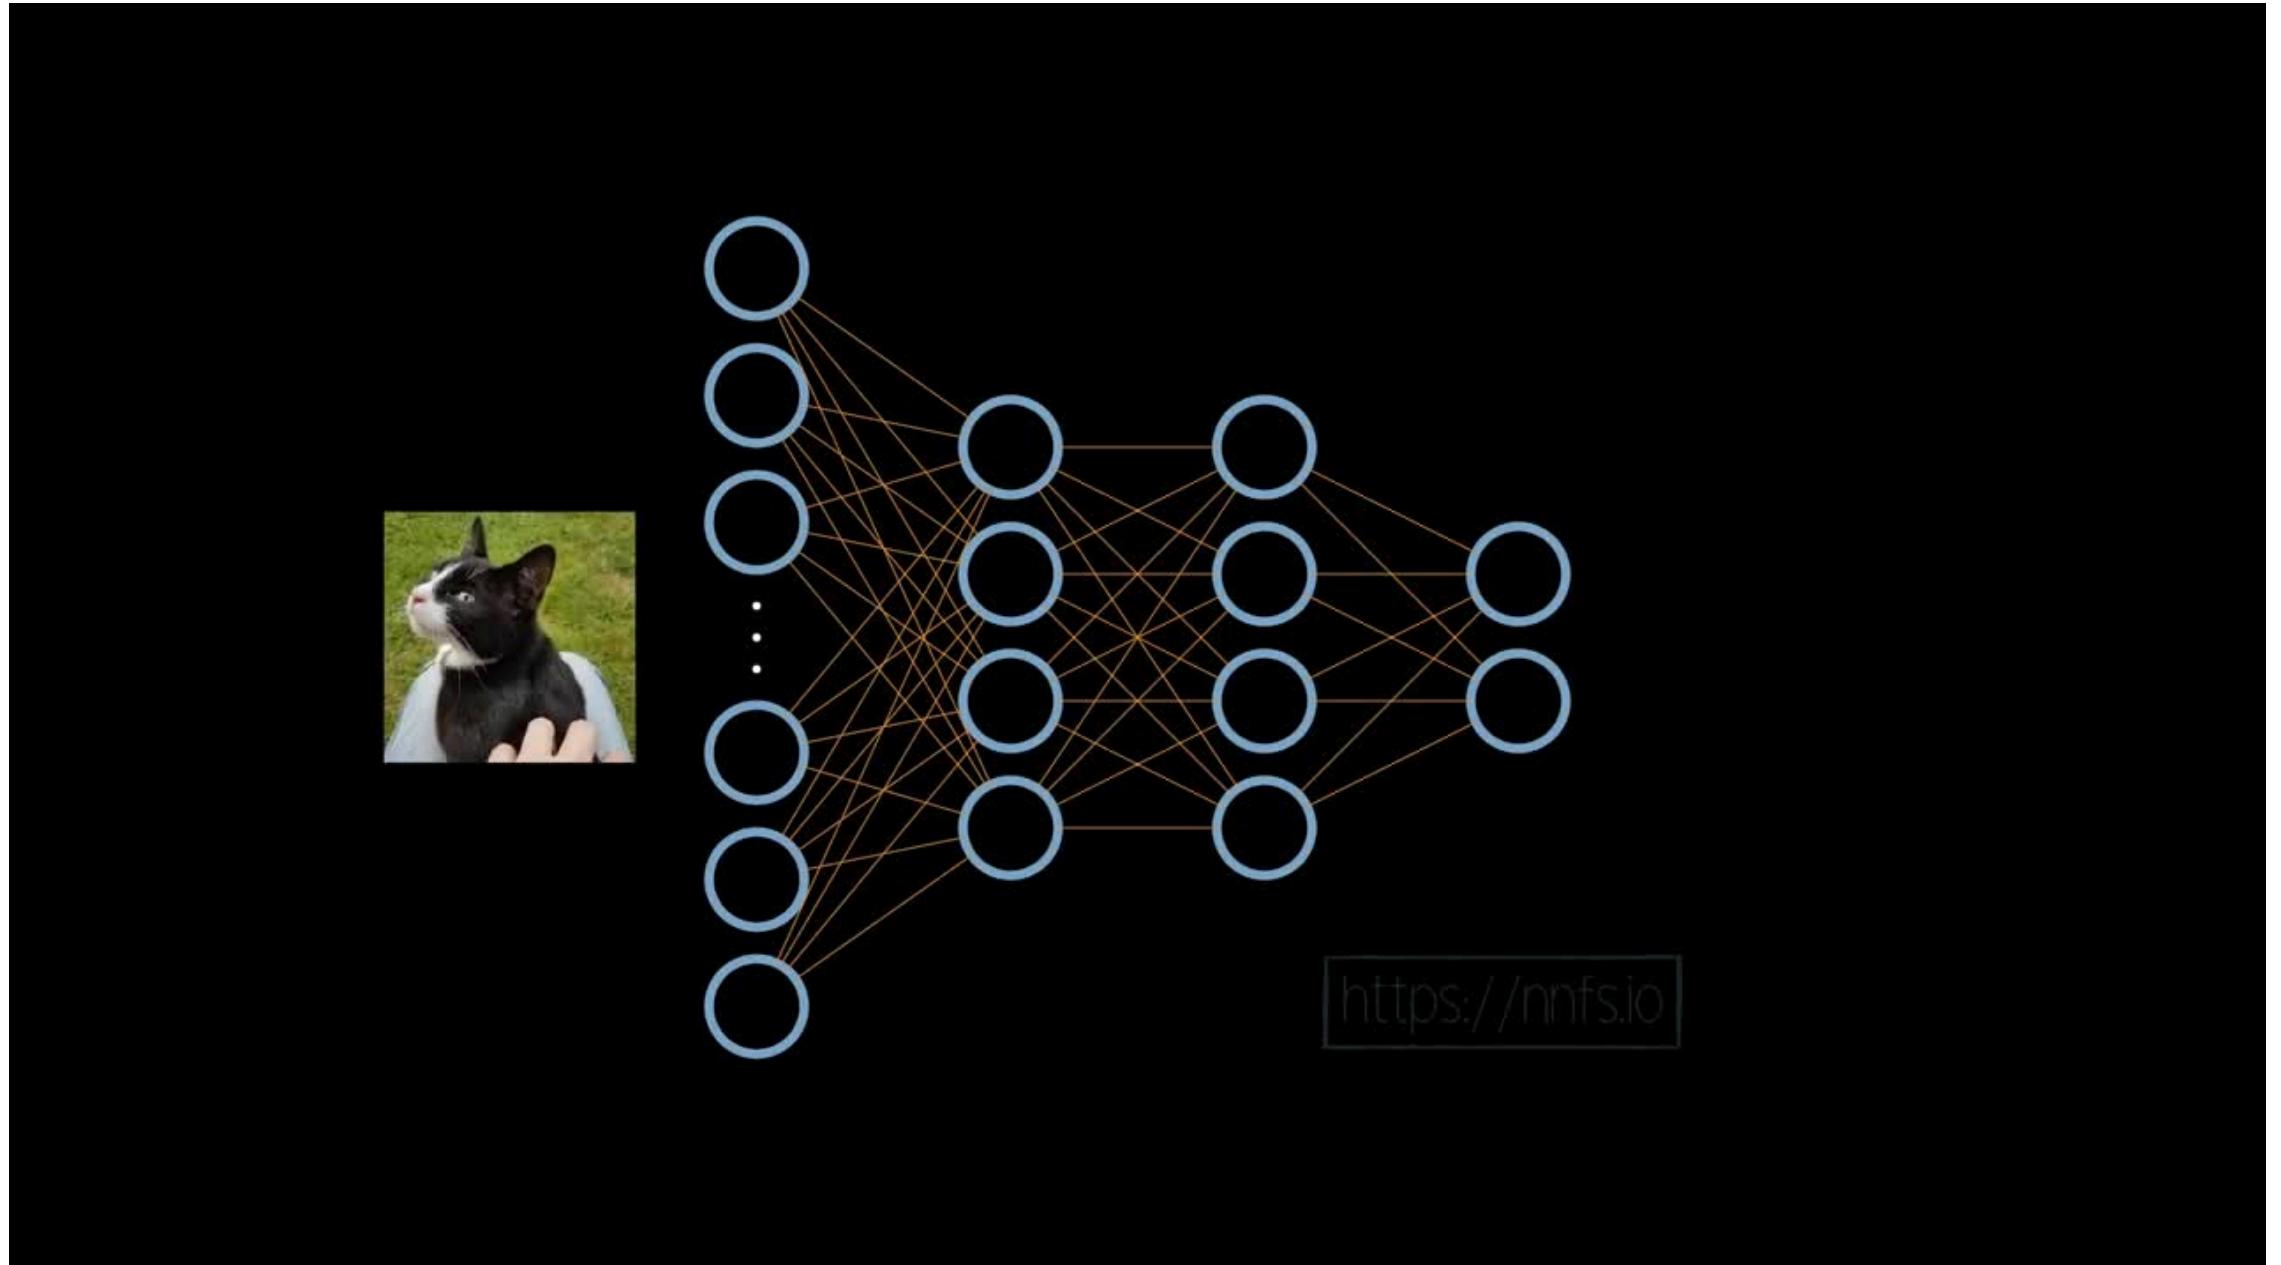
$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$

# Weights, biases on a single neuron.



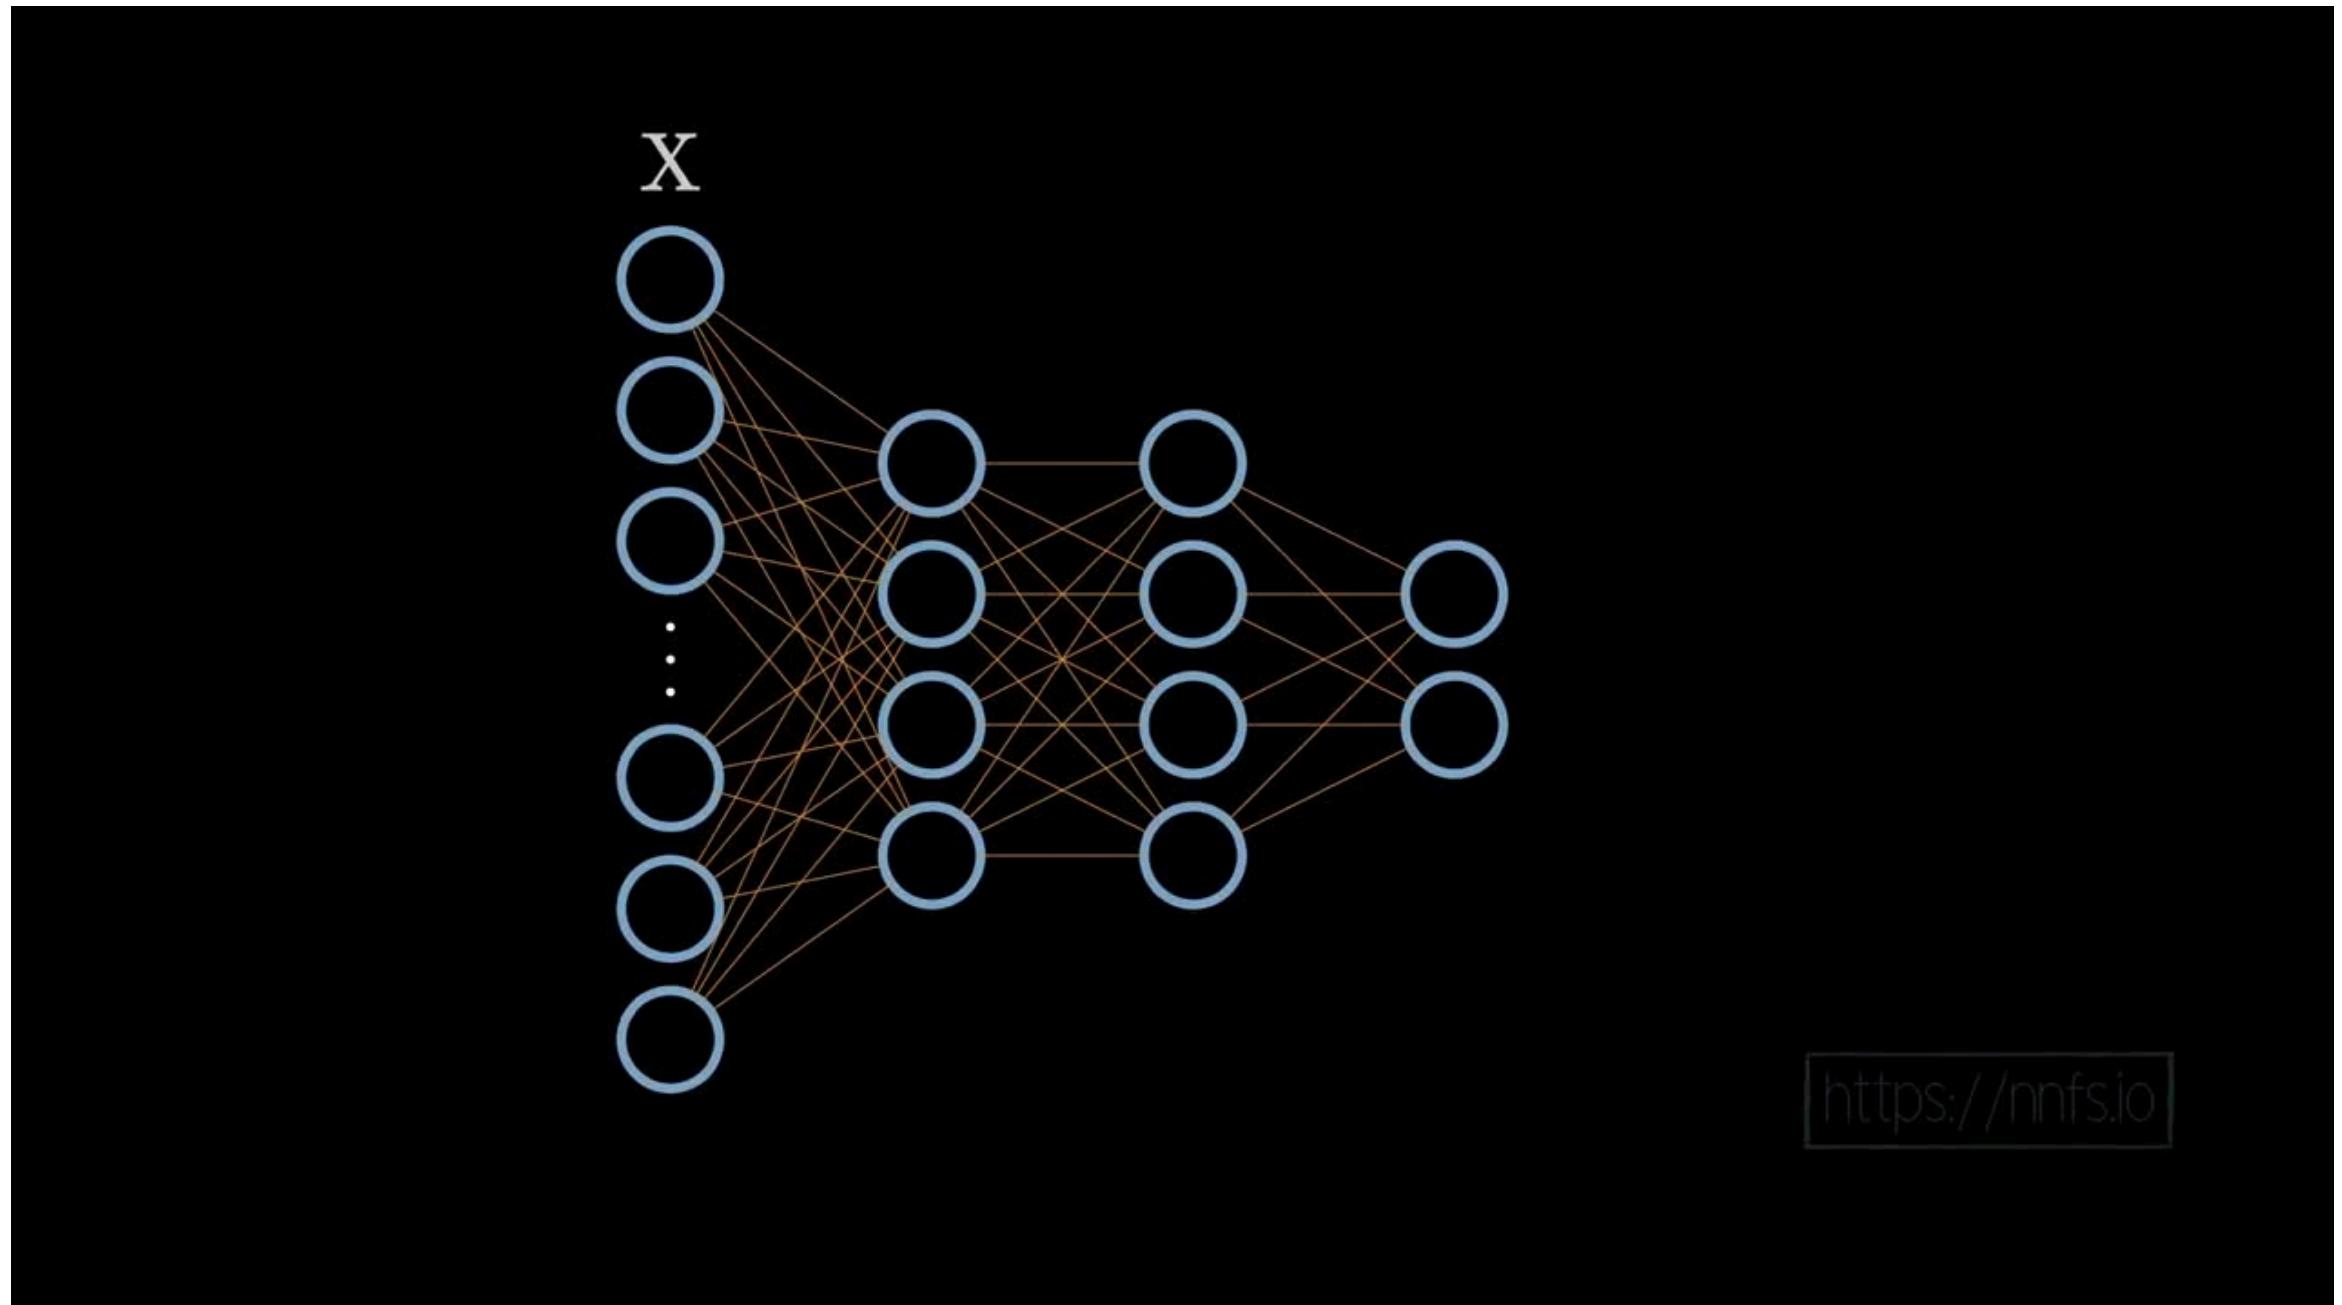
# How do Neural Networks perform classification?

- Input Layer: Holds the input data (image pixels)
- Data is often preprocessed (scaled/normalized) to fit a range like 0–1 or -1–1.
- Each neuron will produce one output based on input.



# Lots of mathematics is behind these networks.

- Neural networks are computationally intensive models!
- The process of transforming the input into a label or a number is called **forward propagation** or **model inference**.



# Let's implement a Neural Network

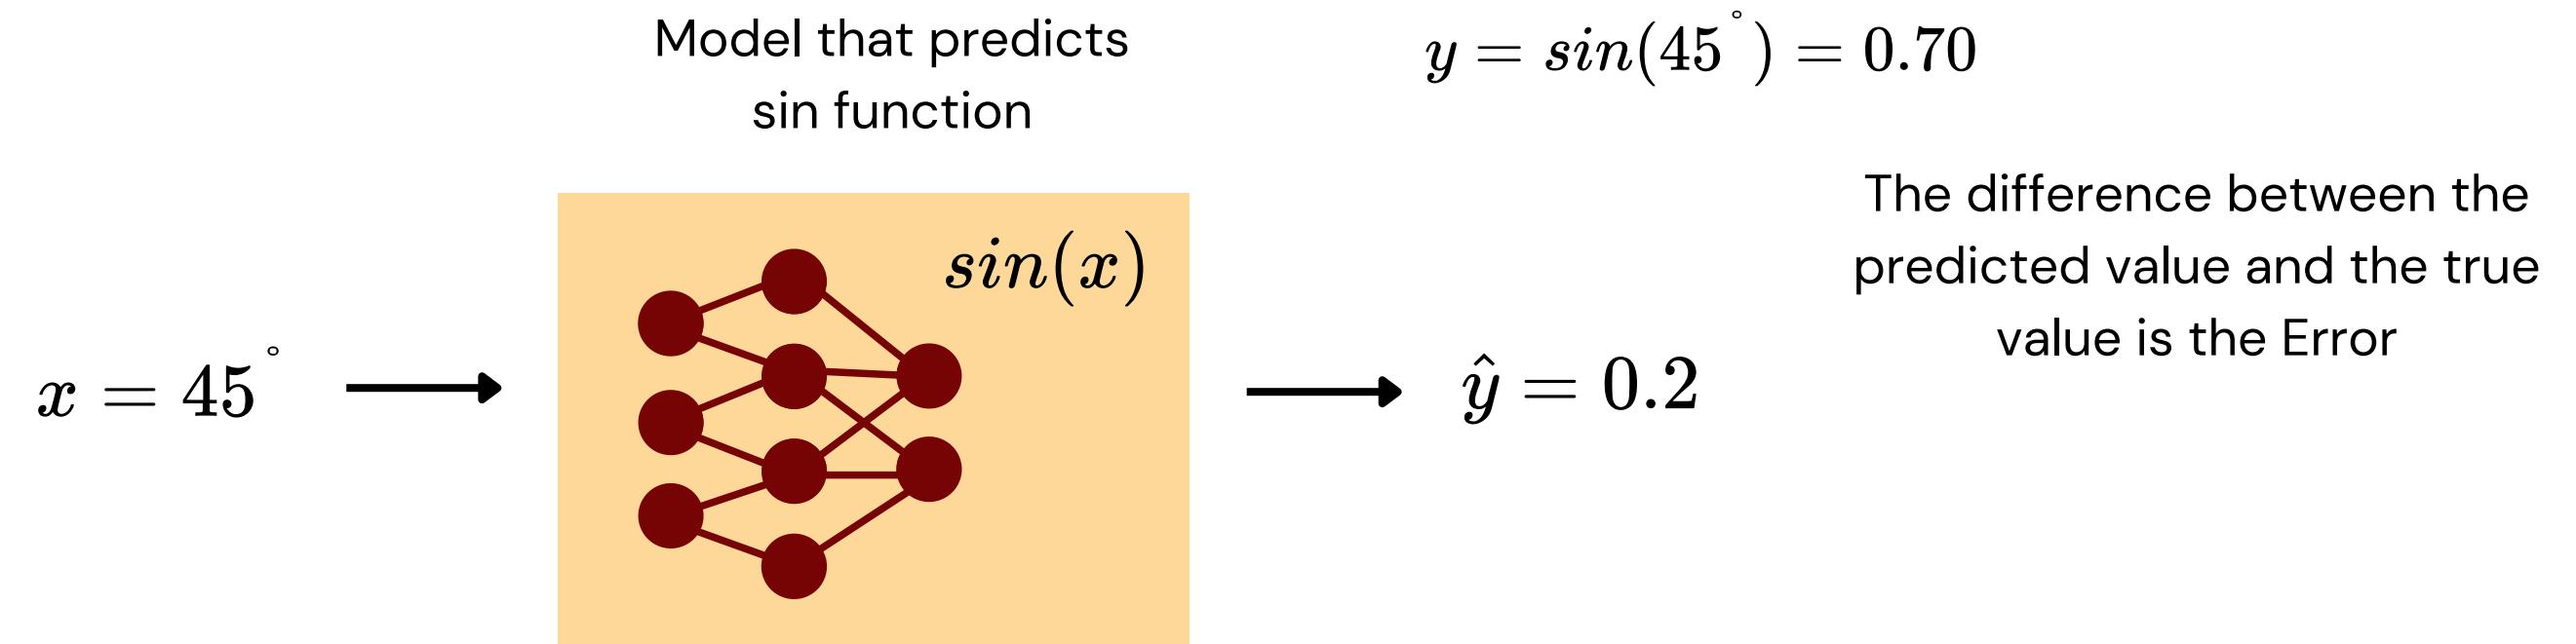
Scan the Drive



# How Do Neural Networks Learn?

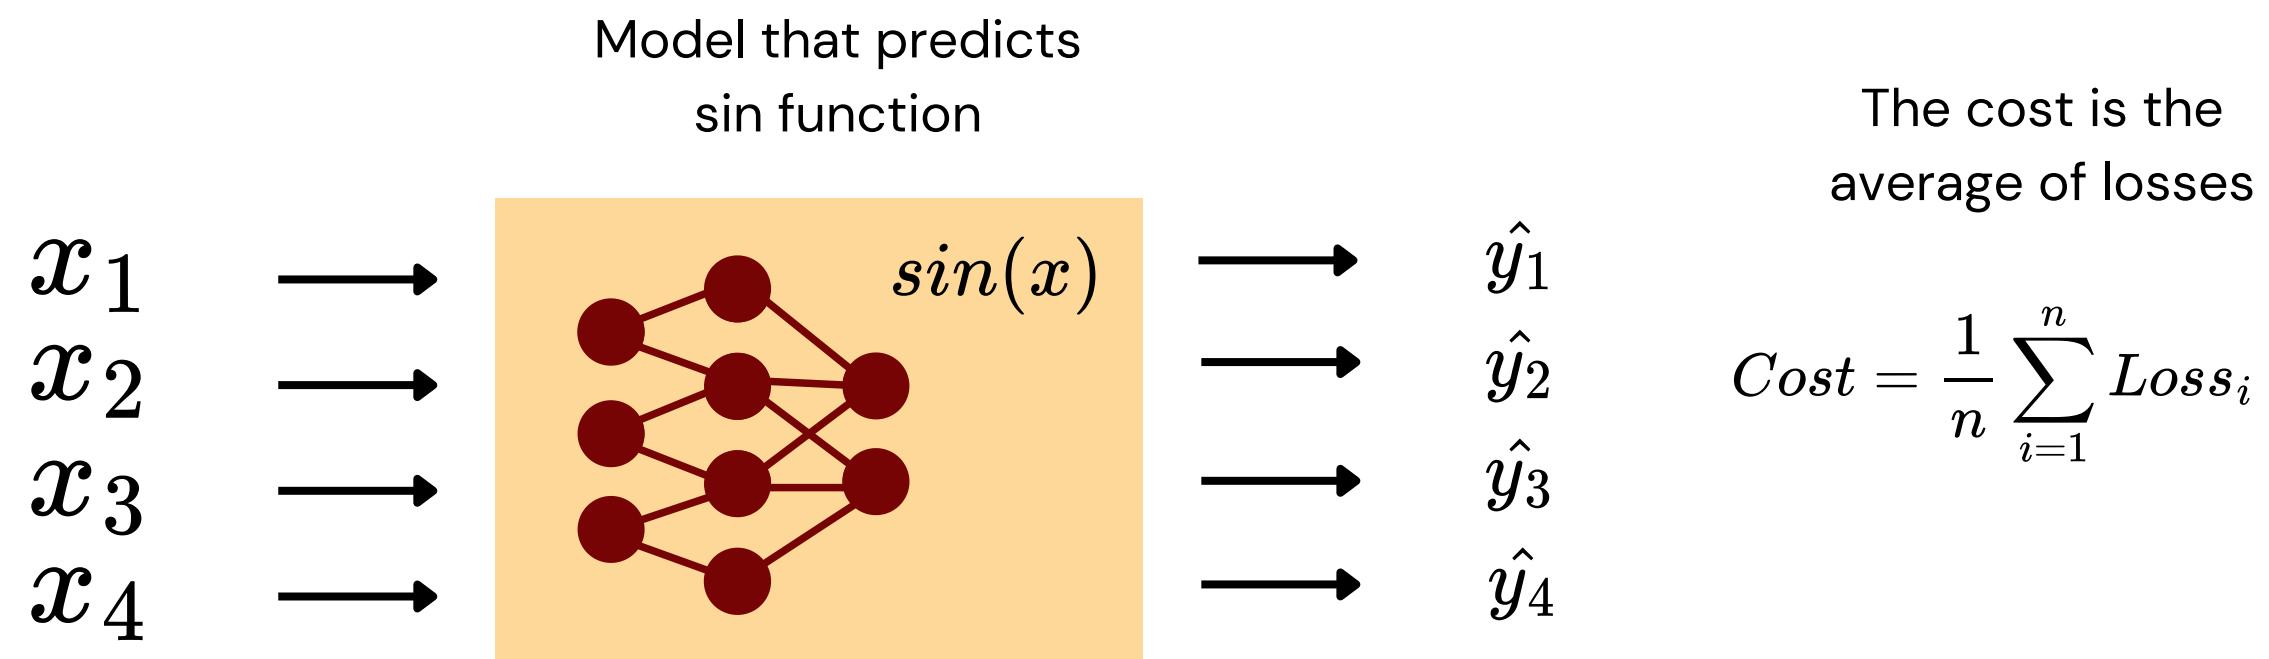
# Calculating Network Error with Loss

- The loss measures how far the neural network's prediction is from the actual correct answer for one example.
- It's like a "score" for how bad the prediction was for that single input.
- The error is assessed using various methods and formulas that depend on the nature of the problem.



# The Cost Function

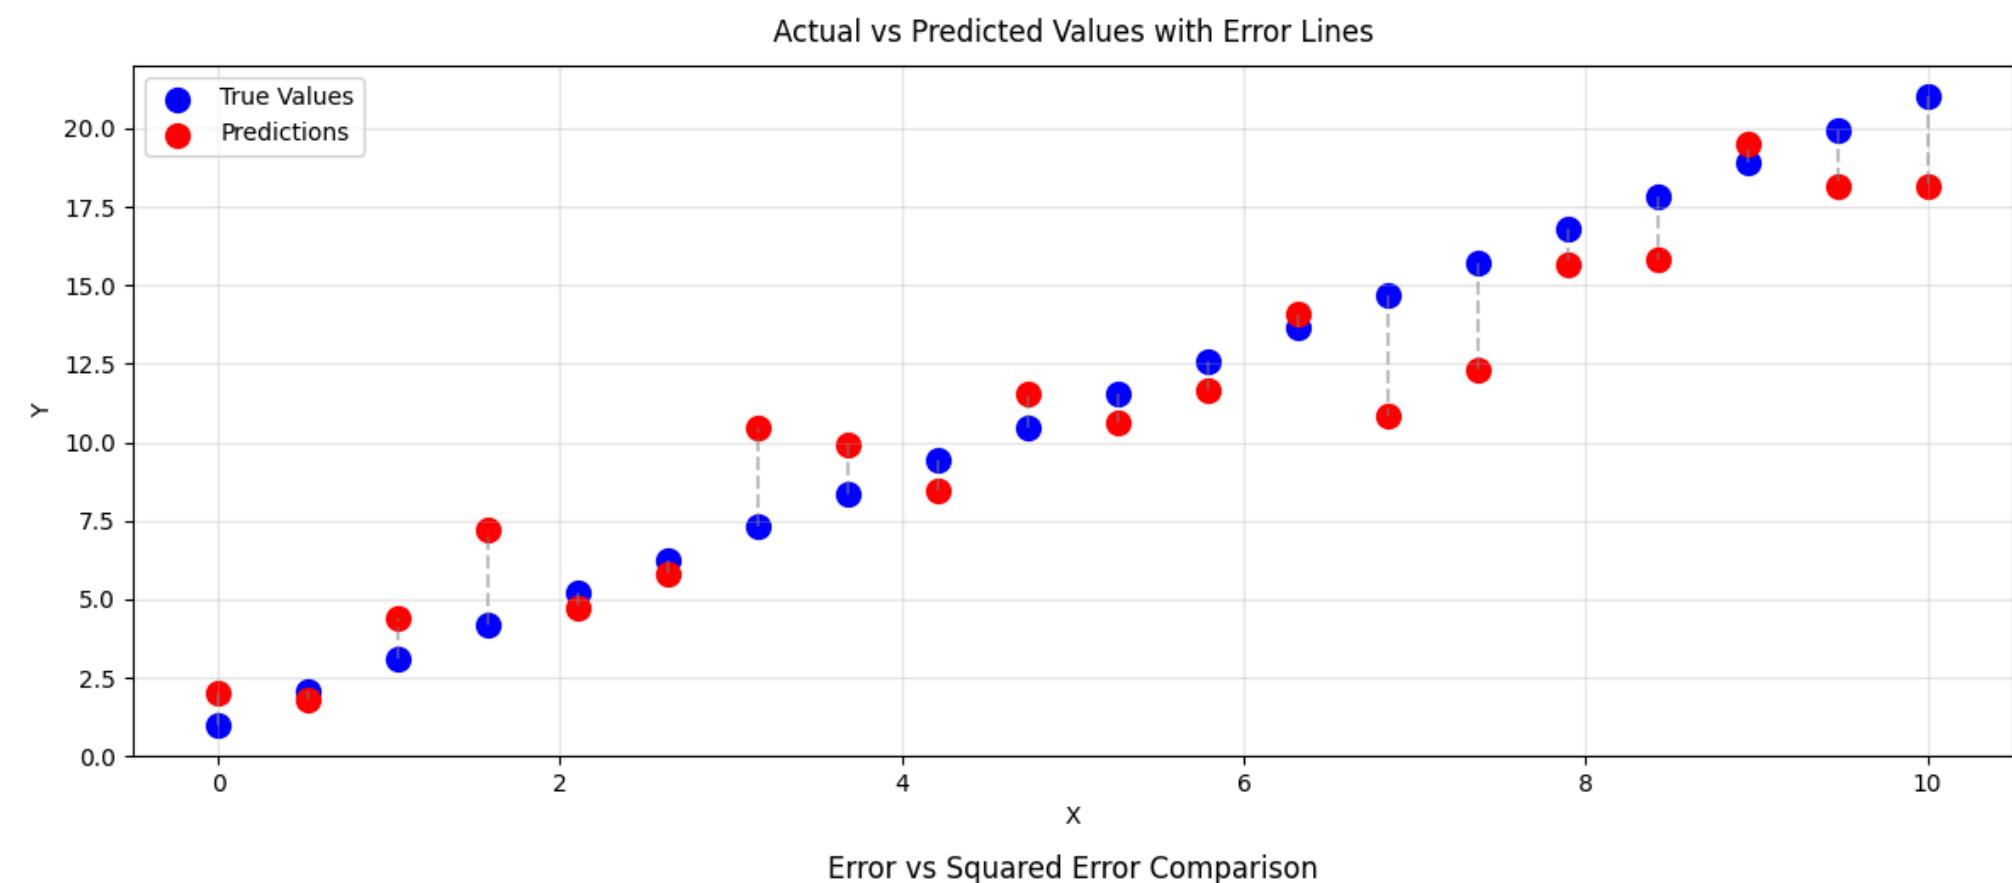
- The cost function is the average of all the losses across all examples in the dataset.
- It gives an overall "score" for how well the network is doing on the entire dataset.
- The goal of training is to minimize this cost



# Mean Squared Error (MSE)

- Measures the average squared difference between predicted and actual values.
- Use Case: Predicting numbers, e.g., house prices.
- Here is its mathematical formula:

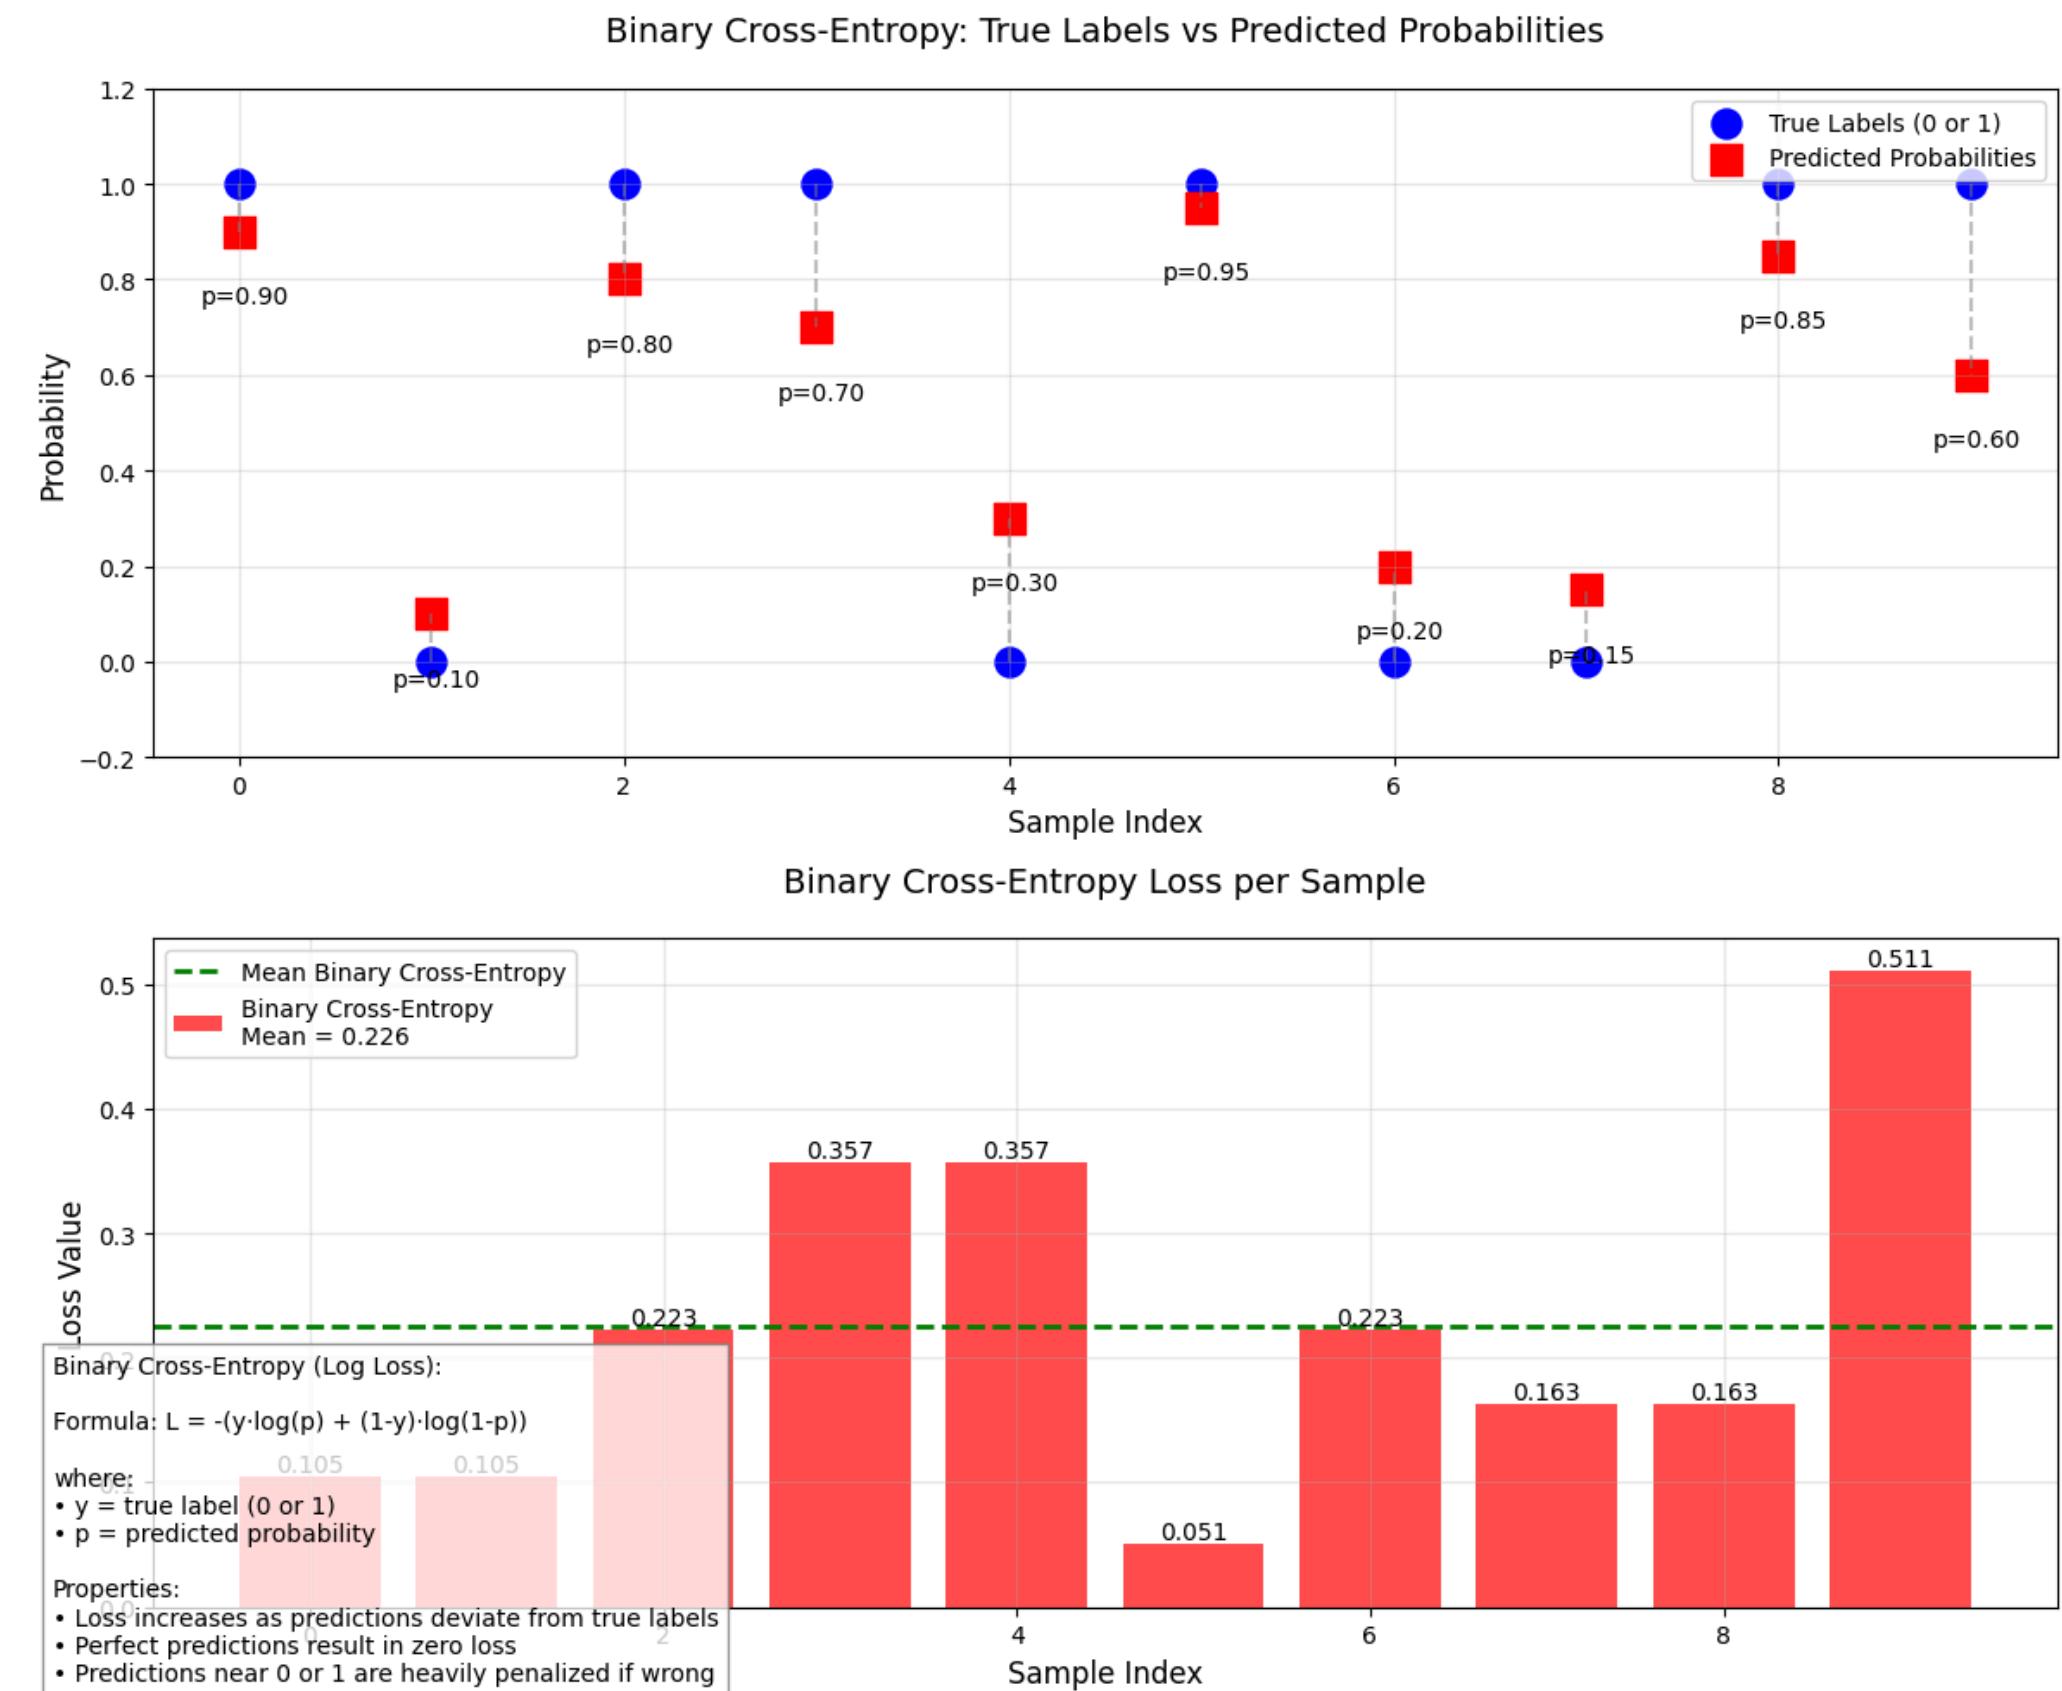
$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y})^2$$



# Binary Cross-Entropy (Log Loss)

- Used for binary classification tasks, where the output is either 0 or 1 (e.g., "dog" vs. "not dog").
- The predicted output is treated as a probability (between 0 and 1) of belonging to class 1
- The loss penalizes incorrect predictions more heavily when the model is very confident but wrong.

$$BCE = \frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

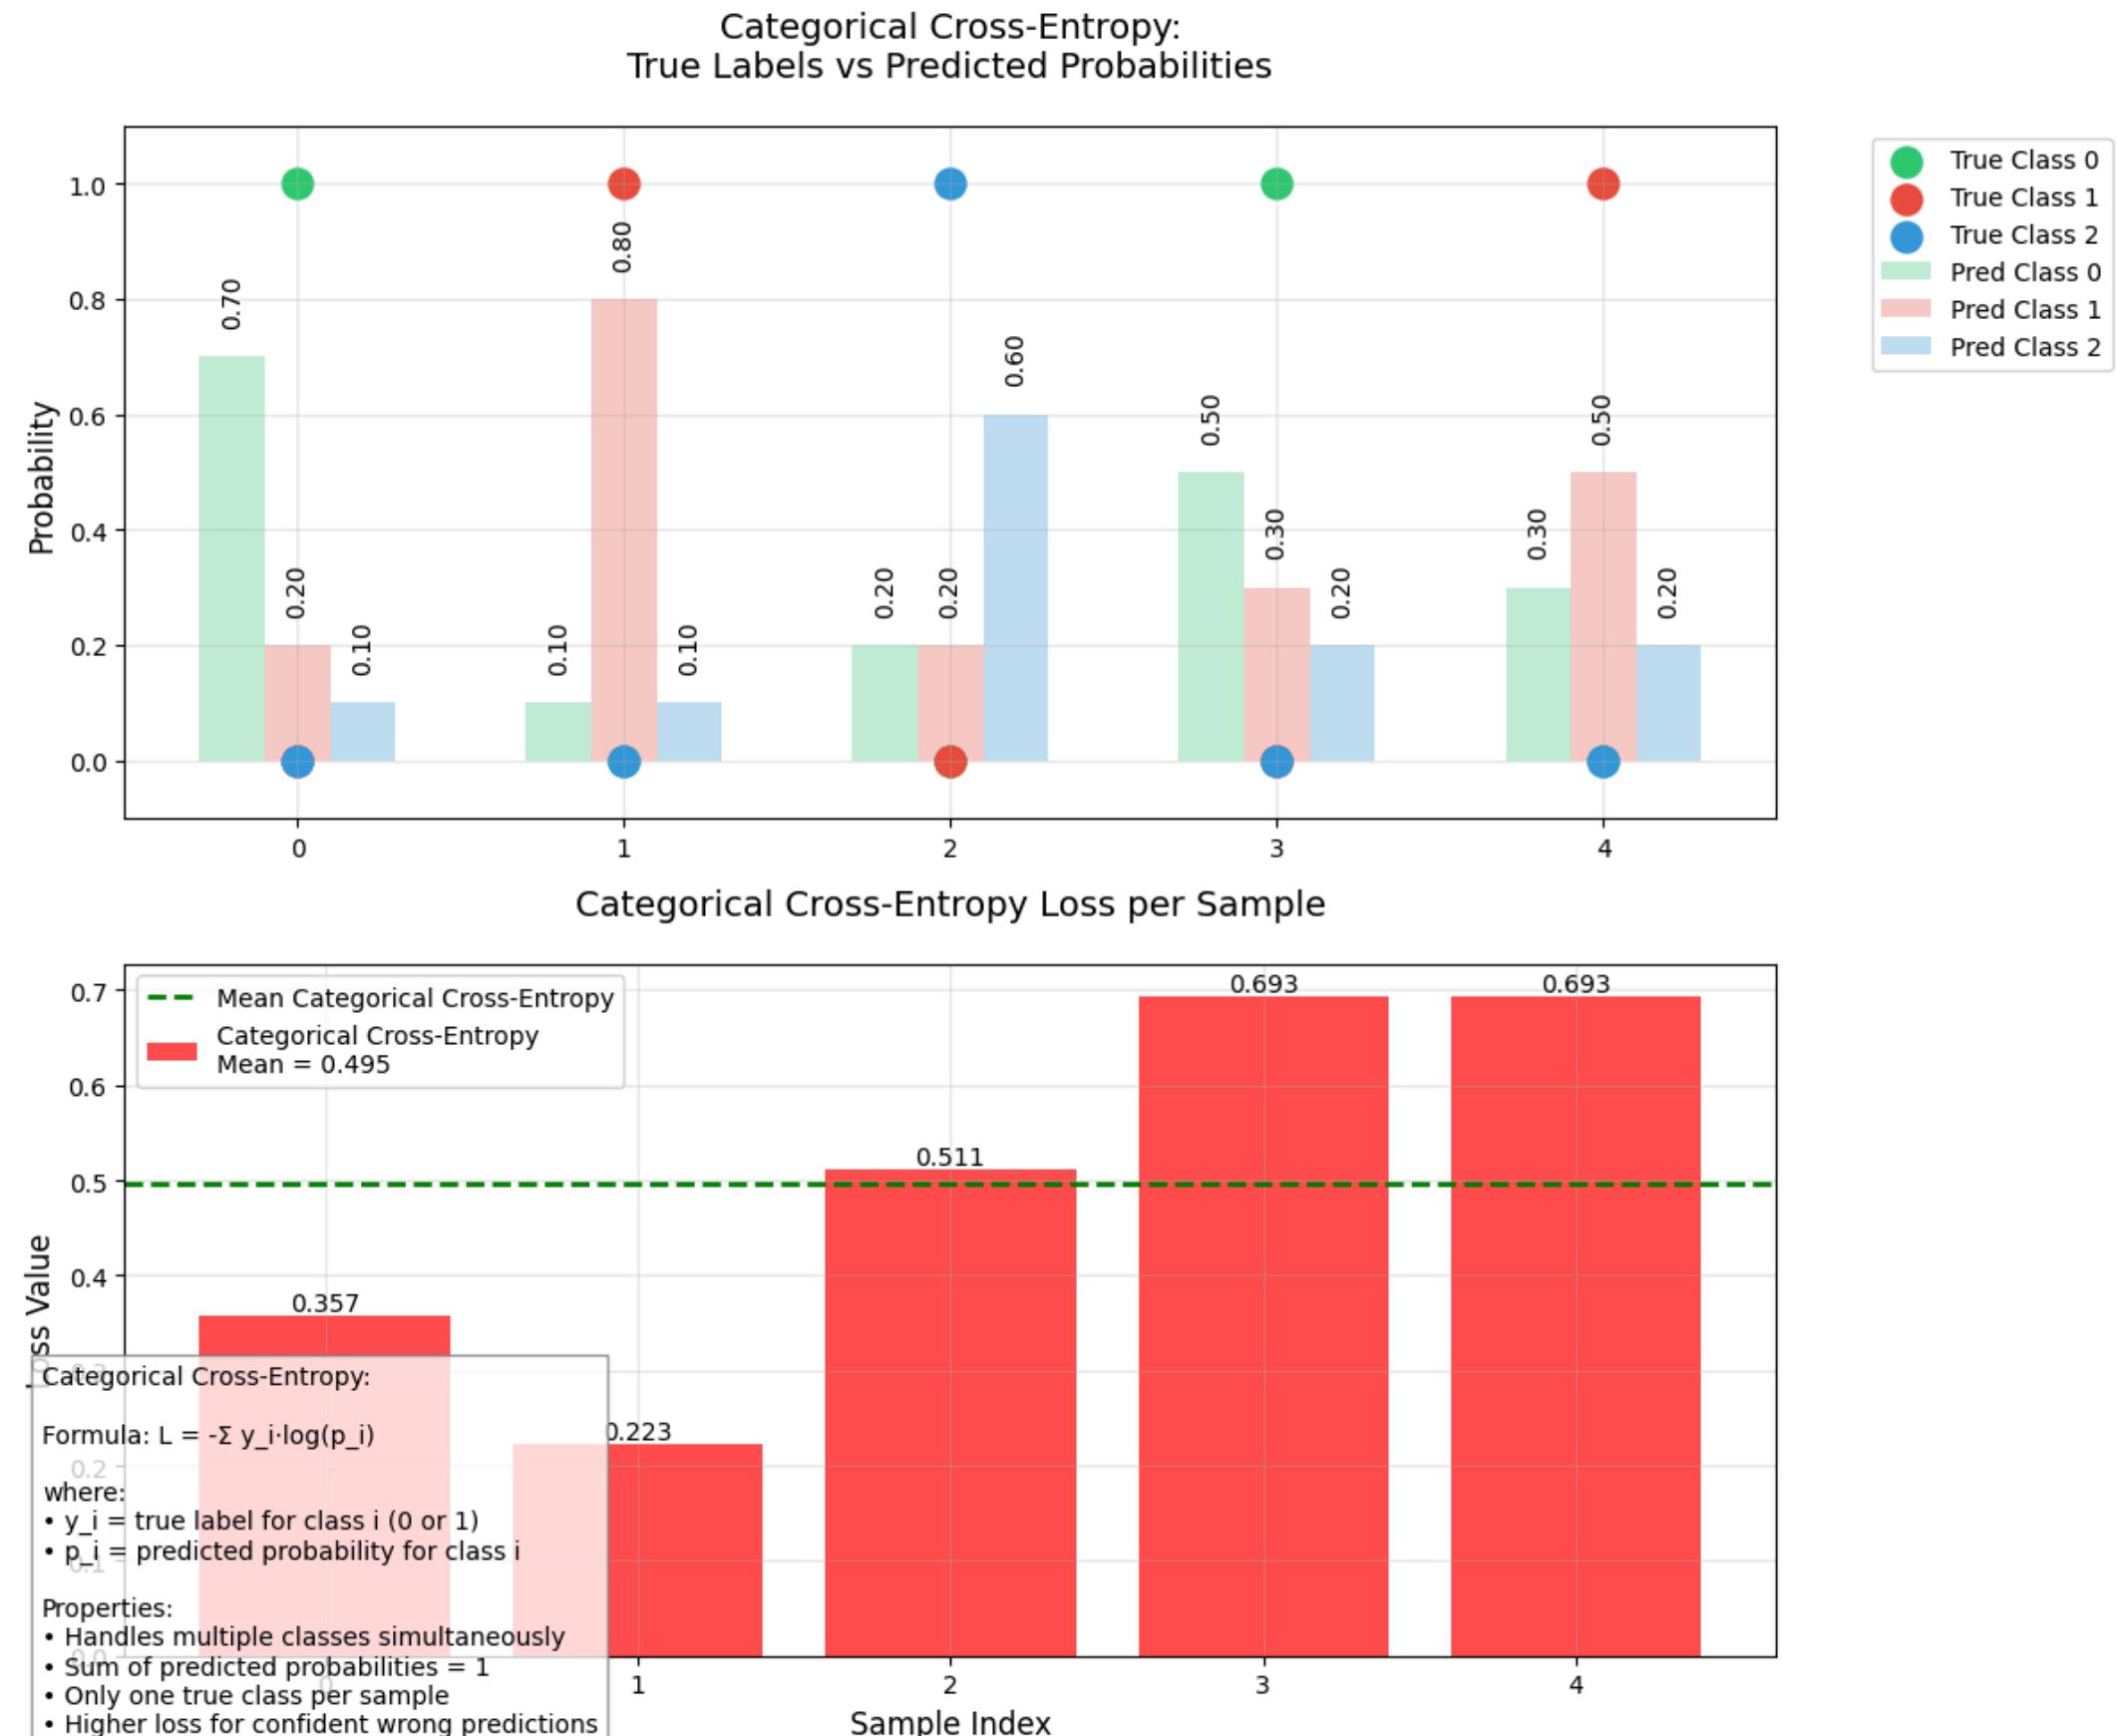


# Categorical Cross-Entropy

- Used for multi-class classification tasks, where the output can belong to one of  $k$  classes (e.g., dog, cat, bird).
- Each output neuron corresponds to a class.
- For each example, only the predicted probability of the correct class contributes to the loss

$$CCE = -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k y_{ij} \log(\hat{y}_{ij})$$

$$\hat{y}_{ij} = \frac{e^{z_j}}{\sum_{l=1}^k e^{z_l}}$$

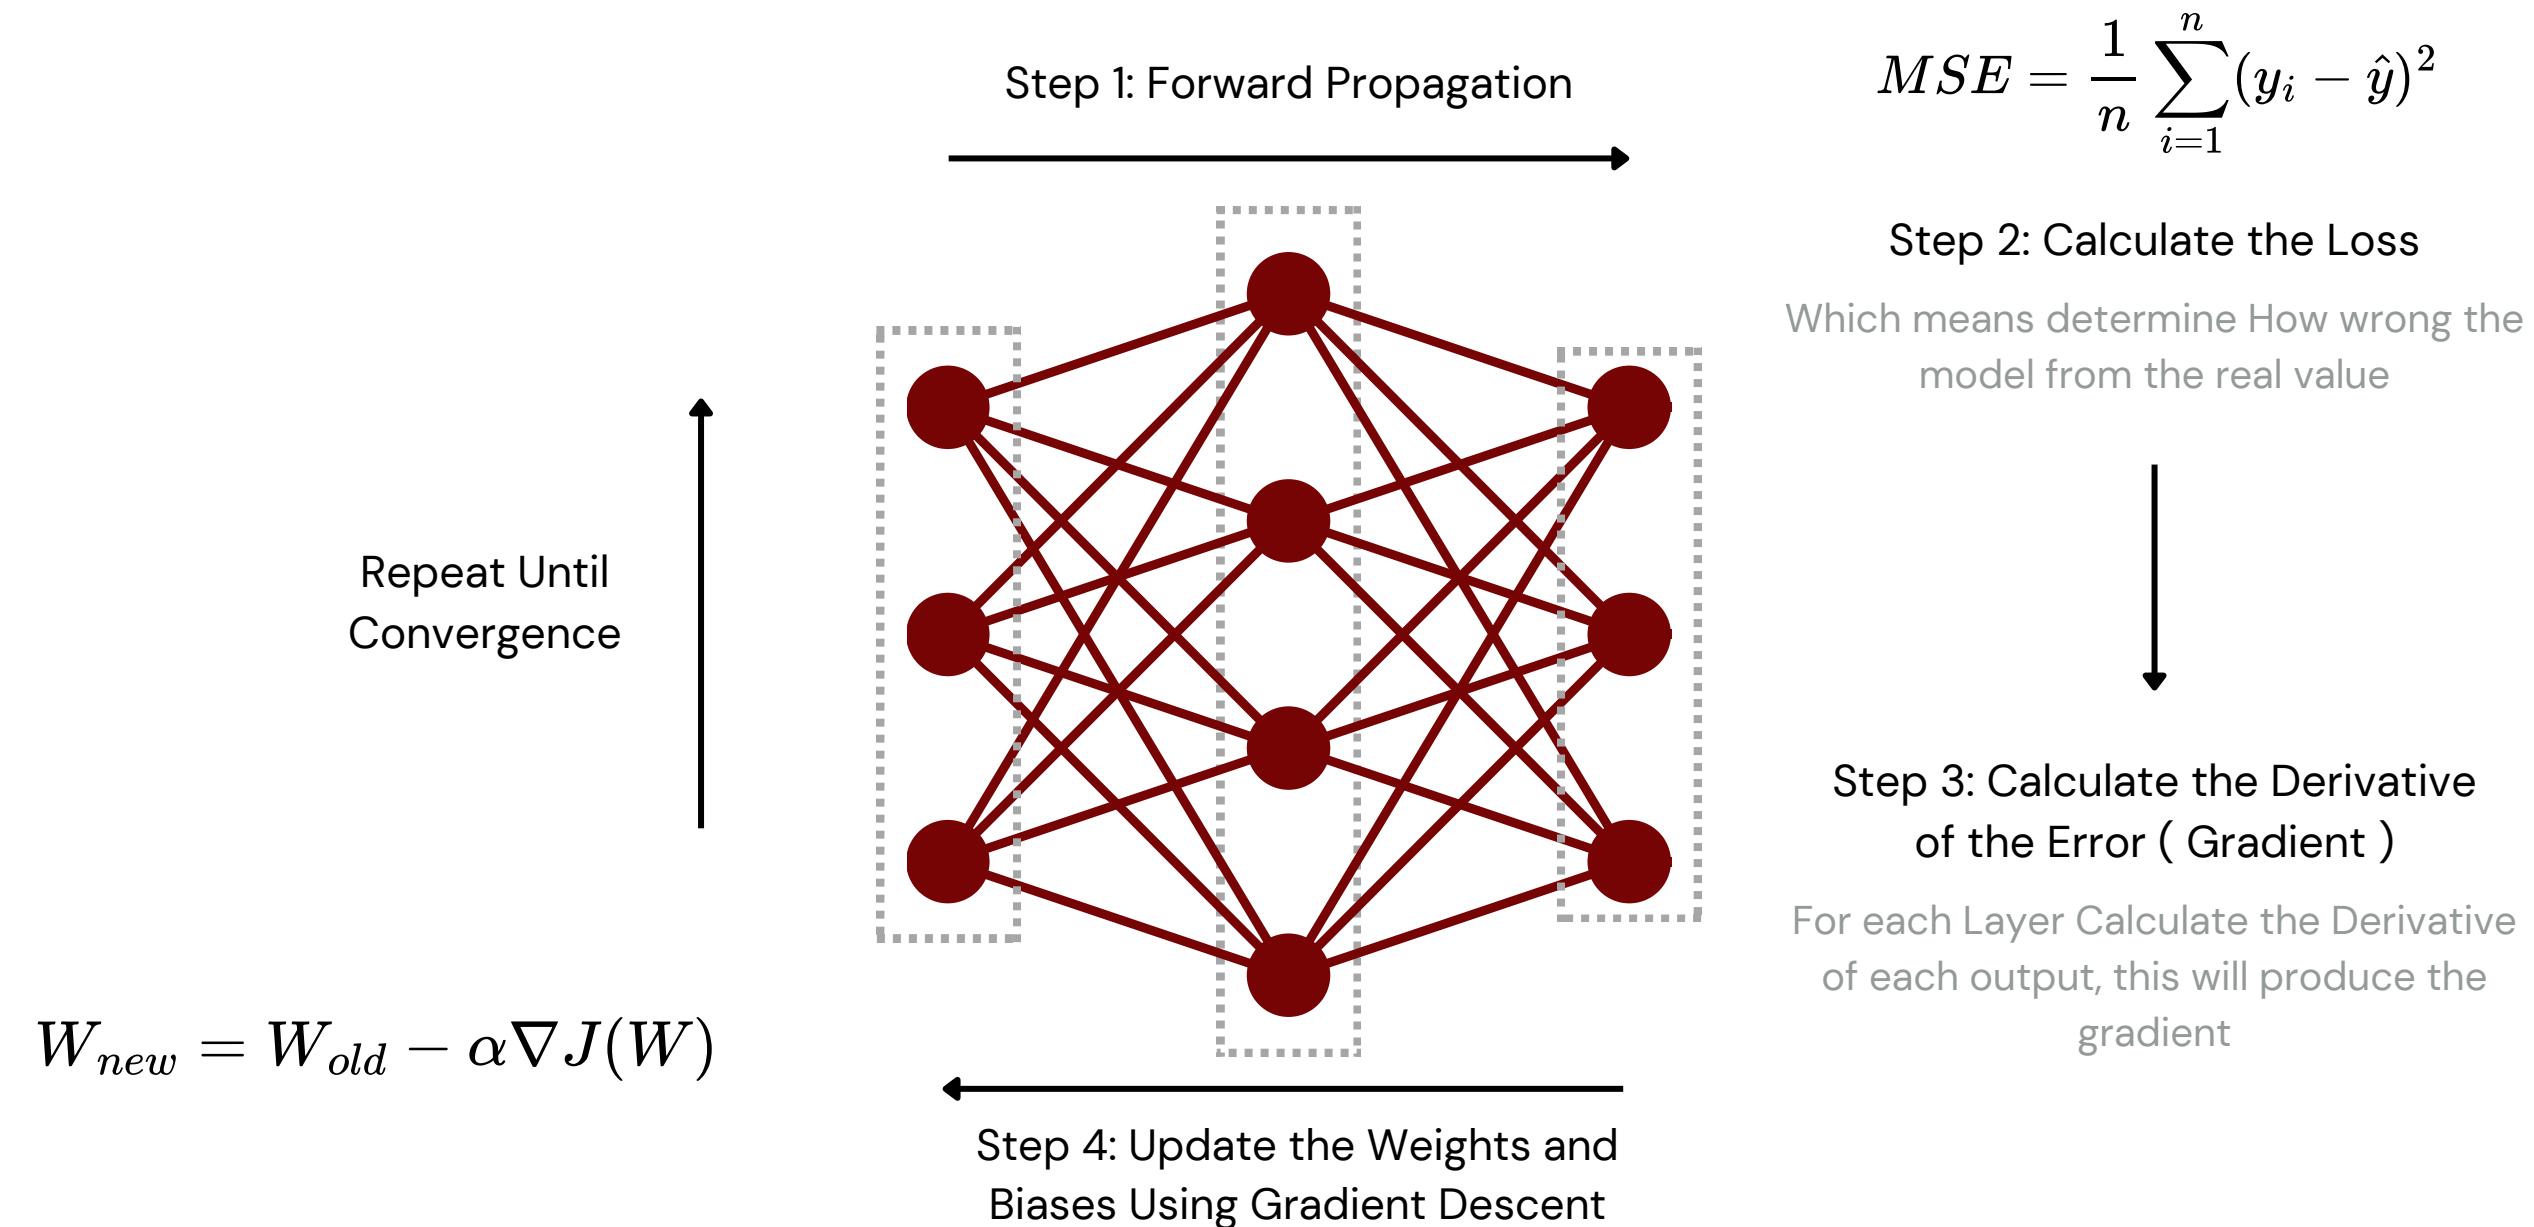


# Let's implement a Cost function

scan the notebook

# Where the magic unfolds: Backpropagation

- Backpropagation is a process that helps a neural network learn by adjusting the weights and biases in the right direction to reduce the error (loss).
- The goal of backpropagation is to minimize this error by tweaking the weights and biases.



# Backpropagation

## Forward Pass

- The input data passes through the network, layer by layer.

## Calculate the Error

- Using MSE, compute how far the network's prediction is from the actual value

## Backward Pass

- Starting with the output layer, we calculate **how the error changes with respect to the predicted value**

$$\frac{\partial \text{Error}}{\partial \hat{y}} = \hat{y} - y$$

- Work backward through the layers; for each weight in the network, calculate how much it contributed to the error using the **chain rule**

$$\frac{\partial \text{Error}}{\partial w} = \frac{\partial \text{Error}}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z} \cdot \frac{\partial z}{\partial w}$$

from the previous step                      Derivative of the activation function              How the weight affects the weighted sum

- Update weights by using the gradient.

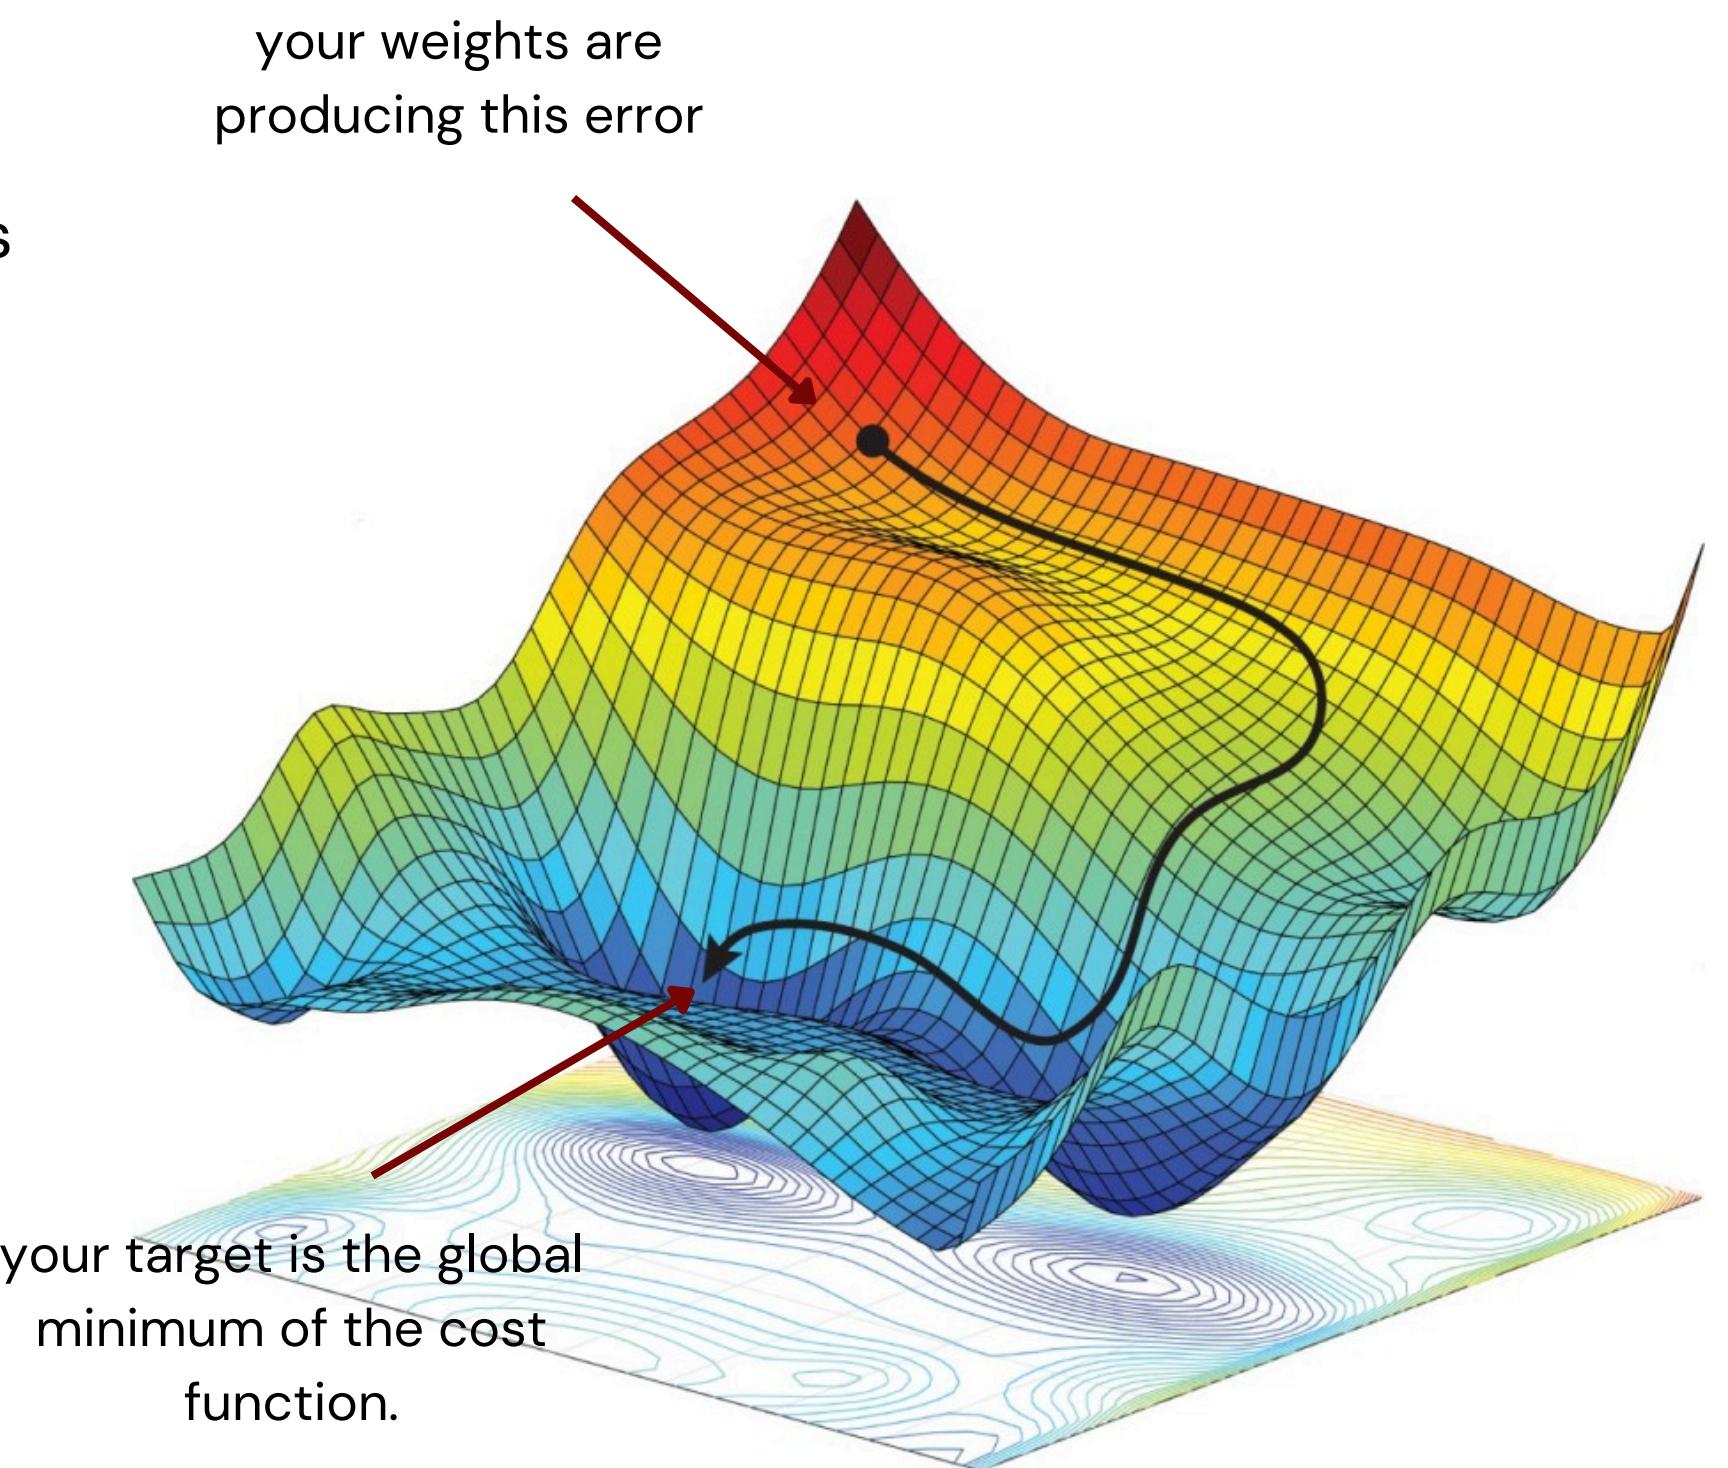
$$w_{\text{new}} = w_{\text{old}} - \alpha \cdot \frac{\partial \text{Error}}{\partial w}$$

The alpha is a hyperparameter called learning rate.

# Let's implement a Back Propagation

# Introducing Optimization

- Optimization is the process of adjusting the weights and biases of a neural network so that it makes better predictions.
- After calculating the error (loss) during training, optimization algorithms decide how much to change the weights and biases to reduce the error.
- This process uses the gradients calculated during backpropagation.
- Optimizers are algorithms or methods used to update the weights and biases in a neural network.



# Optimizers

## Gradient Descent

- Updates weights by moving in the direction of the negative gradient of the loss function.

## Stochastic Gradient Descent (SGD)

- Instead of using all the training data to compute the gradient, it uses one sample at a time.
- Faster but noisier updates.

## Mini-Batch Gradient Descent

- Combines both: it computes gradients using small batches of data instead of the entire dataset or just one sample.
- Balances speed and stability.

## Adam (Adaptive Moment Estimation)

- Most popular optimizer.
- Combines the ideas of Momentum and another method called RMSProp.
- Automatically adjusts the learning rate for each parameter based on the gradient's history.

*Equation 11-9. Adam algorithm*

- $\mathbf{m} \leftarrow \beta_1 \mathbf{m} - (1 - \beta_1) \nabla_{\theta} J(\theta)$
- $\mathbf{s} \leftarrow \beta_2 \mathbf{s} + (1 - \beta_2) \nabla_{\theta} J(\theta) \otimes \nabla_{\theta} J(\theta)$
- $\widehat{\mathbf{m}} \leftarrow \frac{\mathbf{m}}{1 - \beta_1^t}$
- $\widehat{\mathbf{s}} \leftarrow \frac{\mathbf{s}}{1 - \beta_2^t}$
- $\theta \leftarrow \theta + \eta \widehat{\mathbf{m}} \oslash \sqrt{\widehat{\mathbf{s}}} + \epsilon$

# Let's implement an Optimizer

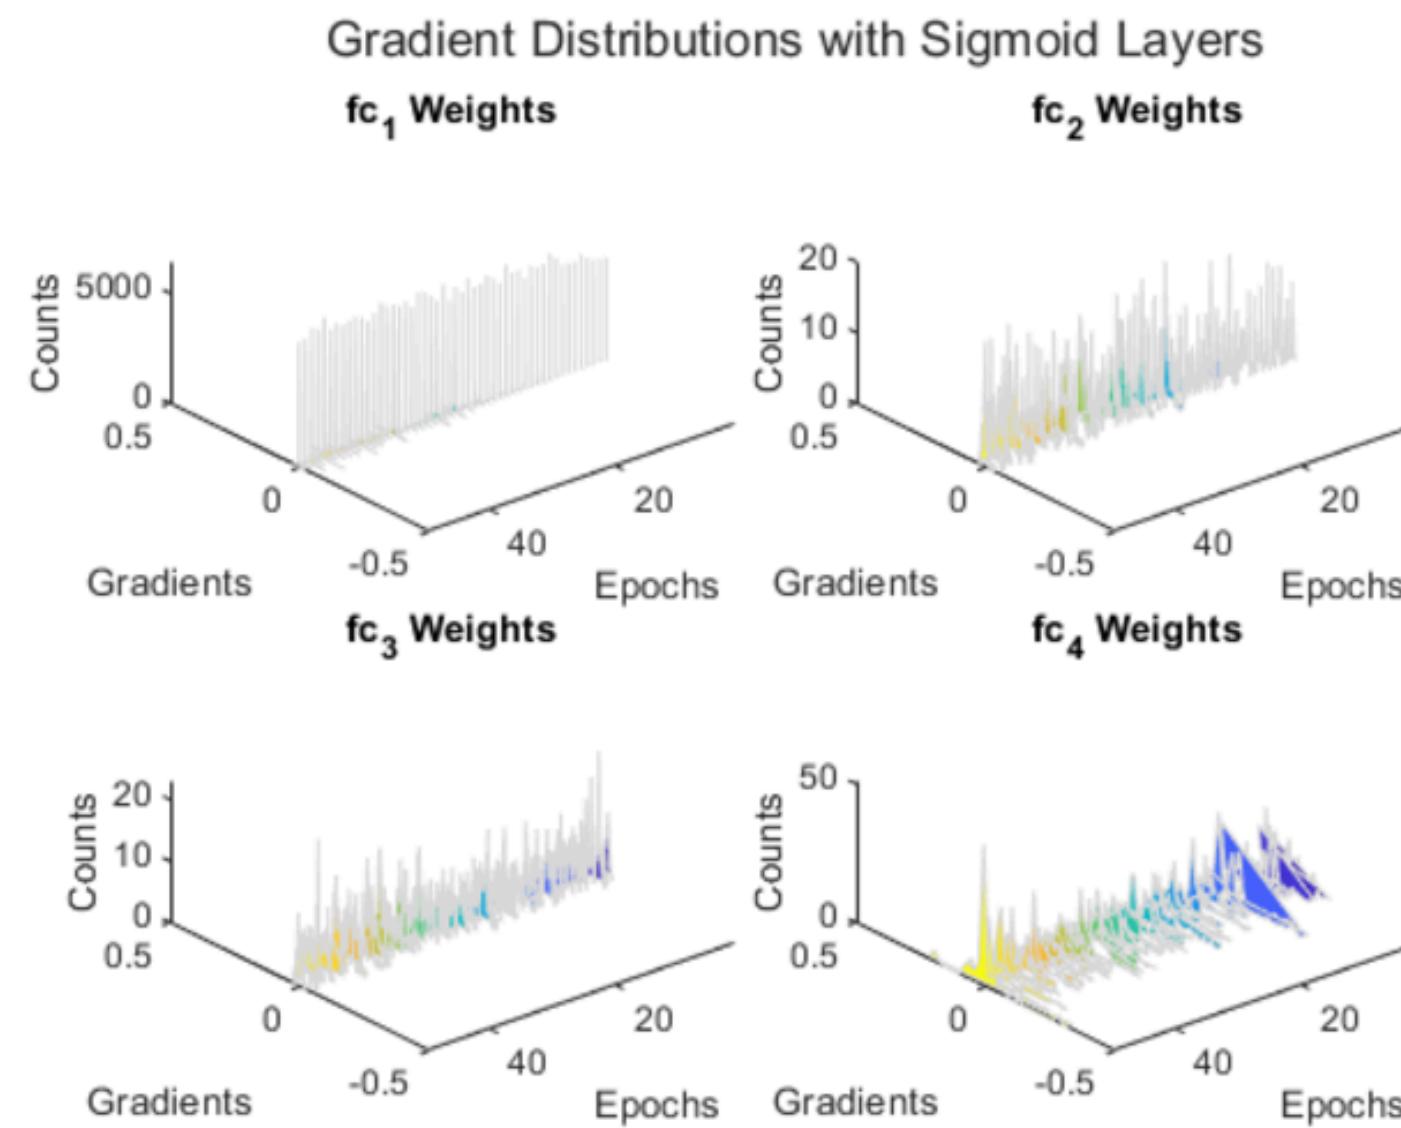
# Training Deep Neural Networks

Training deeper neural networks (with many layers and connections) introduces several problems

- Gradients may shrink (vanish) or grow (explode) during backpropagation, making training difficult.
- Not enough training data or costly labeling can be an issue.
- Large models risk overfitting, especially with noisy or limited data.

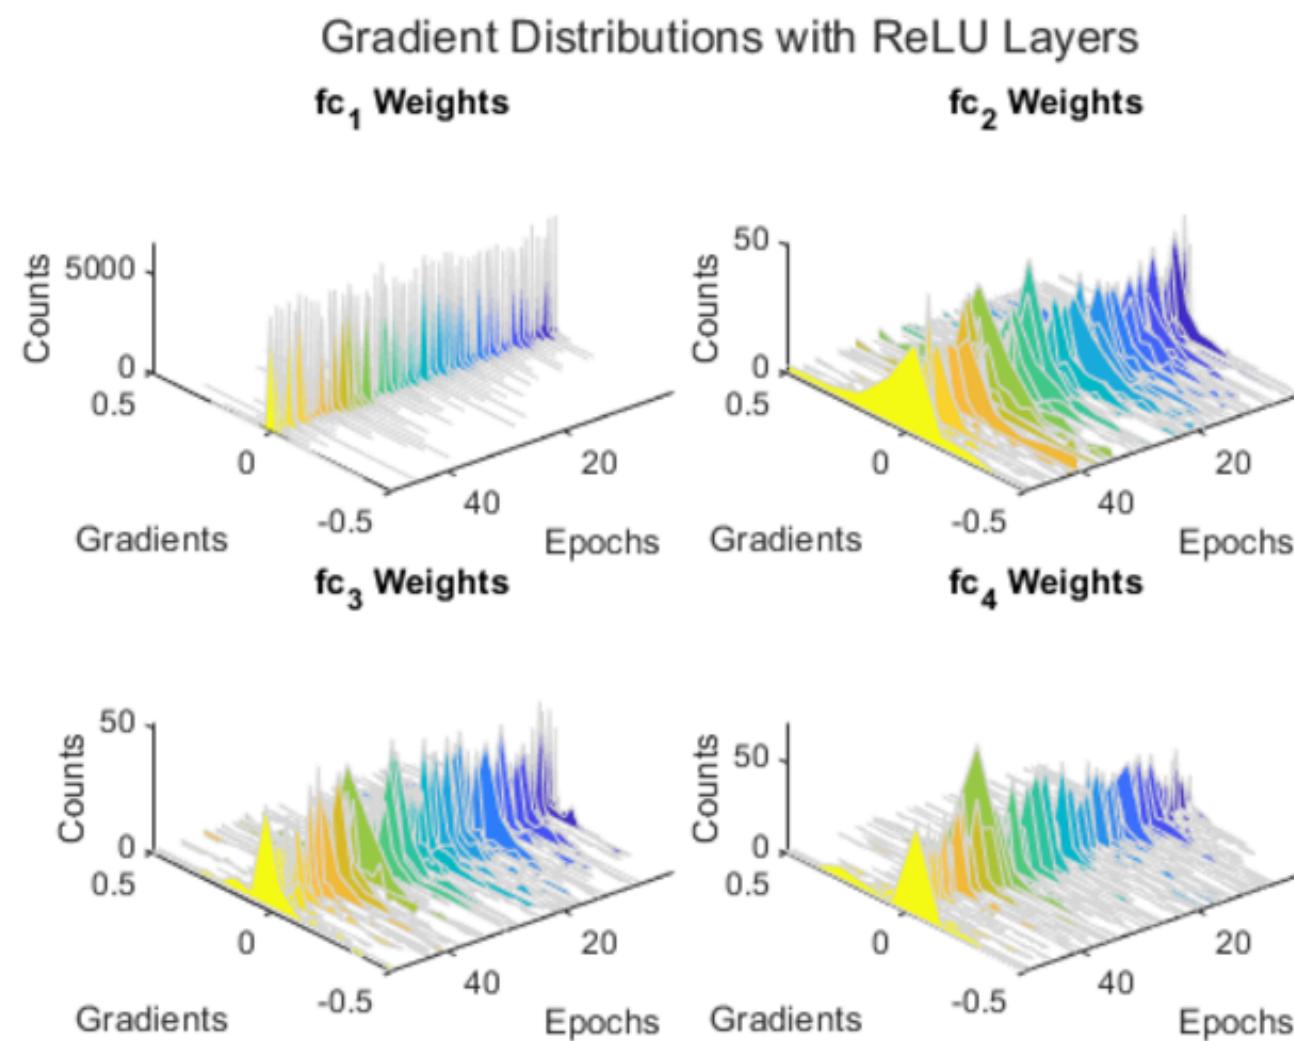
# The Vanishing Gradients Problem

The vanishing gradients problem arises when gradients become very small as they move backward through layers during training, hindering effective learning in the network.



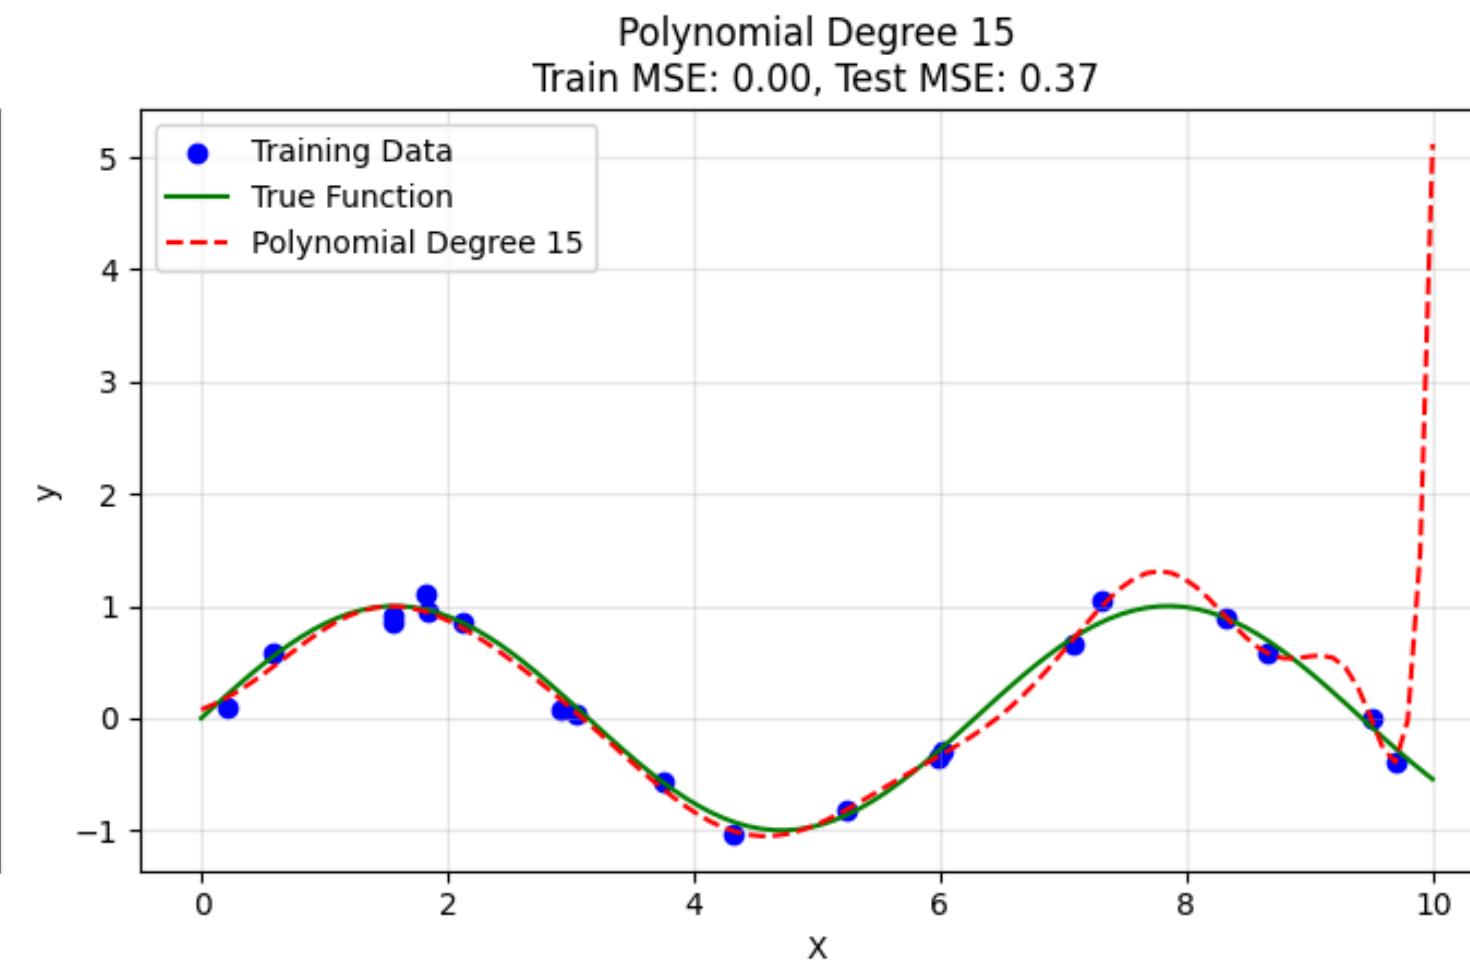
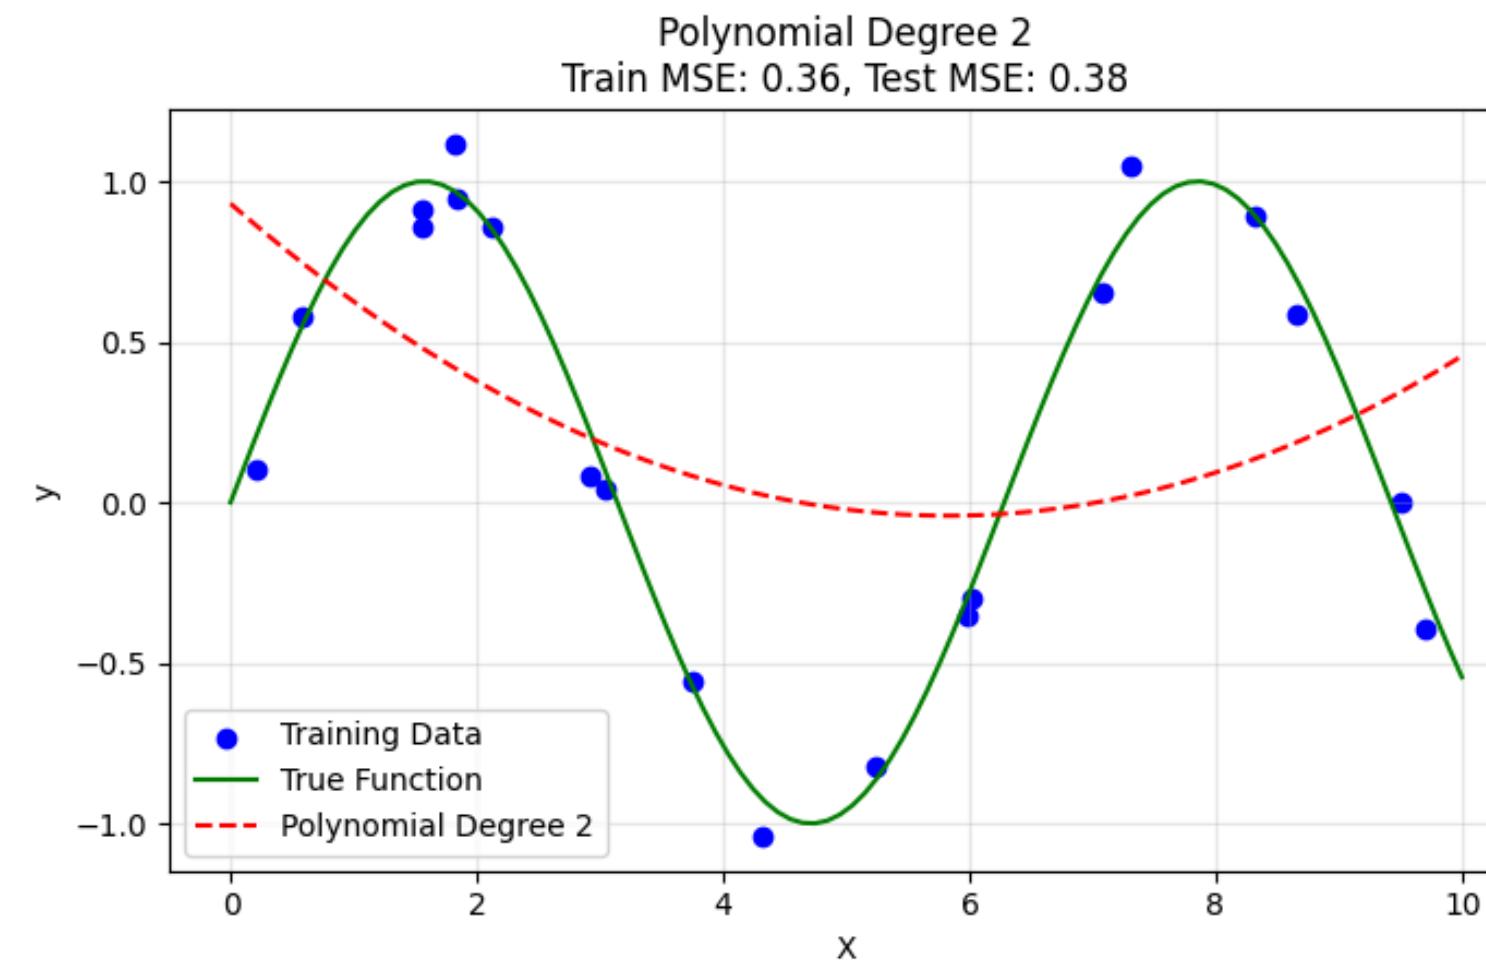
# The Vanishing Gradients Problem

- Use activation functions like ReLU (Rectified Linear Unit), which doesn't squash values, allowing gradients to flow better.
- Use techniques like batch normalization or proper weight initialization to keep the gradients healthy during training.



# Overfitting Problem

- Overfitting occurs when a model learns training data too well, failing to generalize to new, unseen data.
- Performs very well on the training data but poorly on the test or validation data.



# Overfitting Problem

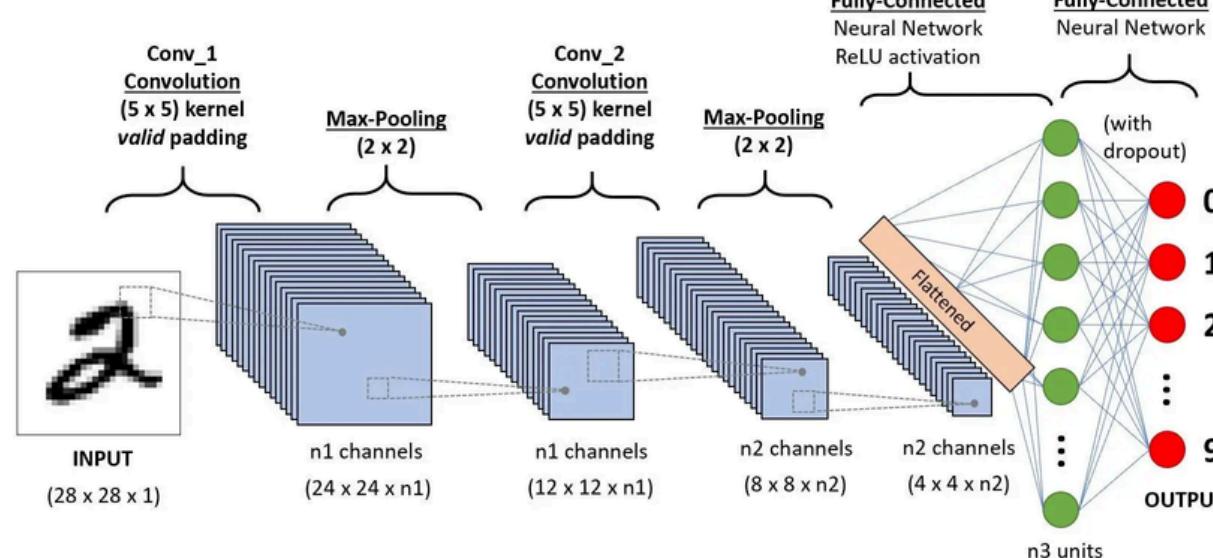
## Solutions

- More data can help the model learn general patterns and reduce the chance of memorizing specific details or noise.
- Use a smaller model (fewer layers or neurons) to prevent the network from learning irrelevant details.
- Stop training early before the model starts to overfit.
- During training, randomly "drop" (set to zero) some of the neurons in each layer.
- Use cross-validation to evaluate the model on various data sets, helping to identify overfitting early.

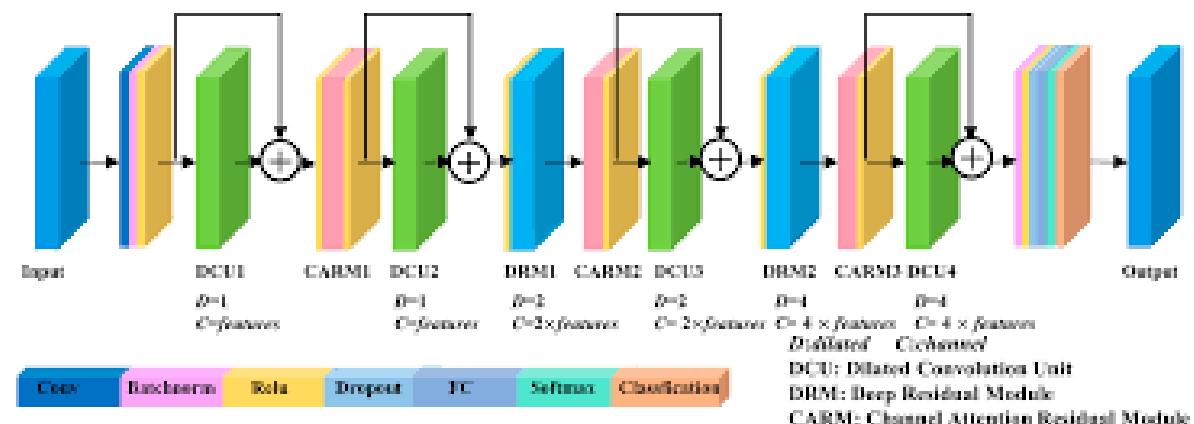
# Starting A Real Deep Learning Project



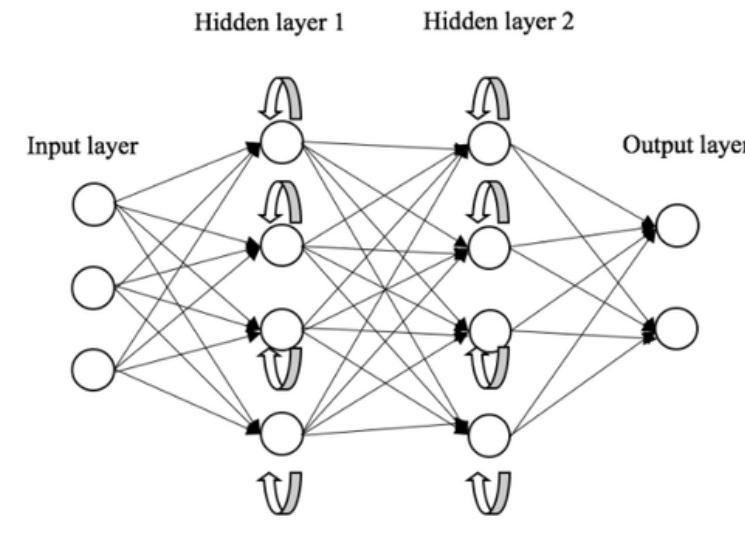
# The world of deep learning is vast!



Convolutional Neural Networks



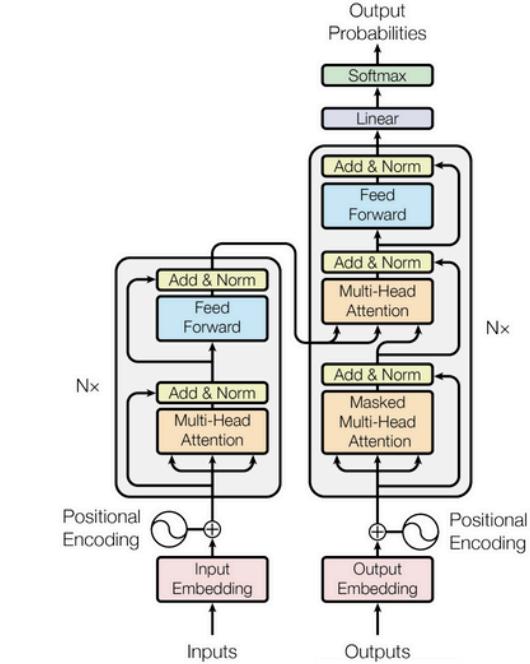
Residual Neural Networks



Recurrent Neural Networks

BERT

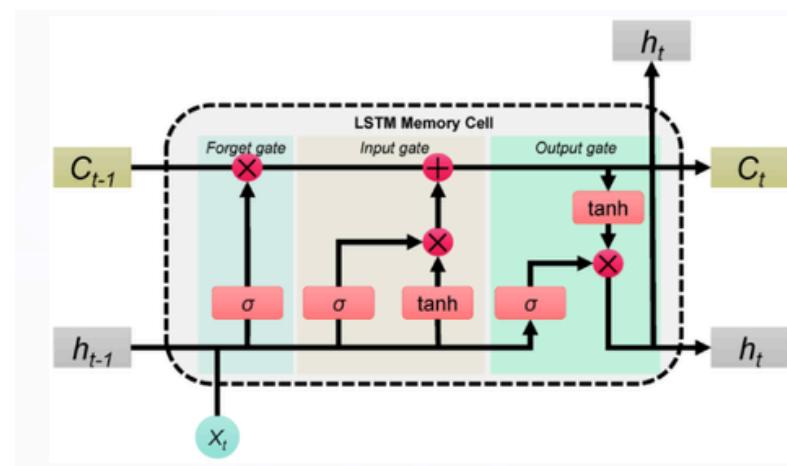
Encoder



GPT

Decoder

Transformer Architecture



LSTM's

And much more....

# Dive deeper

Stanford CS229  
Machine Learning

+

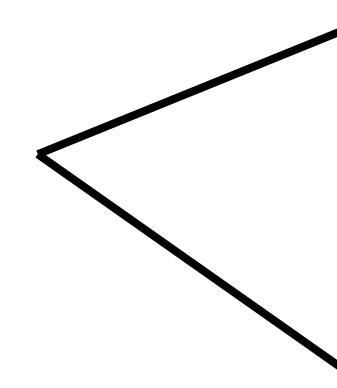
Andre NJ Coursera  
Machine Learning  
Specialization



Stanford CS230 Deep  
Learning

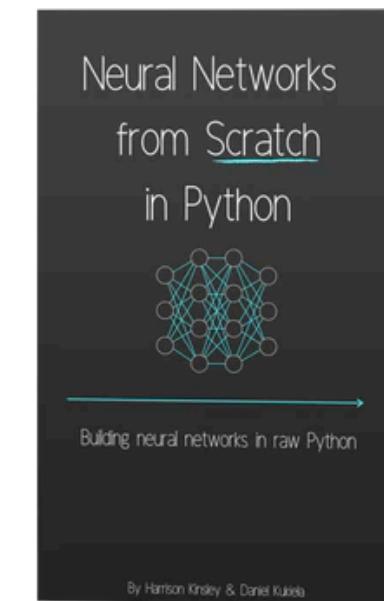
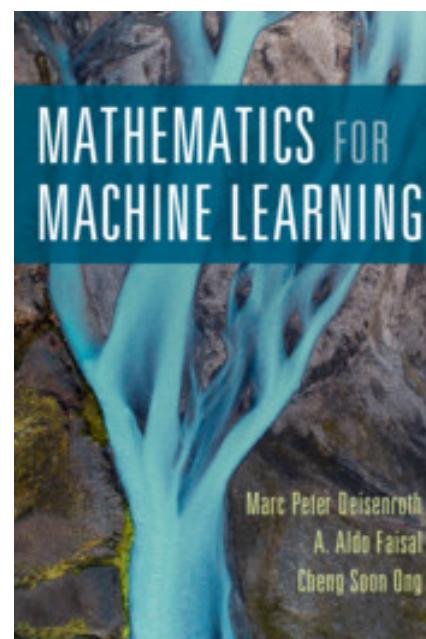
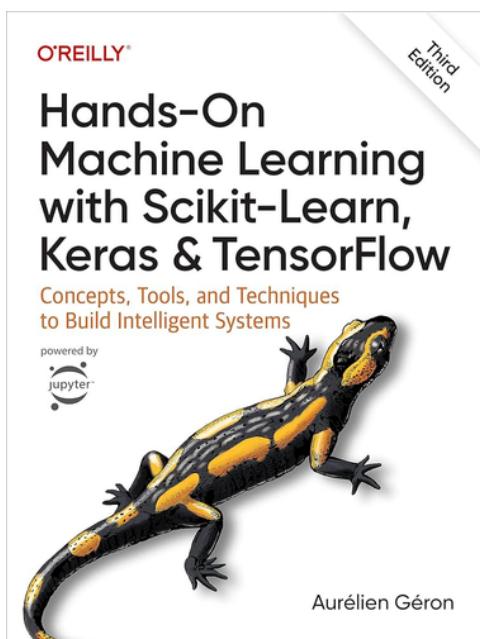
+

Andre NJ Coursera  
Deep Learning  
Specialization



Stanford CS231n  
Computer Vision

Stanford CS224n  
Natural Language  
Processing with Deep  
Learning



The internet

Thank you